

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

DNS: domain name system

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address - used for addressing datagrams
 - 4 bytes or 32 bits
 - e.g., 129.23.4.51
 - used for routing
- “name”,
 - e.g., www.uci.edu
 - used by humans
 - variable length

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

- *distributed database*
 - implemented in hierarchy of many *name servers*
- *application-layer protocol:*
 - hosts, name servers communicate to *resolve* names (address/name translation)
 - runs on top of UDP, port 53
 - note: core Internet function, implemented as application-layer protocol
 - complexity at network's “edge”

DNS services

- hostname to IP address translation
- host aliasing: canonical vs. alias names
- mail server aliasing
- load distribution
 - replicated Web servers: set of IP addresses for one canonical name
 - rotating

nslookup (or host, dig, whois) athina.calit2.uci.edu

Name: athina.calit2.uci.edu
Address: 128.195.177.83

nslookup (or dig) www.cnn.com

www.cnn.com canonical name = www.cnn.com.vgtf.net
www.cnn.com.vgtf.net canonical name = cnn-56m.gslb.vgtf.net.
Name: cnn-56m.gslb.vgtf.net
Address: 157.166.249.11
Name: cnn-56m.gslb.vgtf.net
Address: 157.166.248.10

Nslookup -type=mx stanford.edu

Stanford.edu mail exchanger = 40 mx1.stanford.edu.
stanford.edu mail exchanger = 20 mx2.stanford.edu.
stanford.edu mail exchanger = 20 mx3.stanford.edu.

nslookup (or dig) google.com

Name: google.com
Address: 74.125.227.167
Name: google.com
Address: 74.125.227.168
Name: google.com
Address: 74.125.227.169

.....


DNS structure

why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!


DNS: a distributed, hierarchical database




client wants IP for www.amazon.com; 1st approximation:

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

DNS: a distributed, hierarchical database



DNS: a distributed, hierarchical database

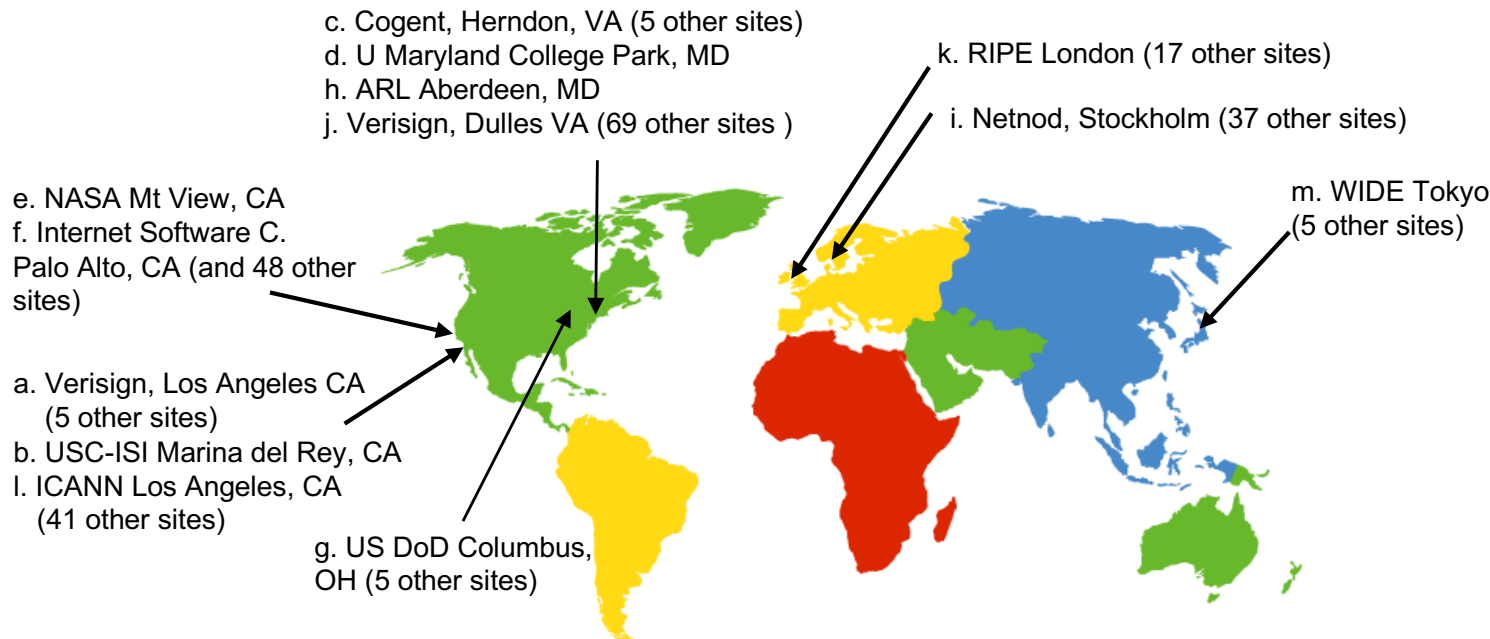


client wants IP for www.amazon.com; 1st approximation:

- client queries **root** server to find com DNS server
- client queries **com** DNS server to get amazon.com DNS server
- client queries **amazon.com** DNS server to get IP address for www.amazon.com

DNS: root name servers

- 13 root name servers worldwide: a, b...m
 - in fact replicated: 247 root servers as of 2011
 - <https://www.iana.org/domains/root/servers>
- contacted by local name server that can not resolve name
- root name server:
 - contacts authoritative name server if name mapping not known
 - gets mapping
 - returns mapping to local name server



TLD, authoritative servers

top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
 - <http://www.iana.org/domains/root/db>
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD: <http://whois.educause.edu>

authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

Local DNS name server

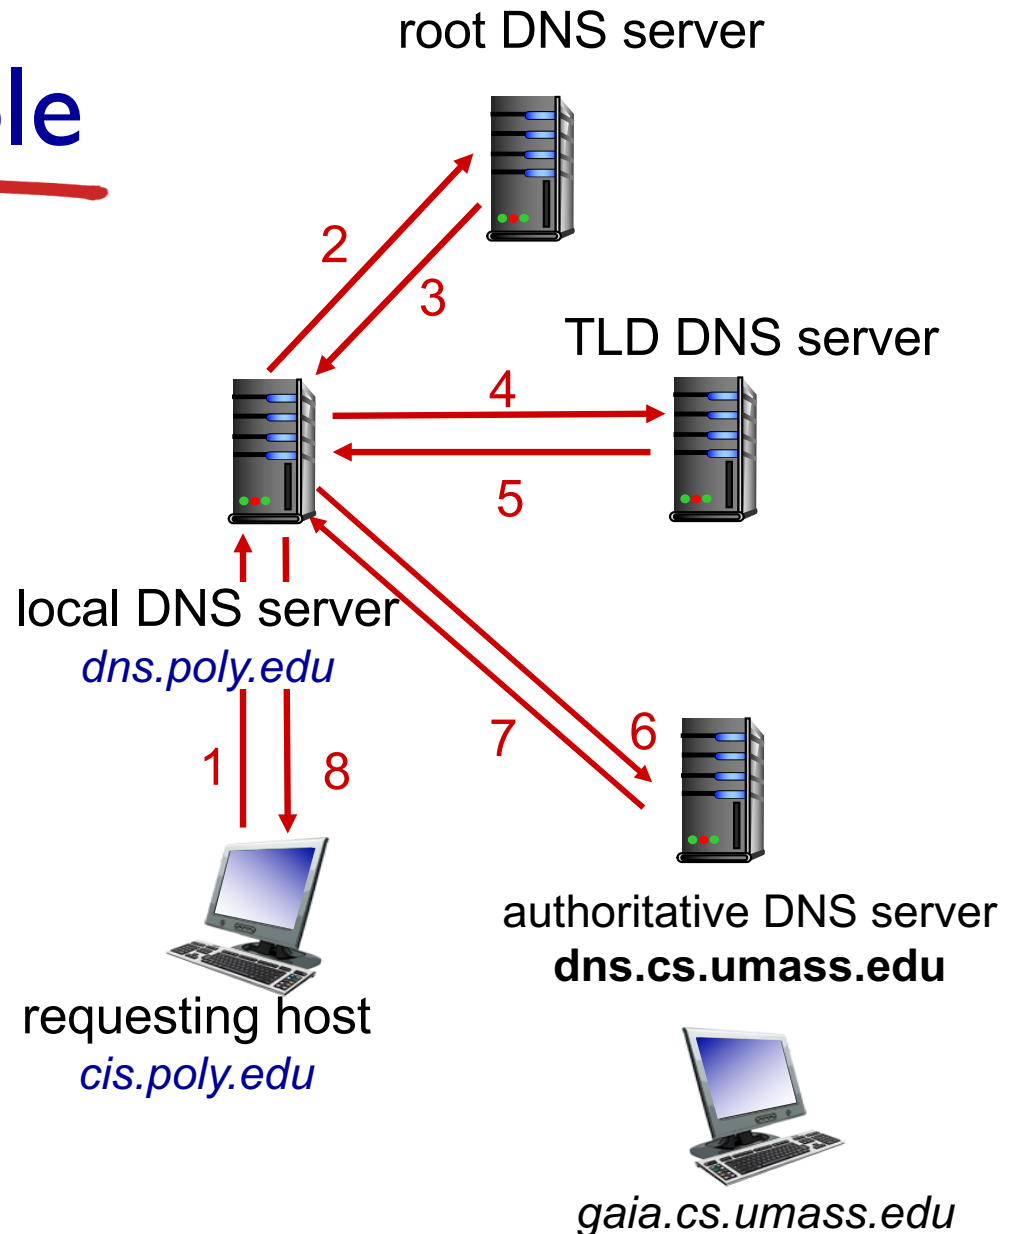
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
- also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
 - has local cache of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy
 - forwards query into hierarchy
 - caches records
- Ex: `more /etc/resolv.conf`

DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:

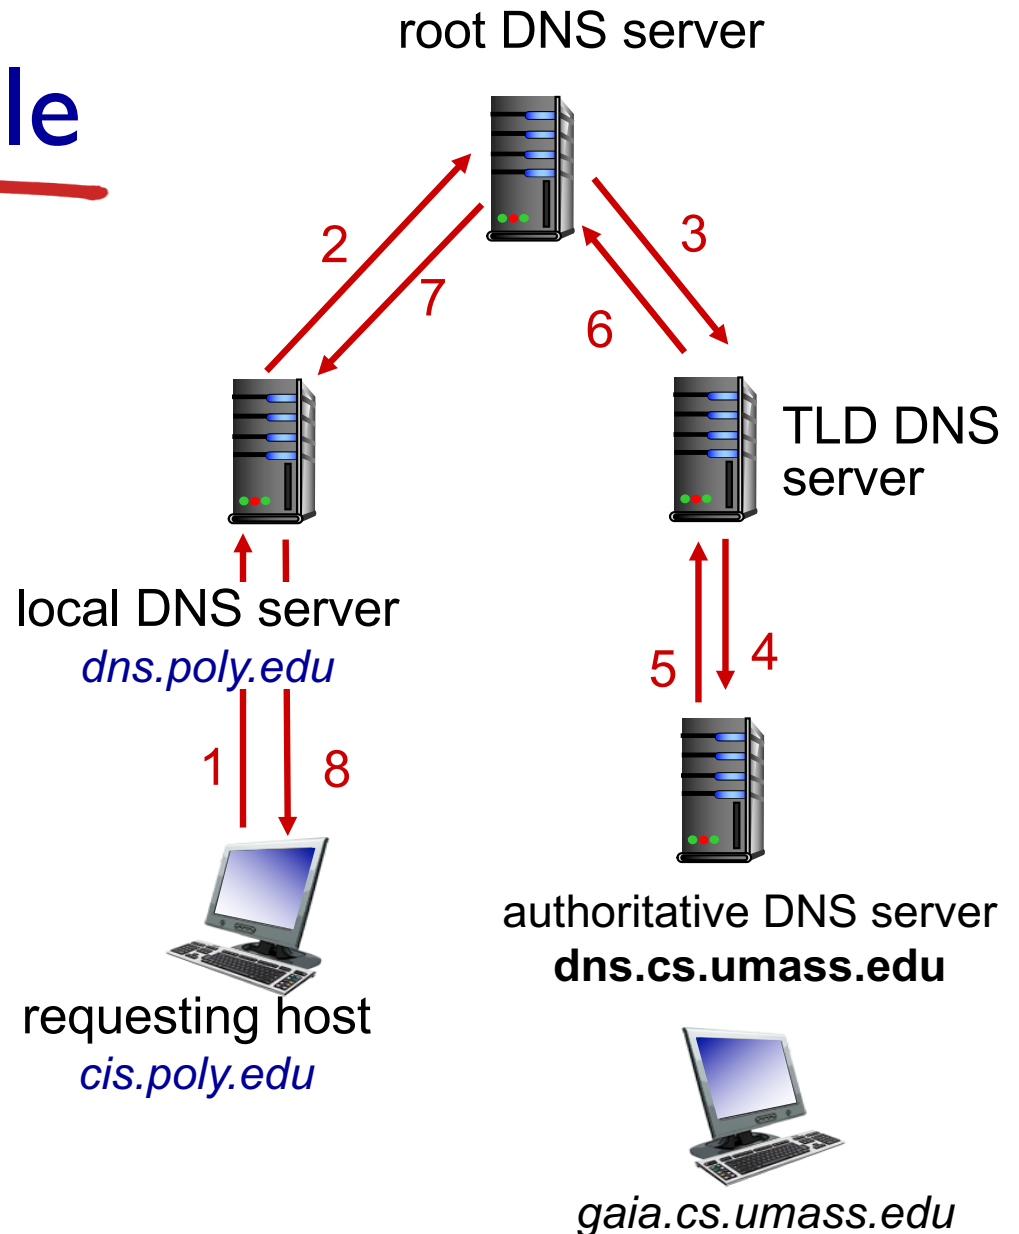
- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



DNS: caching, updating records

- once (any) name server learns a mapping, it *caches it*
 - cache entries timeout (disappear) after some time (TTL)
 - Time-to-live (TTL) by default is 2 days
 - Needed because records change often
 - TLD servers typically cached in local name servers
 - thus root name servers are not visited often
- cached entries may be *out-of-date* (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- How to configure the records in the database
 - statically
 - update/notify mechanisms RFC 2136

DNS records -Summary

DNS: distributed database storing resource records (**RR**)

RR format: (name, value, type, ttl)

type=A

- **name** is hostname
- **value** is IP address

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

type=MX

- **value** is name of mailserver associated with **name**

DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

Type=A

- **name** is hostname
- **value** is IP address
- Stored at authoritative server of that domain

Example

- (**odysseas.calit2.uci.edu**, **128.195.185.112**, **A**)
 - You can lookup this info (both directions)
 - by command line, e.g.: nslookup or dig or host
 - or on the web, e.g.
 - <http://www.kloth.net/services/nslookup.php>
 - <http://www.iana.org/domains/root/db>
 - <http://whois.educause.edu>
 - or in the old days: gethostbyname(), gethostbyaddr()

DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

Type=NS

- name is domain (e.g., foo.com)
- value is hostname of authoritative name server for this domain
- this record is used to route a request further

Example

- (uci.edu, ns1.service.uci.edu, NS)
 - type “nslookup -ty=ns uci.edu”
 - Or simply “nslookup uci.edu”

DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

Type=MX

- **value** is name of mailserver associated with **name**

Example

- (uci.edu, mta.service.uci.edu, **MX**)
 - type “nslookup -ty=mx uci.edu”
 - type “nslookup -ty=mx stanford.edu”
 - Can have multiple NS and MS records
 - several MX records, allow for load balancing

DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

Type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **value** is canonical name

Example

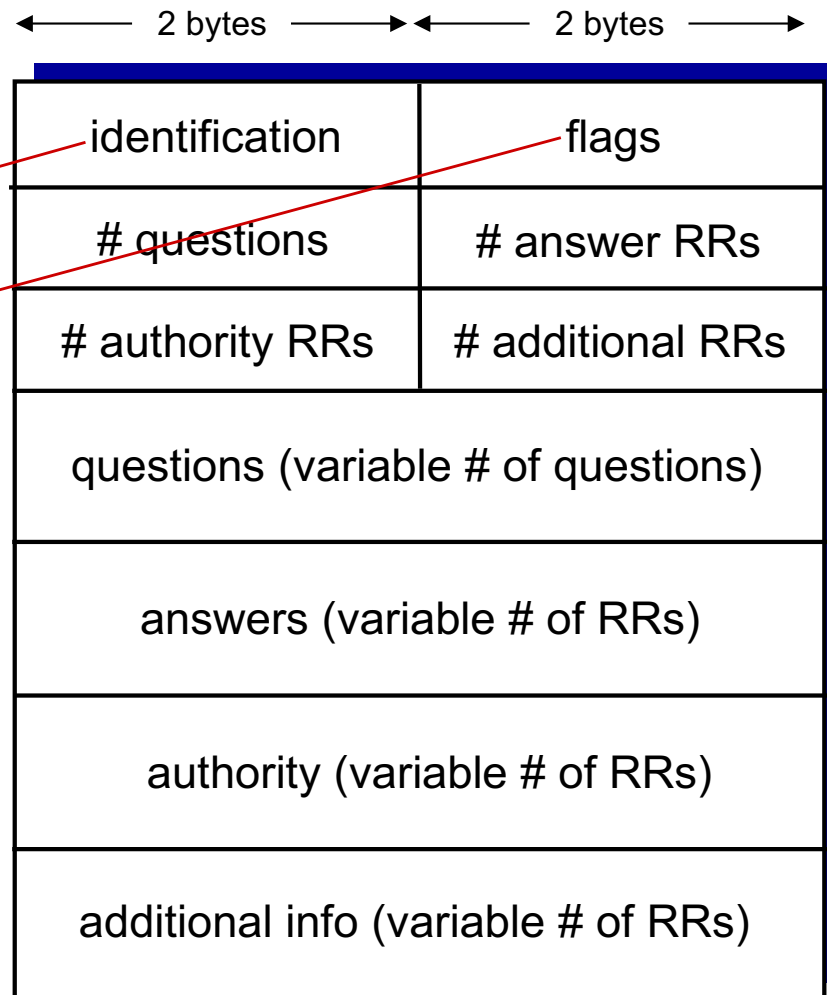
- (www.networkedsystems.uci.edu, odysseas.calit2.uci.edu, CNAME)
 - type “nslookup -type=cname www.networkedsystems.uci.edu”
 - Or simply “nslookup www.networkedsystems.uci.edu”
 - “nslookup -type=cname www.ibm.com”
 - alias, and potential for load balancing
 - a company can have the same alias for several servers...

DNS protocol, messages

- *query* and *reply* messages, both with same *message format*

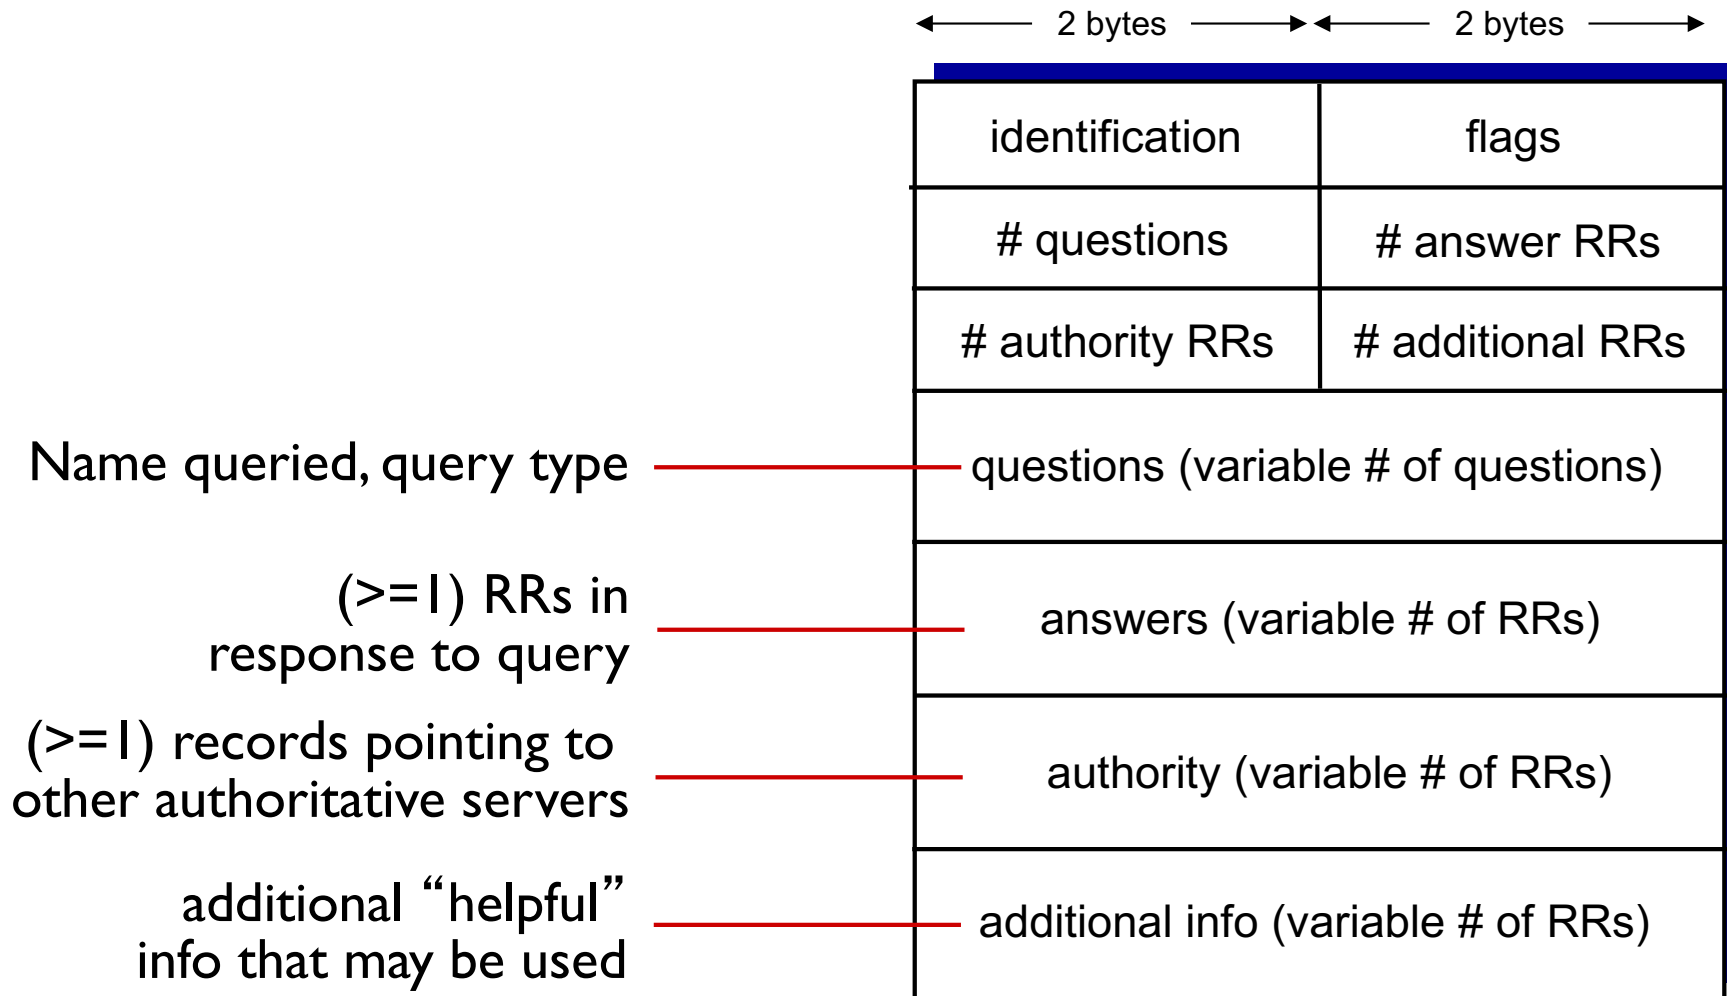
message header

- **identification**: 16 bit # for query, reply to query uses same #
- **flags**:
 - query or reply
 - recursion desired
 - recursion available
 - reply is authoritative



"dig" command returns those fields explicitly

DNS protocol, messages




Inserting records into DNS

- example: new startup “Network Utopia”
- register name networkutopia.com at *DNS registrar(*)*
 - provide names, IP addresses of authoritative name servers (primary and secondary)
 - registrar inserts two RRs into .com TLD server:
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server type A record for www.networkutopia.com; type MX record for networkutopia.com
- (*) E.g. Network Solutions for .com, see www.internic.net for registrars accredited by ICANN)

Inserting records into DNS

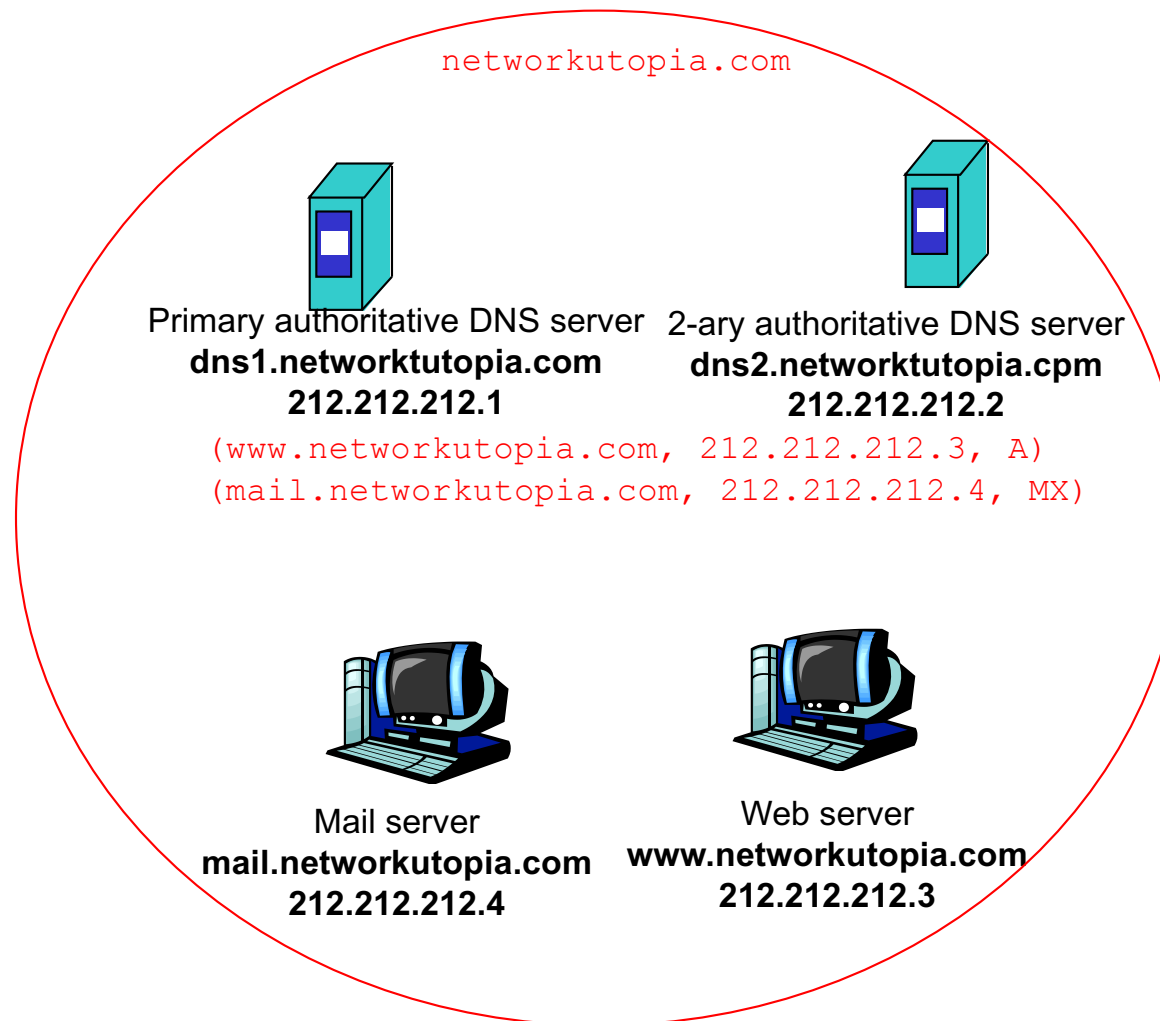
- Example: new startup “Network Utopia”
- Register name networkutopia.com at *DNS registrar* (e.g., Network Solutions, see www.internic.net for approved registrars by ICANN)
 - provide names, IP addresses of authoritative name server (primary and secondary), verifies uniqueness, puts into database for a small fee, accredited by ICANN
 - registrar inserts two RRs into .com TLD server:
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
(networkutopia.com, dns2.networkutopia.com, NS)
(dns2.networkutopia.com, 212.212.212.2, A)
 - Create Type A record in your own authoritative server
(www.networkutopia.com, 212.212.212.3, A)
 - Create Type MX record in your own authoritative server
(mail.networkutopia.com, 212.212.212.4, MX)

Inserting records into DNS



TLD DNS server for .com

```
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
(networkutopia.com, dns2.networkutopia.com, NS)
(dns2.networkutopia.com, 212.212.212.2, A)
```



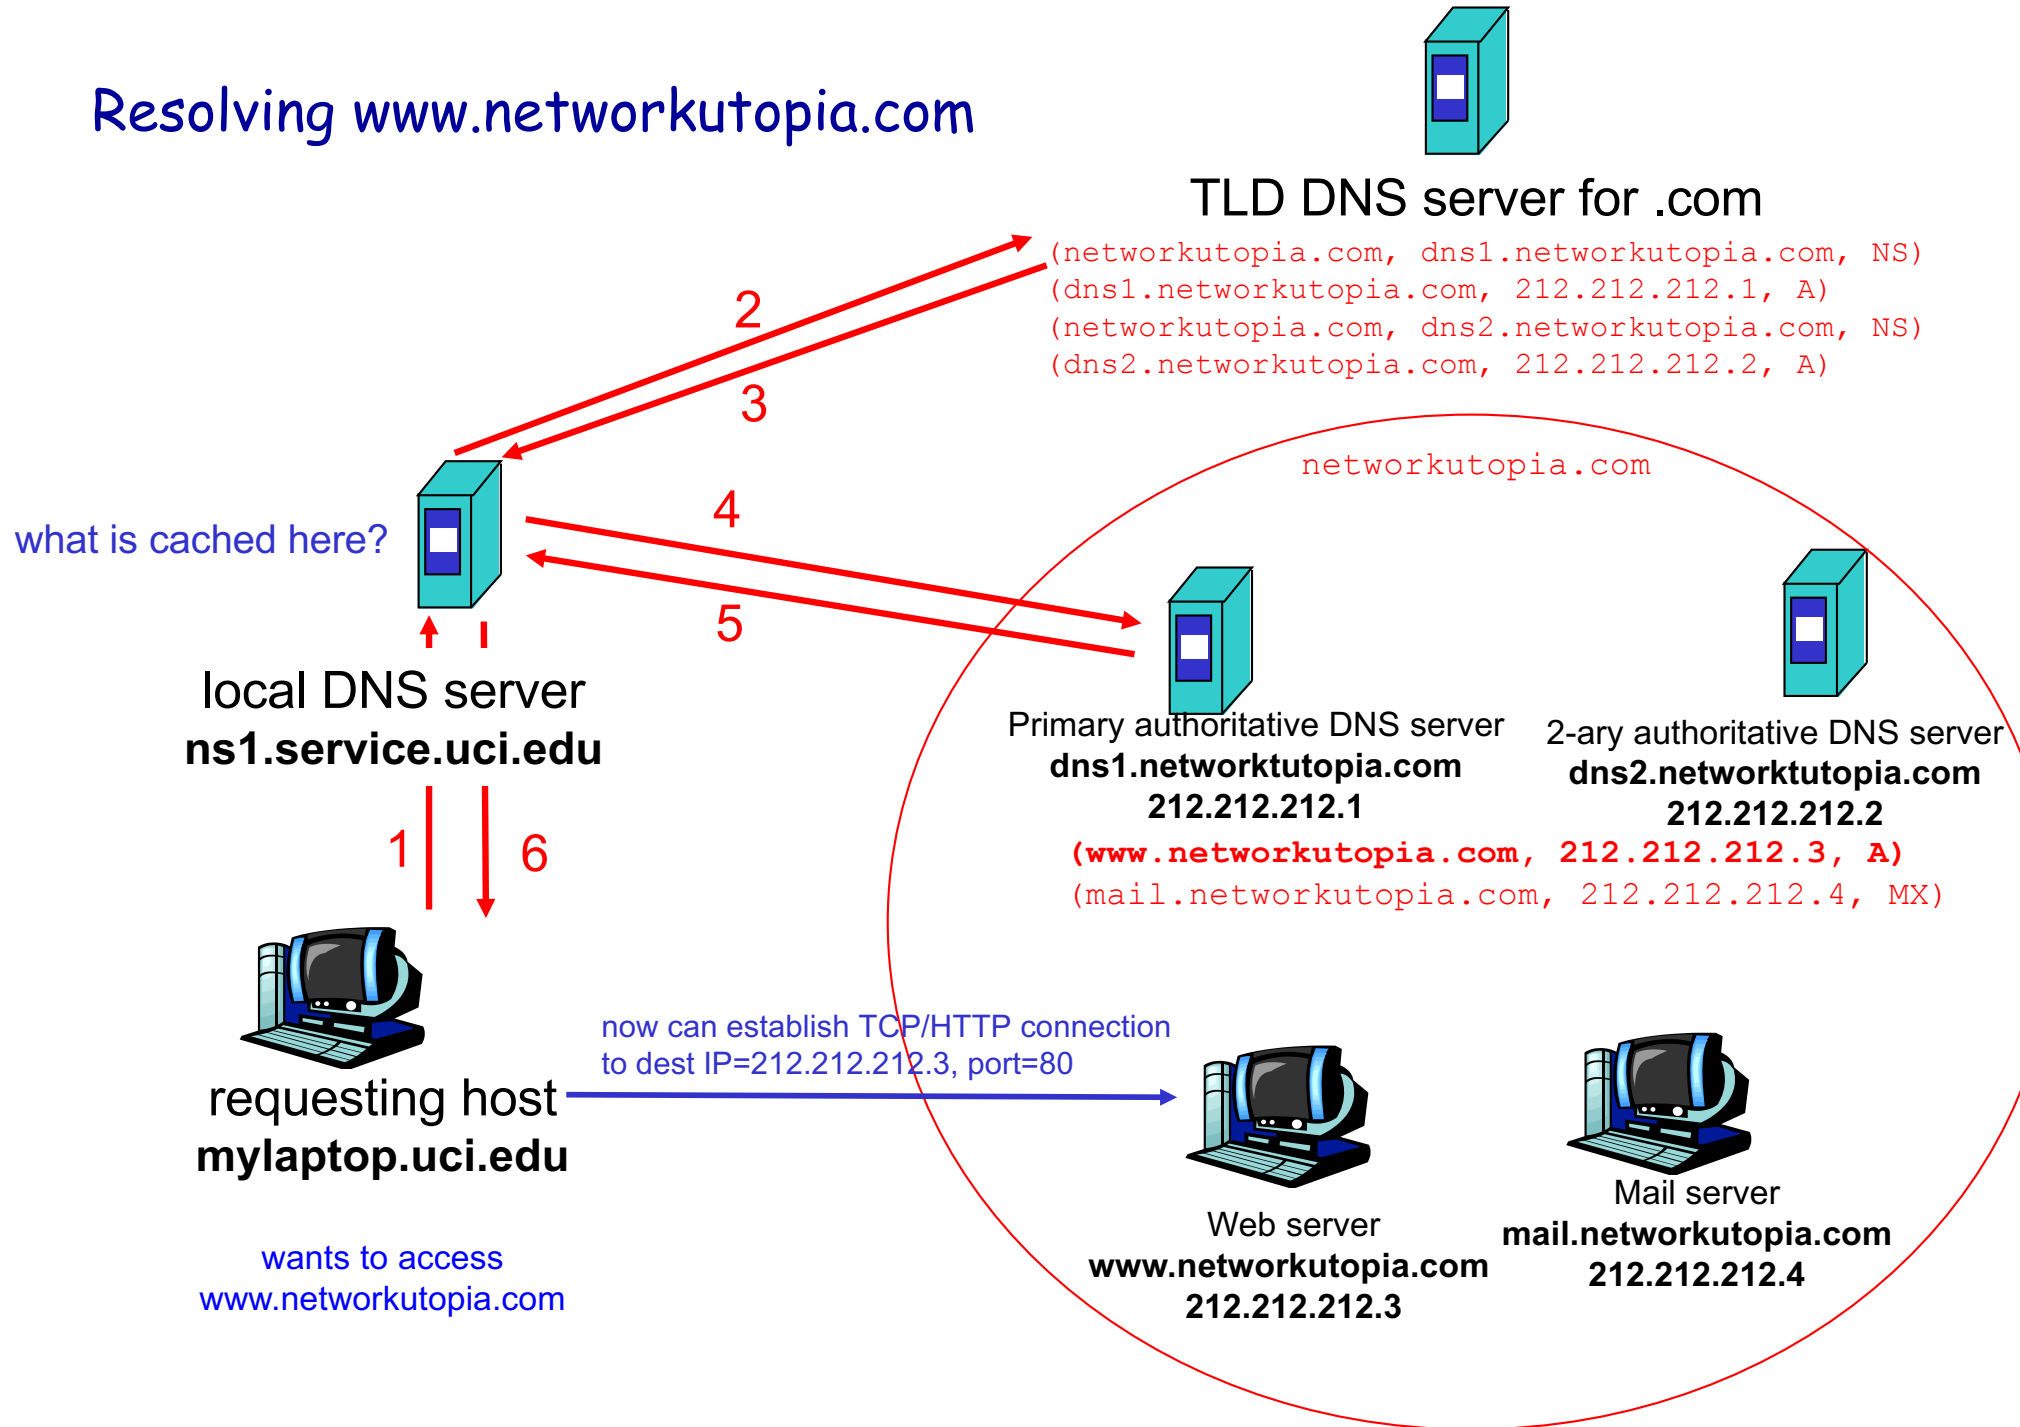
Example cont'd: querying DNS records

❖ Q: How do people visit the website www.networkutopia.com?

❖ A:

- **Host:** sends query to local DNS server
- **Local DNS server:** asks TLD server [or root , if TLD not in cache]
- **TLD server:** provides A and NS records for dns1.networkutopia.com
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
- **Local DNS server:** sends query to authoritative server (212.212.212.1)
- **Authoritative name server:** provides type A record
(www.networkutopia.com, 212.212.212.3, A)
- **Local DNS server:** returns this info to host (and caches RR for future use)
- **Host:** establishes TCP/HTTP connection to (IP: 212.212.212.3, port 80)

Resolving www.networkutopia.com



DNS Load balancing

- DNS may return many RRs in same response
 - E.g. try “dig www.amazon.com” multiple times
- Clients:
 - by default, they pick the first one
 - could also choose not to – this part is not standard
- Order of multiple records: **Unspecified**
 - most often: **Round Robin**
 - or static or preference to numerically “closer” networks
 - or taking into account load or RTT, or other metric computed by the client or by other (non DNS) servers
- Some references
 - http://en.wikipedia.org/wiki/Domain_Name_System
 - http://en.wikipedia.org/wiki/Round-robin_DNS
 - RFC1794: <http://tools.ietf.org/html/rfc1794>
 - [http://technet.microsoft.com/en-us/library/cc787484\(v=ws.10\).aspx](http://technet.microsoft.com/en-us/library/cc787484(v=ws.10).aspx)

Attacking DNS

- ICANN: <http://www.icann.org/>
- Attacks against root servers (2002, 2007):
 - DNS root servers proved robust (to pings or queries). Traffic filtering, caching, anycast load balancing: http://www.icann.org/en/announcements/factsheet-dns-attack-08mar07_v1.1.pdf
 - DDoS attacks to TLD more dangerous
 - 98% of TLD DNS queries are invalid: http://www.caida.org/publications/papers/2008/root_internet/root_internet.pdf
- Redirect/Man in the middle
 - Cache poisoning: send bogus replies to DNS servers that cache
 - Using DNS to redirect traffic
 - 2008: Kaminsky vulnerability: <http://www.unixwiz.net/techtips/iguide-kaminsky-dns-vuln.html>
 - 2009: Twitter Blackout: http://www.theregister.co.uk/2009/12/18/dns_twitter_hijack/
 - 2015: Tesla's DNS hacked (CPS threat): <http://www.tripwire.com/state-of-security/security-data-protection/teslas-dns-hacked-leads-website-and-twitter-hijack/>
 - Intercept queries: E.g. to block access to Facebook in China
- Using DNS to launch DDoS attacks
 - Send requests with spoofed source address (target) – responses flood target
 - Requires amplification

DNS Summary

- Core Internet functionality
- Implemented as a network application
 - On top of UDP (or TCP) port 53
 - Defined in RFCs: 1034, 1035 (1987)
 - <http://www.ietf.org/rfc/rfc1035.txt>
 - Proposed by Mockapertis (UCI PhD 1982)
 - http://en.wikipedia.org/wiki/Paul_Mockapetris
 - Many extensions, e.g. DNSSEC

Practice

- Interactive Exercise:
 - Delay when browsing: DNS+TCP+HTTP:
 - https://gaia.cs.umass.edu/kurose_ross/interactive/DNS_HTTP_delay.php
- Interactive Animation
 - Recursive vs iterative DNS Queries
 - https://media.pearsoncmg.com/aw/ecs_kurose_compnetwork_7/cw/content/interactiveanimations/recursive-iterative-queries-in-dns/index.html
- Problems from the book