


Chapter 3

Transport Layer



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross
Pearson/Addison Wesley

© Based on materials developed by J.F.Kurose and K.W.Ross.
All right reserved.

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP


- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
 - send side: breaks app messages into *segments*, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - Internet: TCP and UDP



Transport vs. network layer

- *application layer*
- *transport layer*: logical communication between processes
 - relies on + enhances, network layer services
- *network layer*: logical communication between hosts

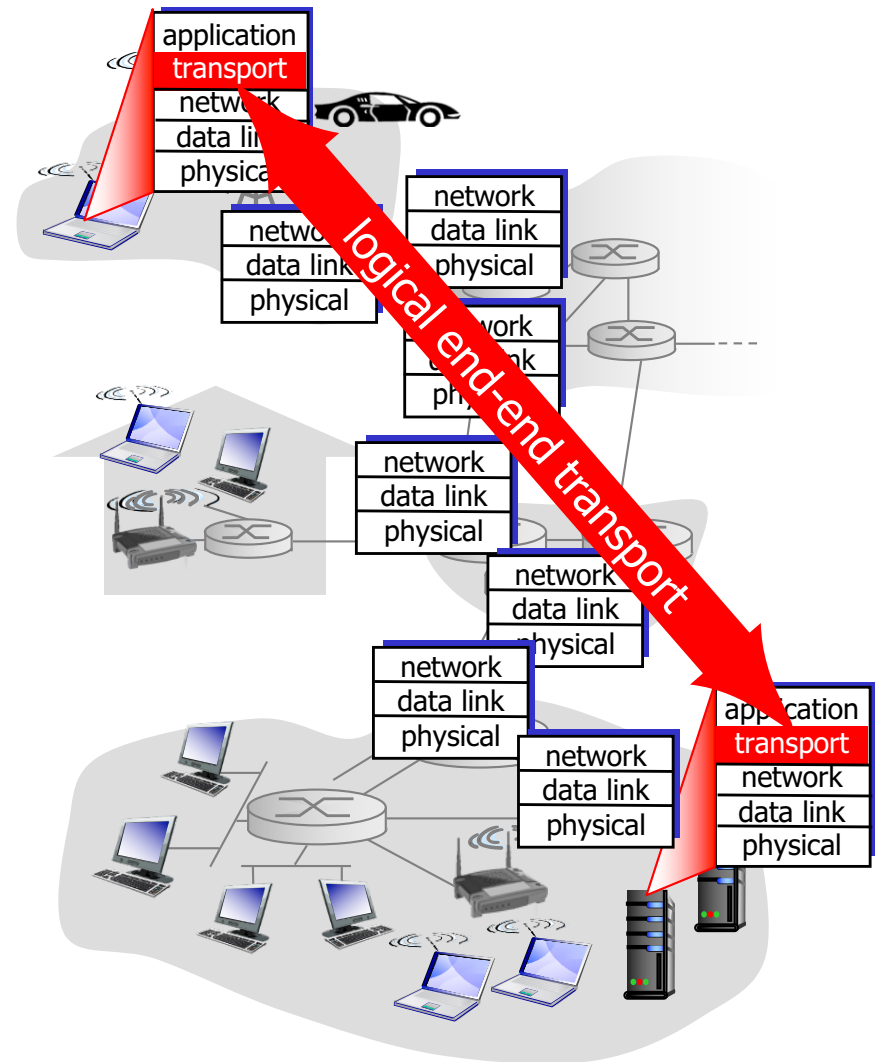
household analogy:

12 kids in Ann's house sending letters to 12 kids in Bill's house:

- hosts = houses
- processes = kids
- Applications: write a book, exchange letters, photos
- app messages = letters in envelopes
- transport protocol = Ann and Bill who demux to in-house siblings
- network-layer protocol = postal service
- additional transport services = reliability, encryption
- another transport protocol = younger kids demux to siblings

Services by Internet Transport?

- reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- unreliable, unordered delivery: UDP
 - no-frills extension of “best-effort” IP
- services not available:
 - delay guarantees
 - bandwidth guarantees



Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

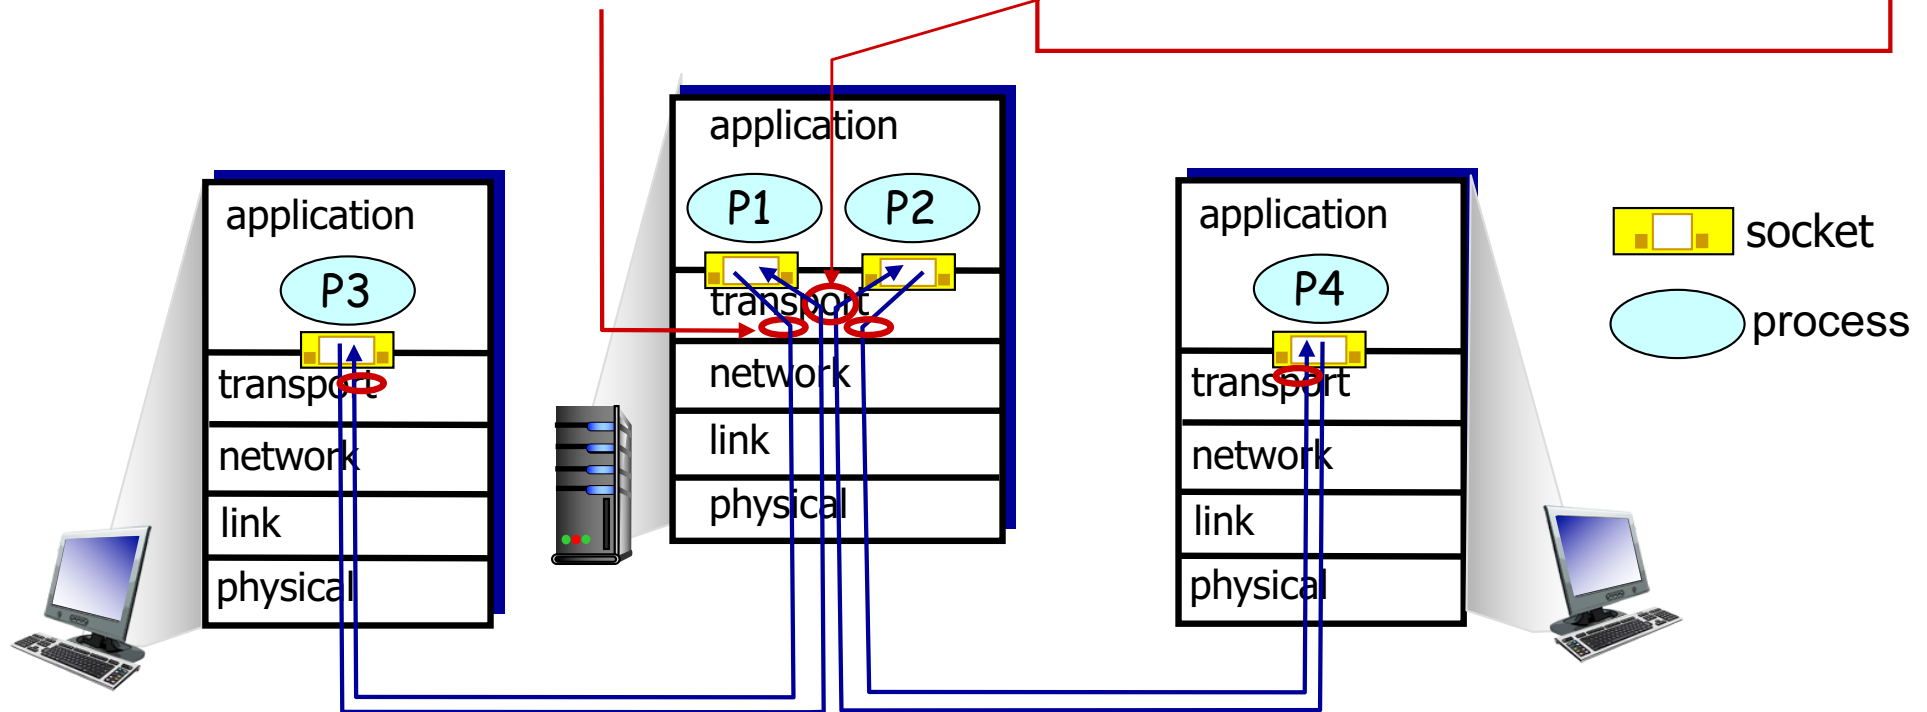
3.6 principles of congestion control

3.7 TCP congestion control

Multiplexing/demultiplexing

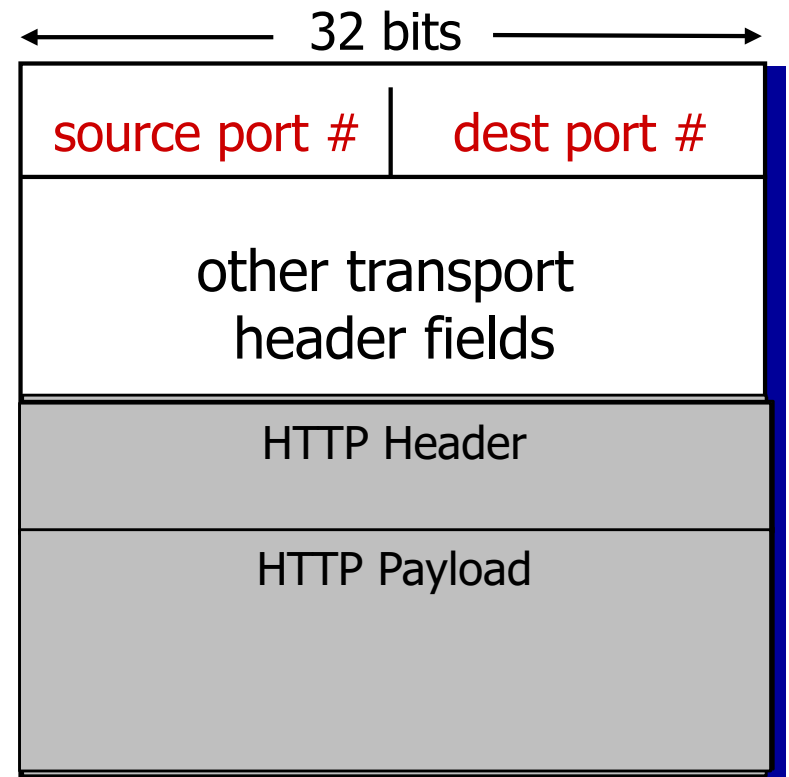
multiplexing at sender:
handle data from multiple sockets, **add transport header** (later used for demultiplexing)

demultiplexing at receiver:
use header info to deliver received segments to correct socket



How demultiplexing works

- host receives IP datagrams
 - each datagram has source IP address, destination IP address
 - each datagram carries one transport-layer segment
 - each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket
- recall: **encapsulation**



TCP/UDP segment format

Connectionless demultiplexing

- *recall*: created socket has host-local port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534) ;
```

- *recall*: when creating datagram to send into UDP socket, must specify

- destination IP address
- destination port #

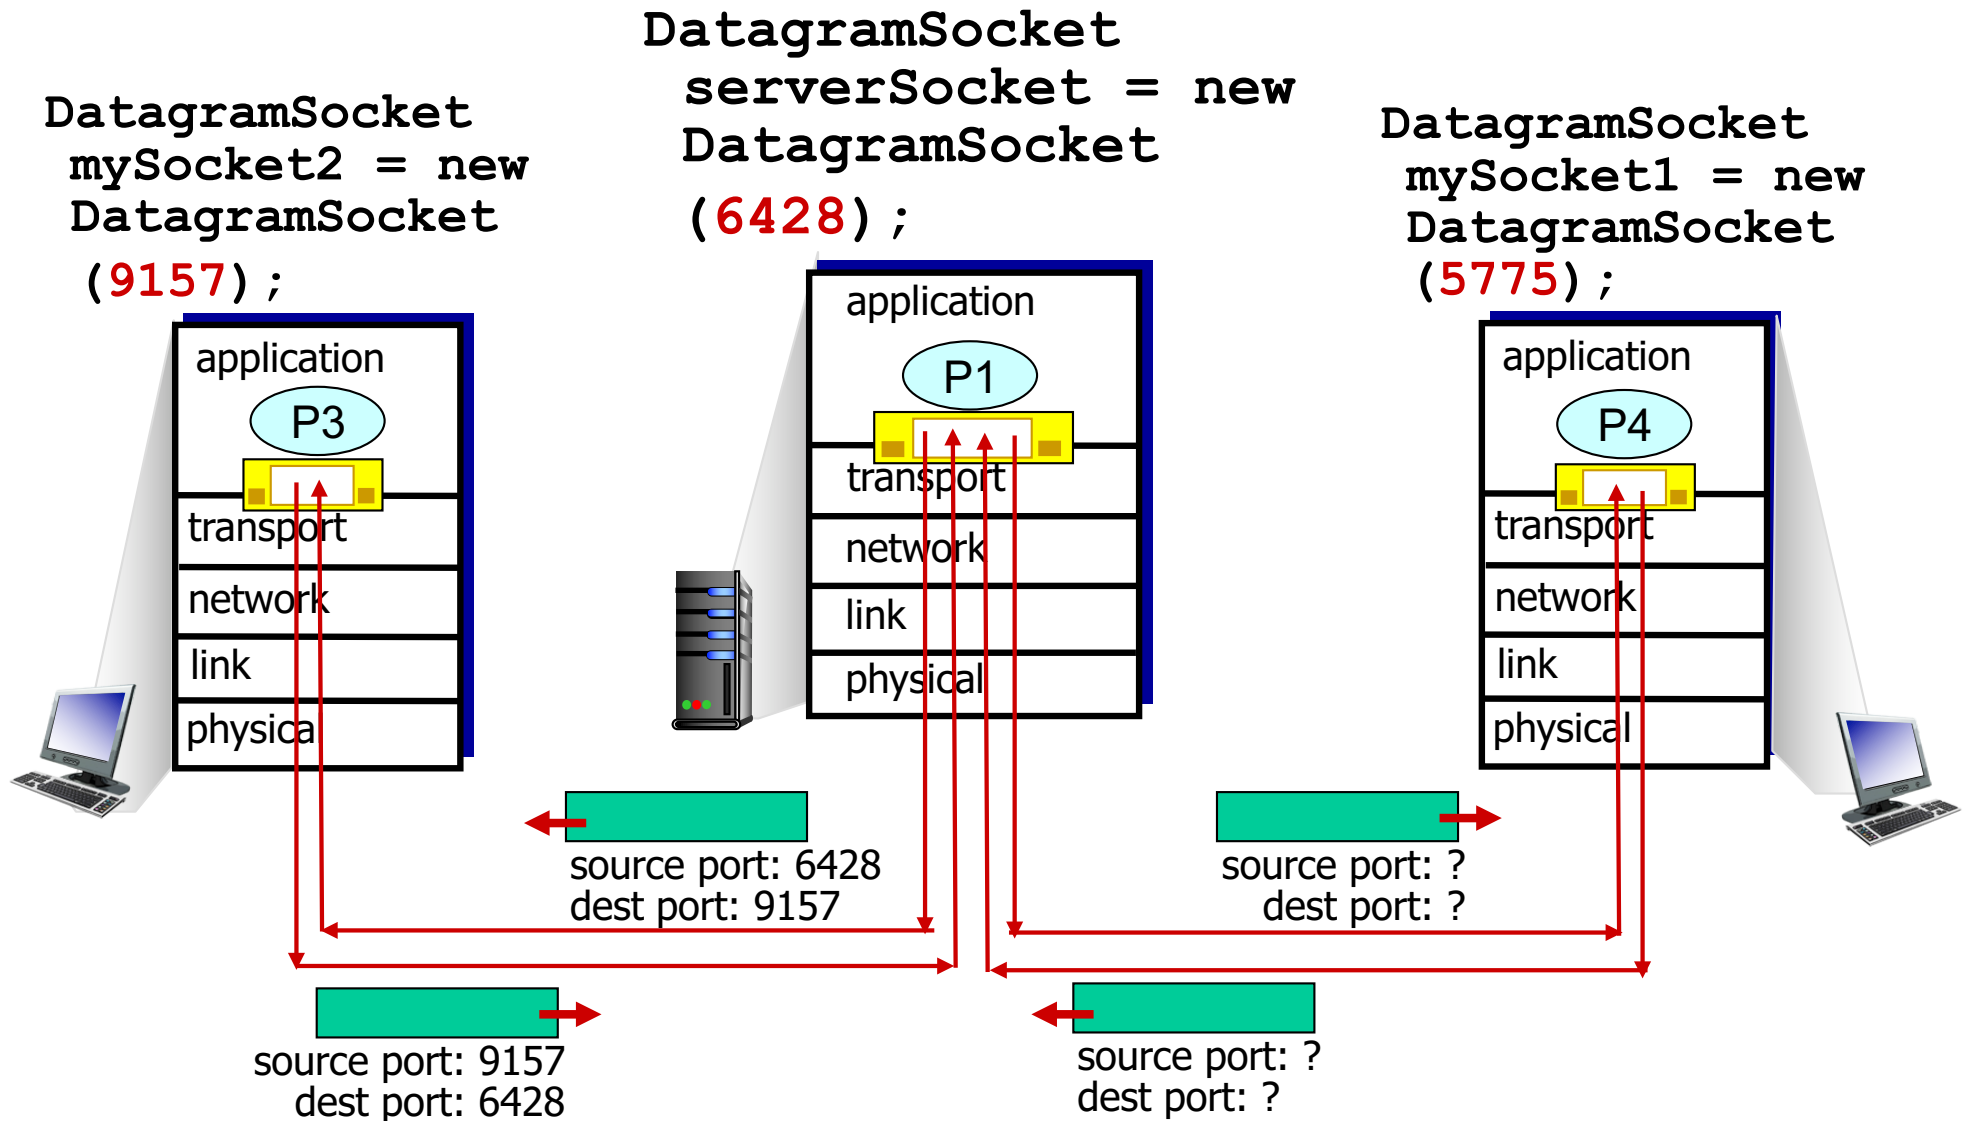
-
- when host receives UDP segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

Connectionless demux: example



Recall: UDP client

Python UDPClient

```
from socket import *
serverName = 'hostname'
serverPort = 6428
clientSocket = socket(AF_INET,
                      SOCK_DGRAM)
message = raw_input('Input lowercase sentence:')
clientSocket.sendto(message.encode(),
                    (serverName, serverPort))
modifiedMessage, serverAddress =
    clientSocket.recvfrom(2048)
print modifiedMessage.decode()
clientSocket.close()
```

create UDP socket
(know own IP and local
port=9157, assigned
automatically)

dest IP= DNS lookup of
server name
dest port=serverPort =6428

Demux received UDP
segment based on dest
(here 9157) port on UDP
header

Recall: UDP server

Python UDPServer

```
from socket import *
```

```
serverPort = 9157
```

```
serverSocket = socket(AF_INET, SOCK_DGRAM)
```

```
serverSocket.bind(('', serverPort))
```

```
print ("The server is ready to receive")
```

```
while True:
```

```
    message, clientAddress = serverSocket.recvfrom(2048)
```

```
    modifiedMessage = message.decode().upper()
```

```
    serverSocket.sendto(modifiedMessage.encode(),  
                        clientAddress)
```

create UDP socket →

bind socket to local port
number 9157 →

loop forever →

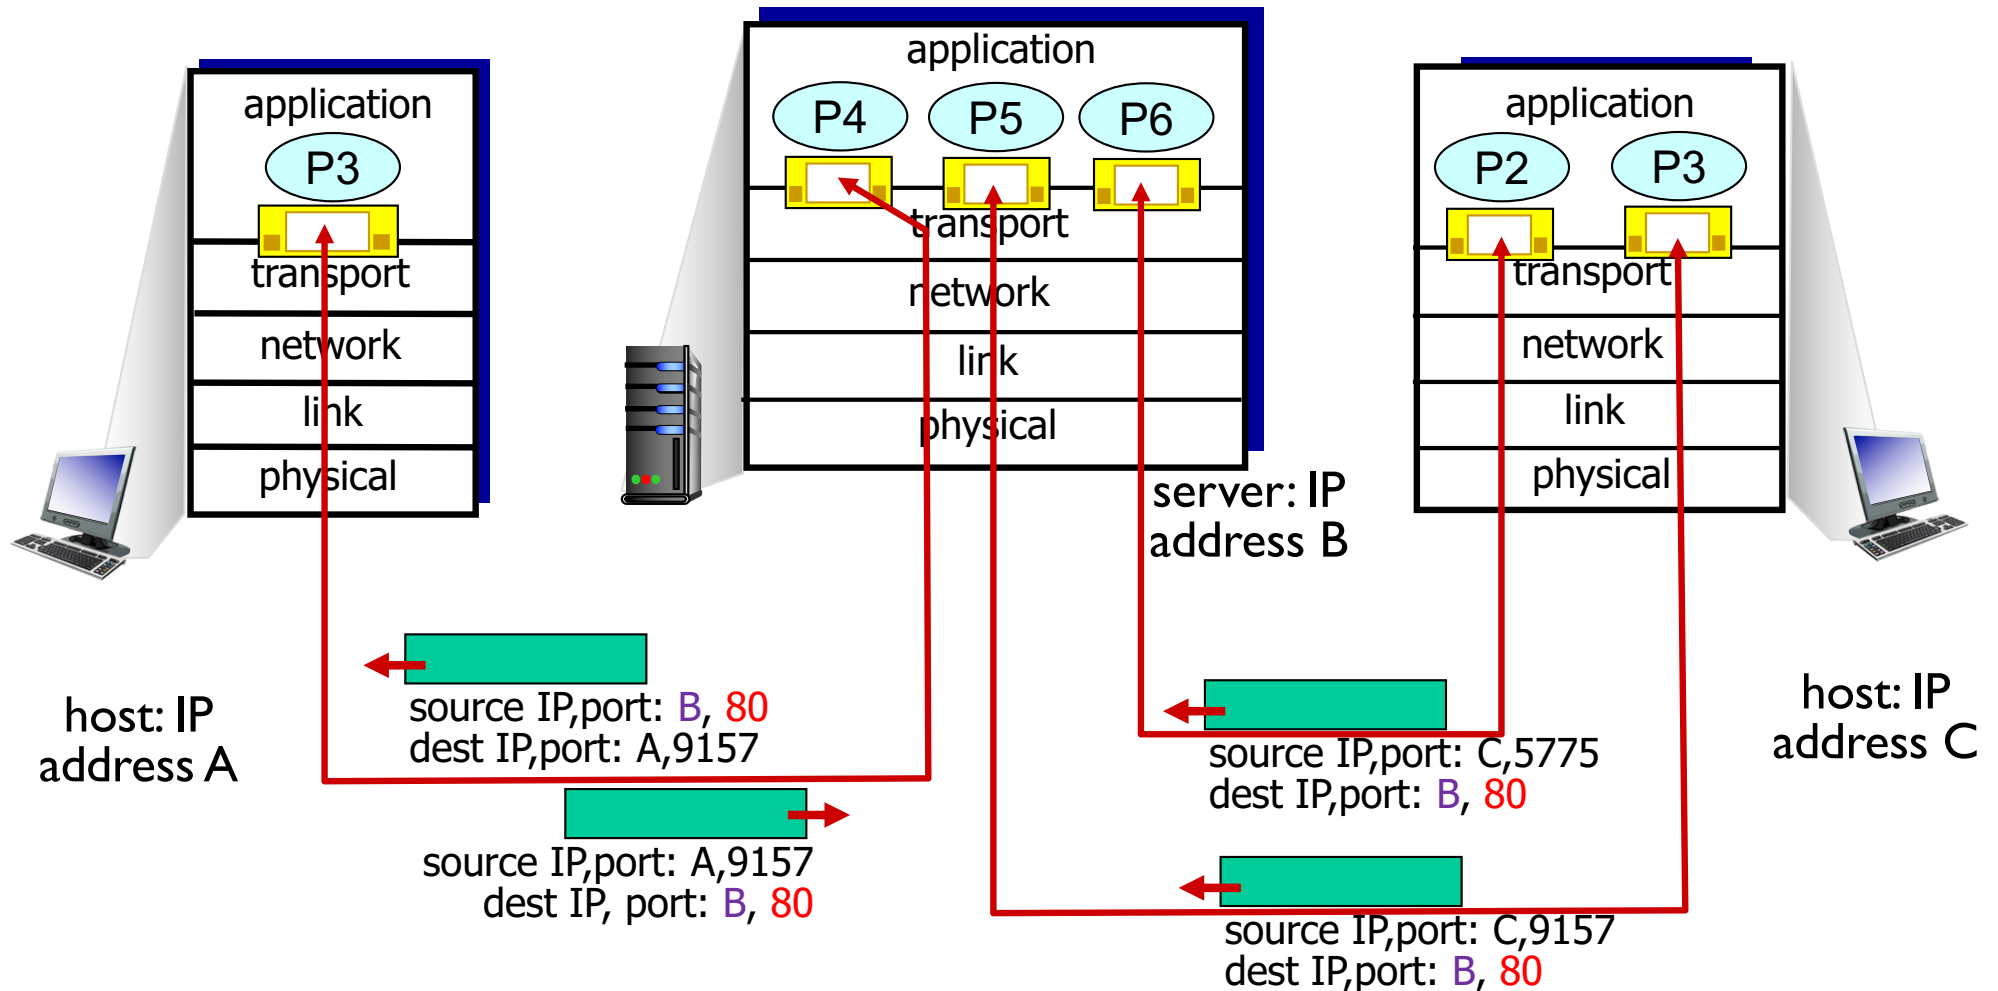
Read from UDP socket into
message, getting client's
address (client IP and port) →

send upper case string
back to this client →

Connection-oriented demux

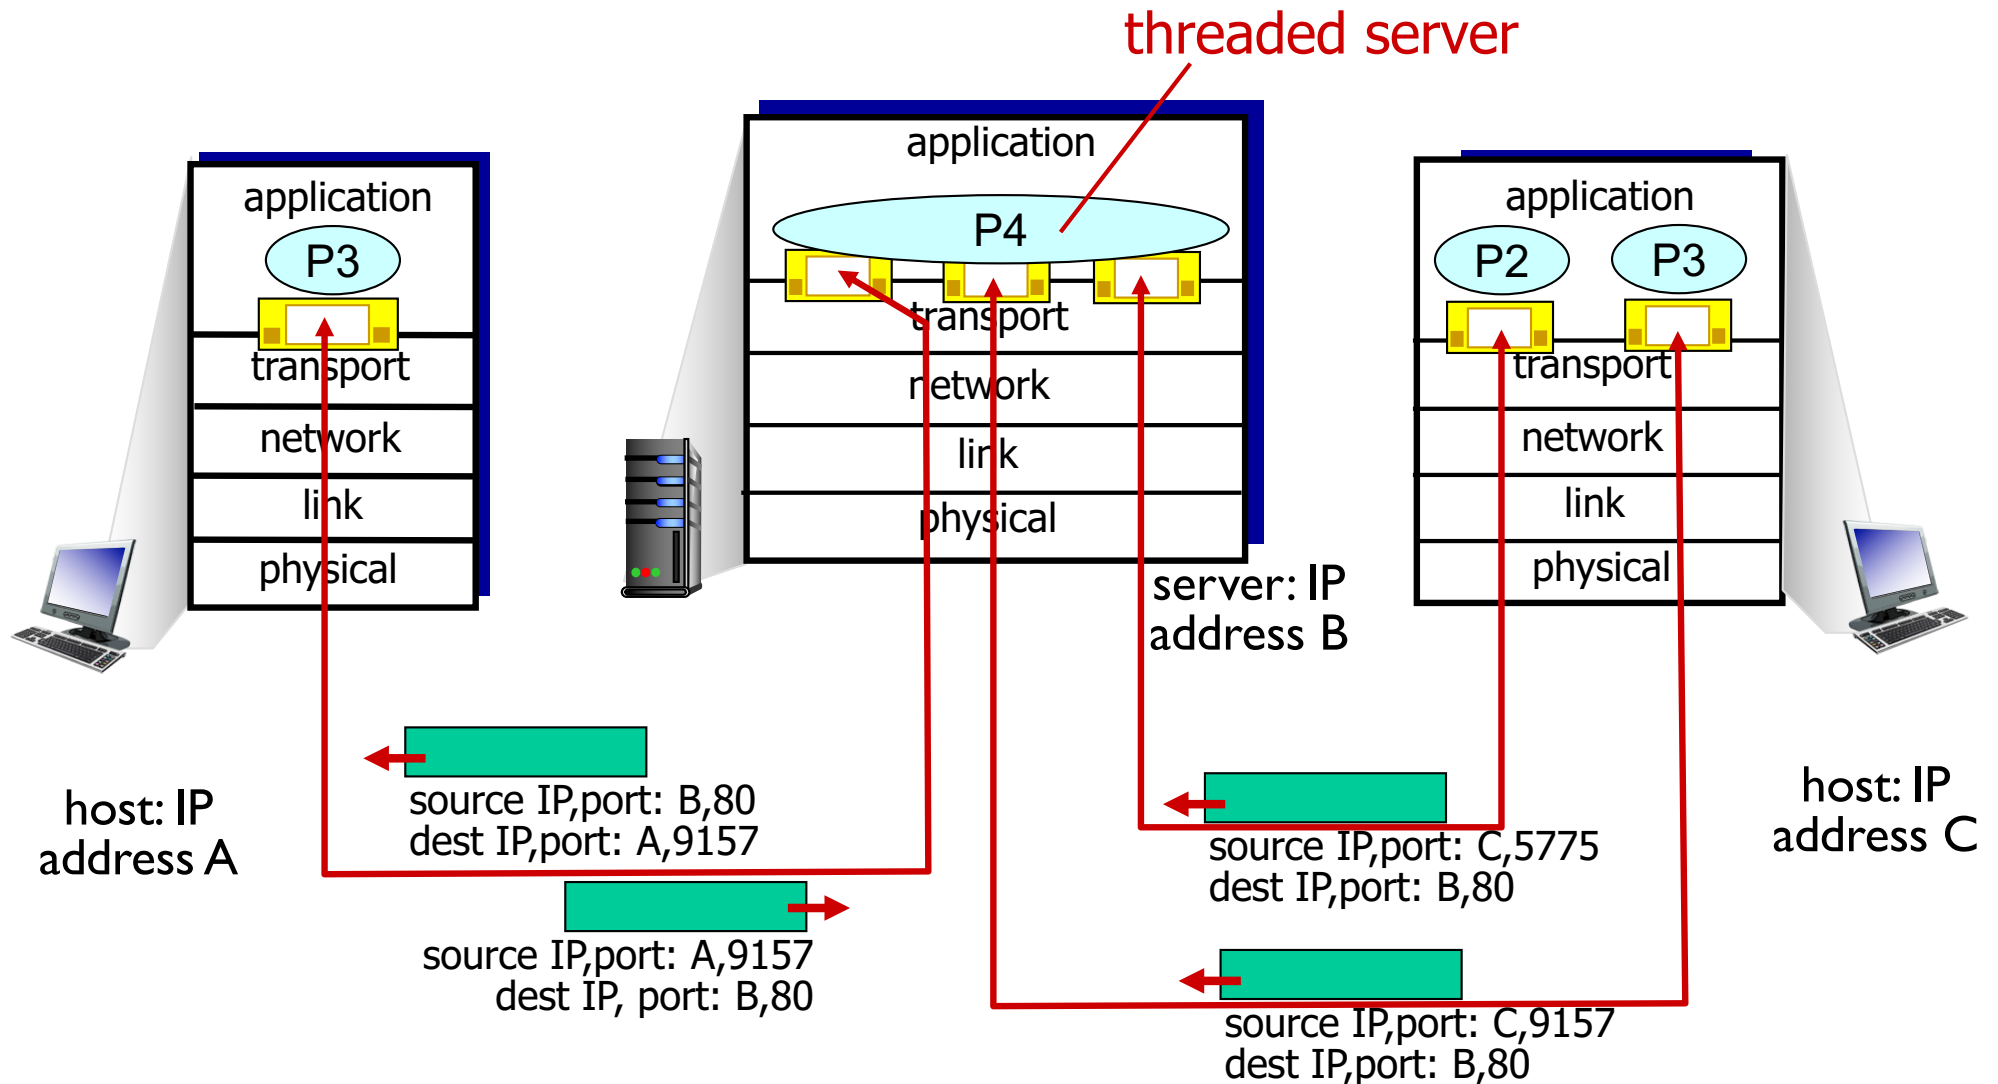
- TCP socket identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- demux: receiver uses all four values to direct segment to appropriate socket
- server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
 - non-persistent HTTP will have different socket for each request

Connection-oriented demux: example



three segments, all destined to IP address: B,
dest port: 80 are demultiplexed to *different* sockets

Connection-oriented demux: example



Example app: TCP client

Python TCPClient A

From now on TCP header
will know to include
(src IP=**A**, src port=**9157**,
dest IP=**B**, dest port=**80**)
in each packet

```
from socket import *
serverName = 'servername'
serverPort = 80
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort))
sentence = raw_input('Input lowercase sentence:')
clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```


Example app: TCP server

Python TCPServer B

```
from socket import *
serverPort = 80
serverSocket = socket(AF_INET, SOCK_STREAM)
serverSocket.bind(('', serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
while True:
```

Maps new socket to
(src IP=A, source port=9157,
dest IP=B, dest Port=80)

```
connectionSocket, addr = serverSocket.accept()
```

Reads these 4 fields from
TCP packet header and
delivers to right socket

```
sentence = connectionSocket.recv(1024).decode()
capitalizedSentence = sentence.upper()
connectionSocket.send(capitalizedSentence.
                        encode())
```

```
connectionSocket.close()
```

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

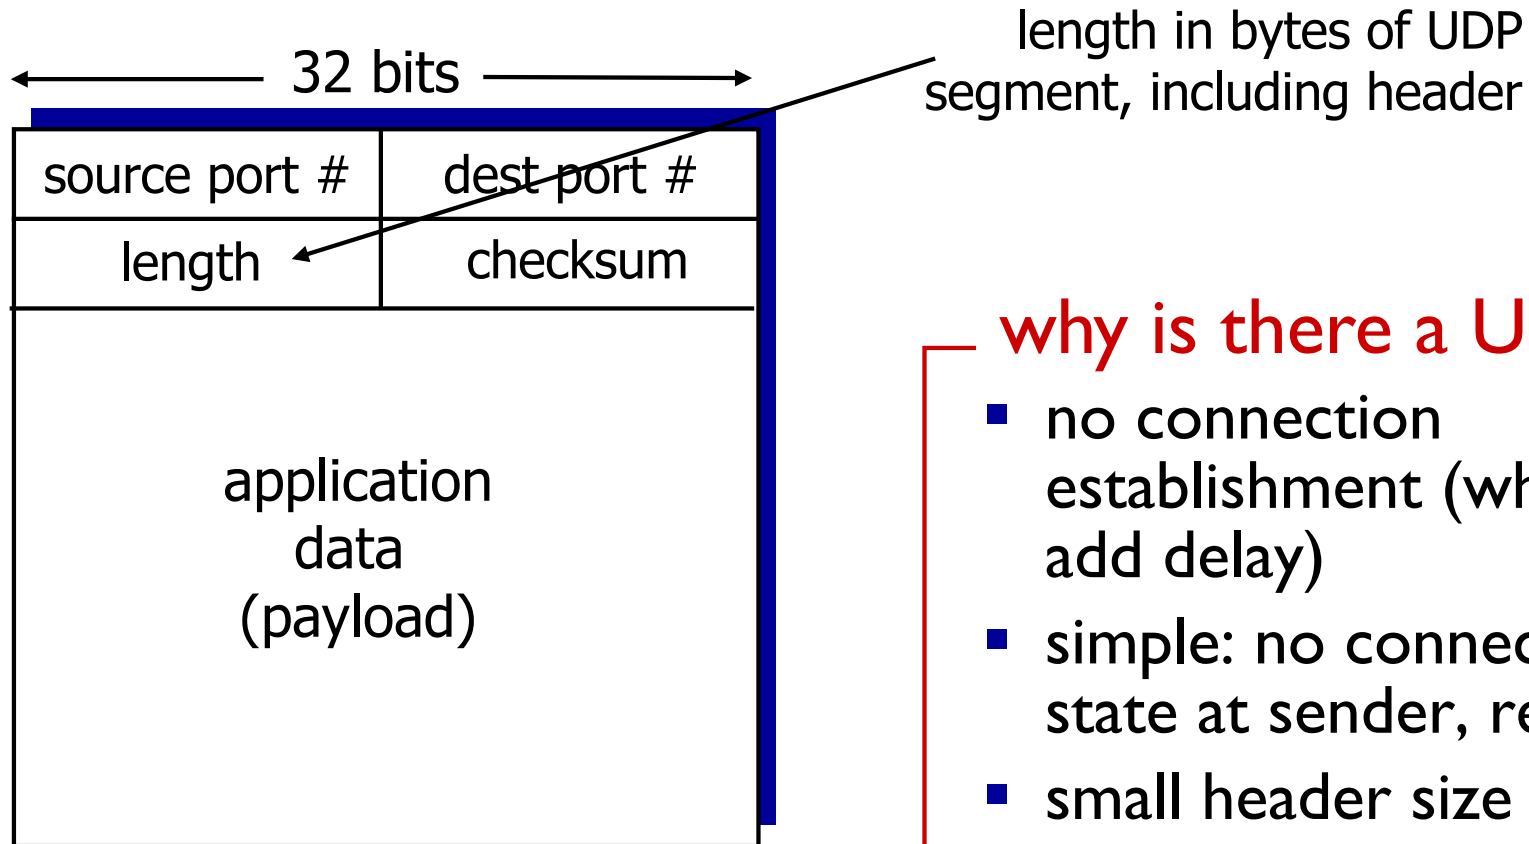
3.6 principles of congestion control

3.7 TCP congestion control

UDP: User Datagram Protocol [RFC 768]

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service from app’s point of view. UDP segments may be:
 - lost
 - delivered out-of-order
- *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others
- UDP use:
 - streaming media apps (loss tolerant, rate sensitive)
 - DNS
 - SNMP
- reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!

UDP: segment header



UDP segment format

why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size : 8B
- no congestion control: UDP can blast away as fast as desired

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

sender:

- treat segment contents, including header fields, as sequence of 16-bit integers
 - a number X
- checksum: addition (one's complement sum) of segment contents
 - Compute $f(x)$
- sender puts checksum value into UDP checksum field
 - Transmit $X, f(X)$

receiver:

- Receive segment X'
- compute checksum of received segment $f(X')$
- check if computed checksum equals checksum field value:
 - Is $f(X)=f(X')??$
 - YES - no error detected, i.e., $X=X$
 - *But maybe errors nonetheless...*
 - NO - error detected, i.e., $X' \neq X$

Internet checksum: example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result