


# rdt3.0: stop-and-wait operation



$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

# rdt3.0: stop-and-wait performance

- rdt3.0 is correct, but performance stinks
- e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microseconds}$$

- $U_{\text{sender}}$  utilization (or “throughput” in your HW1): fraction of time sender busy sending


$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- if RTT=30 msec, 1KB pkt every 30 msec: 33kB/sec throughput over 1 Gbps link
  - Example of a (bad) network protocol limits use of physical resources (fast link)!

# Pipelined protocols

**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver
- [note: different than pipelining in HWI]




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- two generic forms of pipelined protocols:
  - *go-Back-N*
  - *selective repeat*

# Pipelining: increased utilization



**Q:** What is the best choice of the “window” of pipelined messages?

**A:**  $RTT/L/R+1$  keeps the channel always busy

# Pipelined protocols: overview


## Go-back-N:

- sender can have up to N unacked packets in pipeline (**window**)
- receiver only sends *cumulative ack*
  - doesn't ack packet if there is a gap
- sender has **timer for oldest unacked packet**
  - when timer expires, retransmit *all* unacked packets

## Selective Repeat:

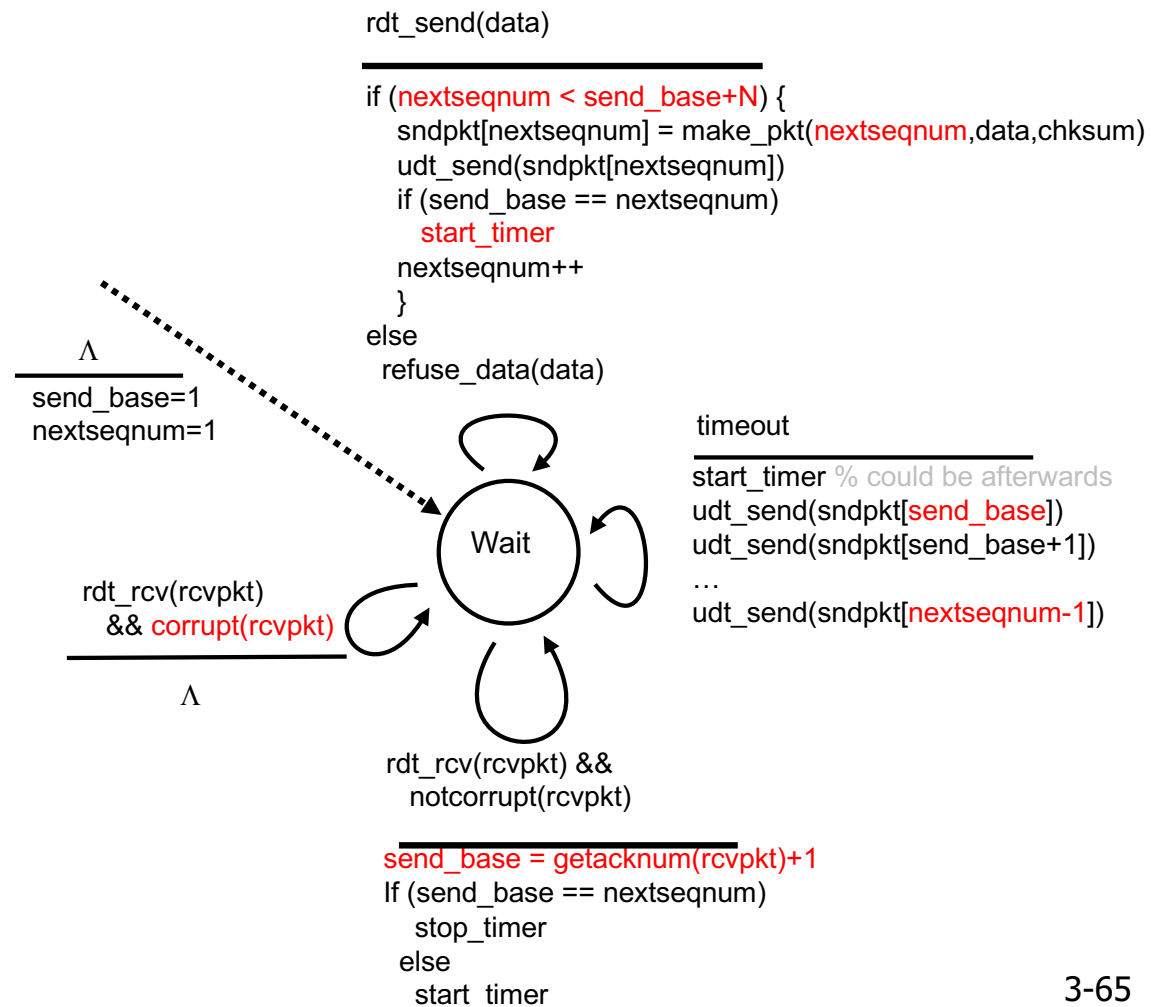
- sender can have up to N unack'ed packets in pipeline (**window**)
- rcvr sends *individual ack* for each packet
- sender maintains **timer for each unacked packet**
  - when timer expires, retransmit only that unacked packet

# GBN Sender



- k-bit sequence # in pkt header
- sliding window of up to  $N$ , consecutive unack'd pkts allowed

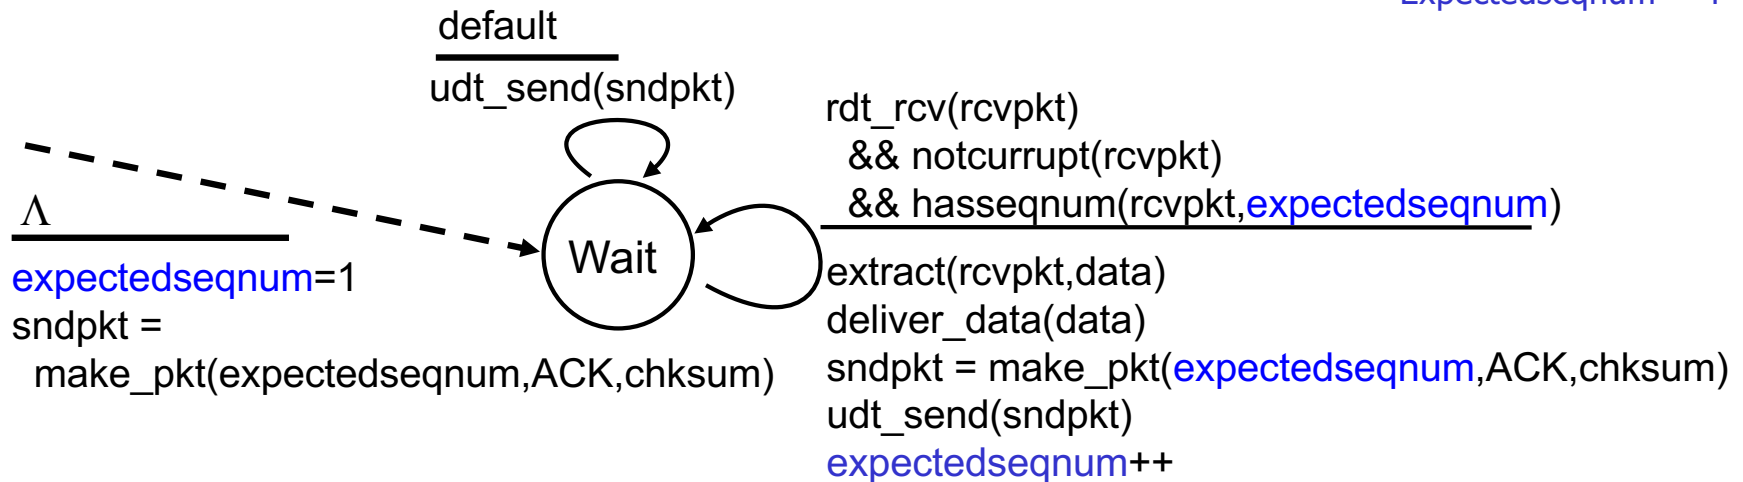
- ❖ Send\_base:
  - ❖ oldest in-flight packet ( $n$ )
  - ❖ nextseqnum: next to send
- ❖ ACK( $m$ ): ACKs all pkts up to, including seq #  $m$  - “cumulative ACK”
  - may receive duplicate ACKs (see receiver)
- ❖ timeout( $n$ ): keep timer for oldest unack-ed packet ( $\text{send\_base}$ ); retransmit all sent but unacked pkts in window – i.e., all yellow packets



# GBN Receiver


0 1 2 3 4 5 6 7 8

Expectedseqnum = 4



- Simple state: need only remember **expectedseqnum**
  - This is the next seq (in order) we expect to receive
  - Everything before, is already received
- ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #
  - may generate duplicate ACKs
- If it receives out-of-order pkt:
  - Discard, *no receiver buffering!*
  - re-ACK pkt with highest in-order seq #

# GBN in action



sender window (N=4)

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

sender

send pkt0  
 send pkt1  
 send pkt2  
 send pkt3  
 (wait)

receiver

receive pkt0, send ack0  
 receive pkt1, send ack1

rcv ack0, send pkt4  
 rcv ack1, send pkt5

receive pkt3, **discard**,  
 (re)send ack1

receive pkt4, **discard**,  
 (re)send ack1

receive pkt5, **discard**,  
 (re)send ack1

ignore duplicate ACK



**pkt 2 timeout**

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

send pkt2  
 send pkt3  
 send pkt4  
 send pkt5


rcv pkt2, deliver, **send ack2**  
 rcv pkt3, deliver, send ack3  
 rcv pkt4, deliver, send ack4  
 rcv pkt5, deliver, send ack5

0 1 2 3 4 5 6 7 8

rcv ack2



# GBN in action



sender window (N=4)

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

sender

send pkt0  
 send pkt1  
 send pkt2  
 send pkt3  
 (wait)

receiver

receive pkt0, send ack0  
 receive pkt1, send ack1

rcv ack0, send pkt4  
 rcv ack1, send pkt5

receive pkt3, **discard**,  
 (re)send ack1

receive pkt4, **discard**,  
 (re)send ack1

receive pkt5, **discard**,  
 (re)send ack1

ignore duplicate ACK



**pkt 2 timeout**

0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8  
 0 1 2 3 4 5 6 7 8

send pkt2  
 send pkt3  
 send pkt4  
 send pkt5

rcv pkt2, deliver, **send ack2**  
 rcv pkt3, deliver, send ack3  
 rcv pkt4, deliver, send ack4  
 rcv pkt5, deliver, send ack5

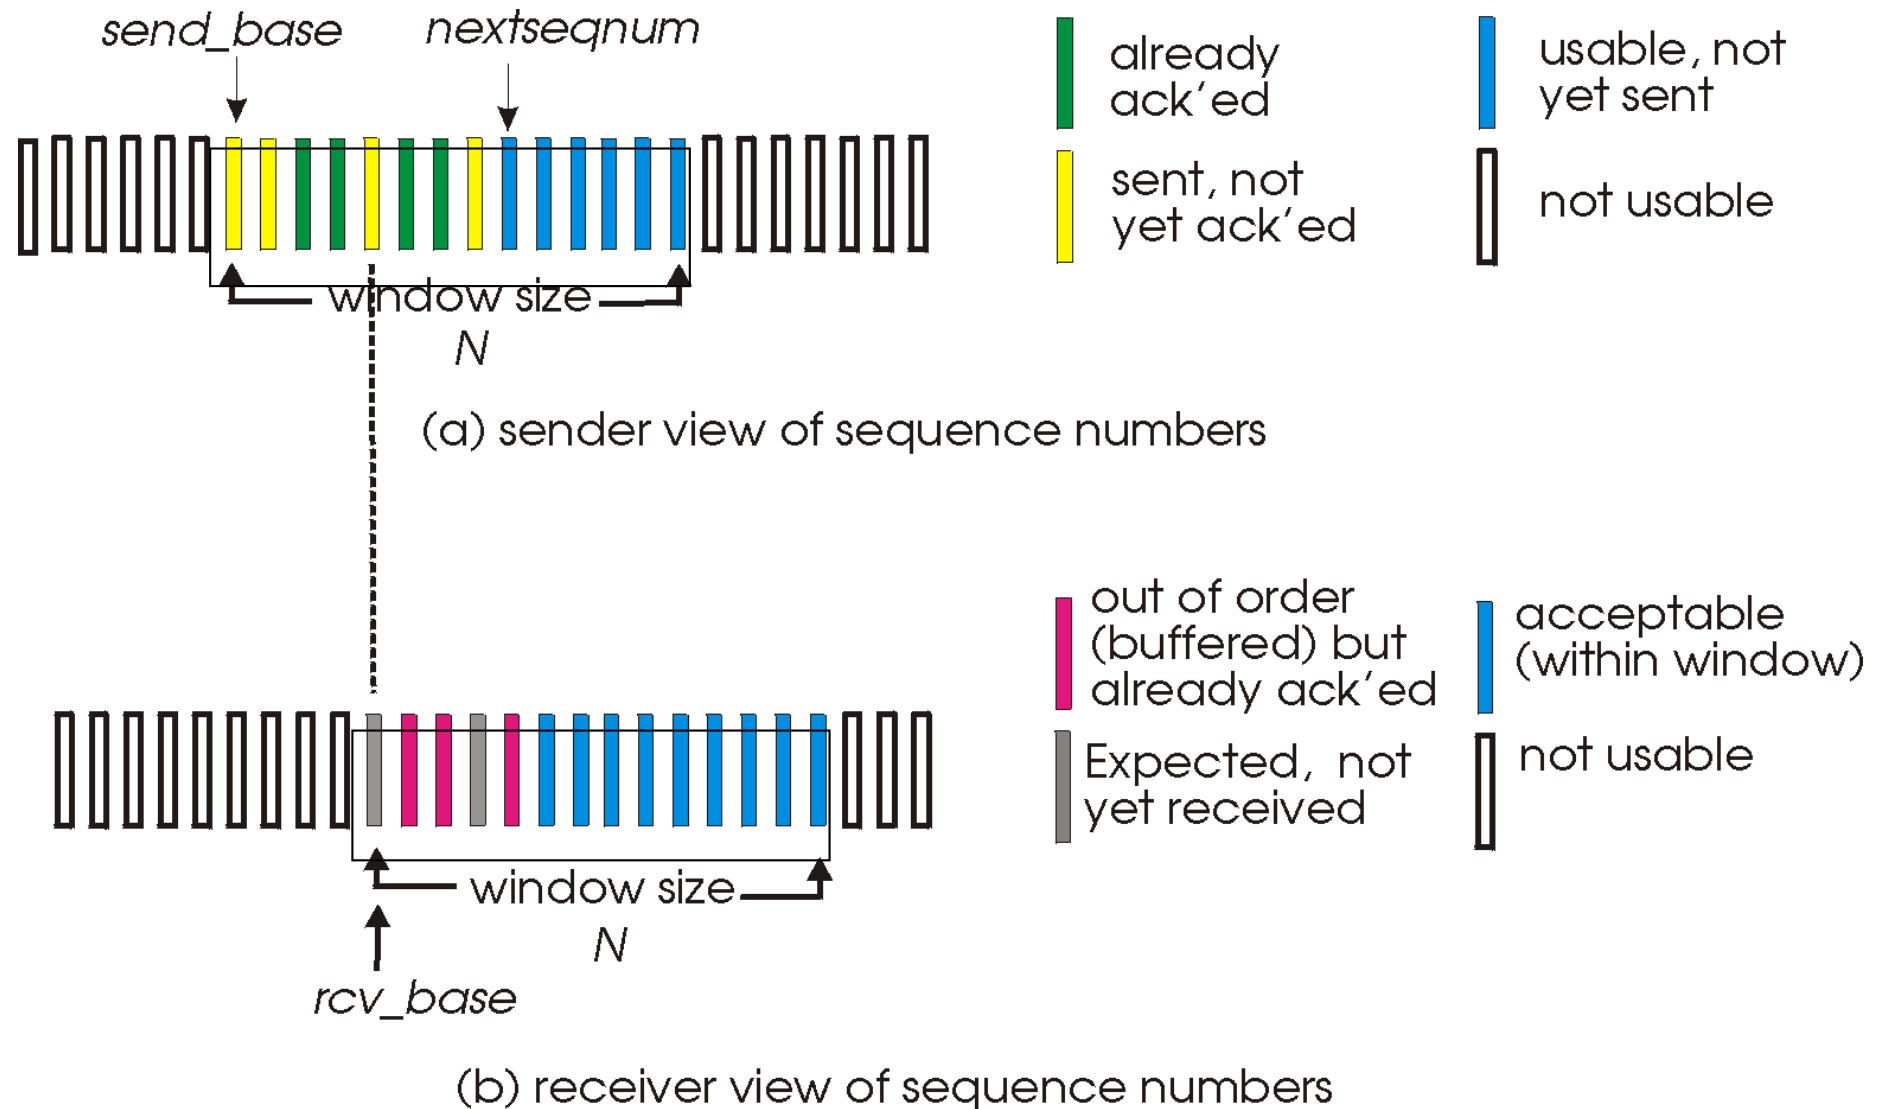
0 1 2 3 4 5 6 7 8

rcv ack2

# Selective Repeat (SR)

- Receiver *individually acknowledges all* correctly received pkts
  - whether they are received in-order or out-of-order
  - buffers pkts, as needed, for eventual in-order delivery to upper layer
- Sender *only retransmits un-ACKed pkts*
  - sender maintains timer for each unACKed pkt
- Both *sender and receiver* maintain *windows*
  - Sender window
    - $N$  consecutive seq #'s
    - limits seq #s of sent, unACKed pkts
  - Receiver window
    - $N$  consecutive seq #'s
    - Limits seq#s of buffered out-of-order packets

# Selective Repeat: sender, receiver windows



# Selective Repeat (FSM not shown)

## sender

### data from above:

- if next available seq # in window, send pkt

### timeout(n):

- resend pkt n, restart timer

### ACK(n) in [sendbase, sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

## receiver

### pkt n in [rcvbase, rcvbase+N-1] % new

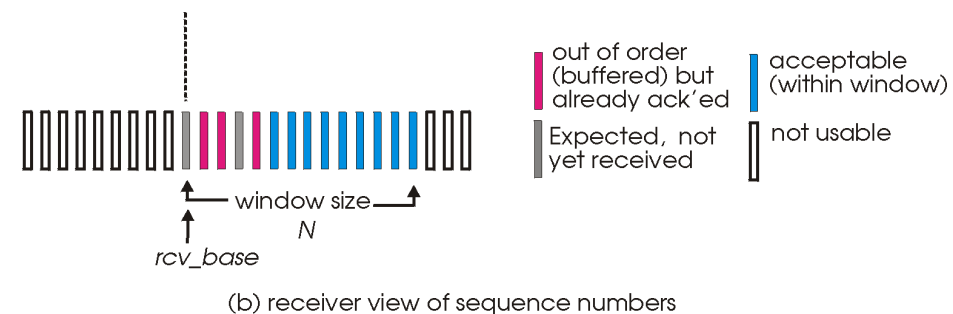

- ❖ send ACK(n)
- ❖ out-of-order: buffer
- ❖ in-order: deliver (also deliver all buffered, in-order pkts), advance window to next not-yet-received pkt

### pkt n in [rcvbase-N, rcvbase-1] % duplicate

- ❖ ACK(n) [Why not ignore?]

### otherwise:

- ❖ ignore



# Selective repeat in action

sender window (N=4)

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8 -

sender

send pkt0

send pkt1

send pkt2

send pkt3

(wait)

rcv ack0, send pkt4

rcv ack1, send pkt5

record ack3 arrived



*pkt 2 timeout*

send pkt2

record ack4 arrived

record ack5 arrived

*Q1: what happens when ack2 arrives?*

receiver

0 1 2 3 4 5 6 7 8

receive pkt0, send ack0

0 1 2 3 4 5 6 7 8

receive pkt1, send ack1

0 1 2 3 4 5 6 7 8

receive pkt3, buffer, send ack3

0 1 2 3 4 5 6 7 8

receive pkt4, buffer, send ack4

0 1 2 3 4 5 6 7 8

receive pkt5, buffer, send ack5

0 1 2 3 4 5 6 7 8

rcv pkt2; send ack2

deliver pkt2, pkt3, pkt4, pkt5;

0 1 2 3 4 5 6 7 8

## Lack of sync + finite seq#: ambiguity

### Example:

- seq #'s: 0, 1, 2, 3
- window size=3


**Problem:** receiver sees no difference in 2 scenarios!

- ❖ new data arrive in (a)
- ❖ duplicate data accepted as new in (b)


**Q1:** would the problem occur if seq# up to 5?

**Q2:** what relationship between seq # size (SN) and window size (N) to avoid problem?

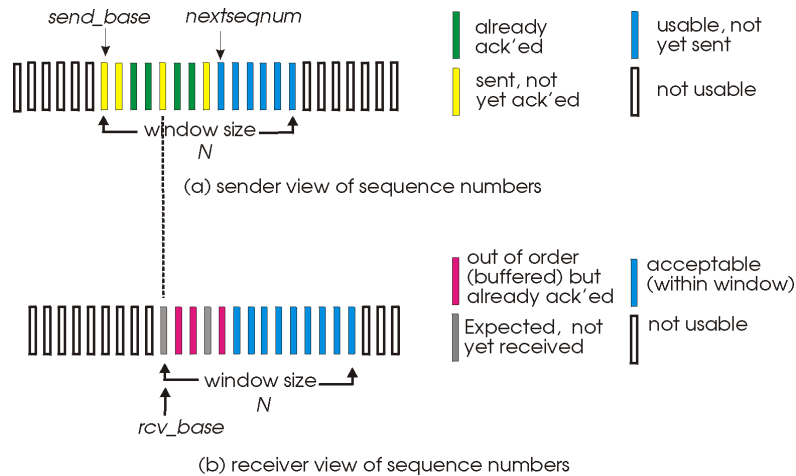
**A2:**  $SN \geq 2N$  because sender and receiver window must have overlap of at least 1



*receiver can't see sender side.  
receiver behavior identical in both cases!  
something's (very) wrong!*




# SR: sender, receiver windows & Seq No



N= **Sender** window  
= **Receiver** window

←-----> SN: sequence numbers for packets and acks



**Sender** & **receiver** windows overlap. At most consecutive.  
 $SN \geq 2N \rightarrow$  seqnums for packets and acks are unique



**Sender** & **receiver** windows in perfect sync



Impossible



Impossible



Special Case: Stop and Wait: Sender  $W=2$ , Receiver  $W=1$

# Go-back-N vs SR: Mechanisms

## Go-back-N:

- Sender can have up to N unack'ed packets in pipeline
  - sliding window
- Rcvr sends *cumulative* ack for last in-order packet
  - maintains *expectedseqnum*
  - doesn't accept or ack out-of-order packet
- Sender maintains timer for *oldest unacked* packet
  - if timer expires, *retransmit all* unack'ed packets

## Selective Repeat:

- Sender can have up to N unack'ed packets in pipeline
  - sliding window
- Rcvr sends *individual ack* for each packet
  - maintains *Rcvr window*
  - buffers and acks all packets within Rcvr window
- Sender maintains timer for *each unacked* packet
  - when timer expires, *retransmit only that one* unack'ed packet



## GBN vs SR: Performance

- Compared to Stop-and-Wait
  - They both fill the pipeline
- Loss rate
  - Light loss:
    - SR: selectively retransmits what is needed
    - GBN: a single packet lost causes unnecessary retransmission of all packets in the window
  - Heavy loss
    - GBN: ok
- Complexity:
  - GBN is simpler – less state