Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

TCP flow control

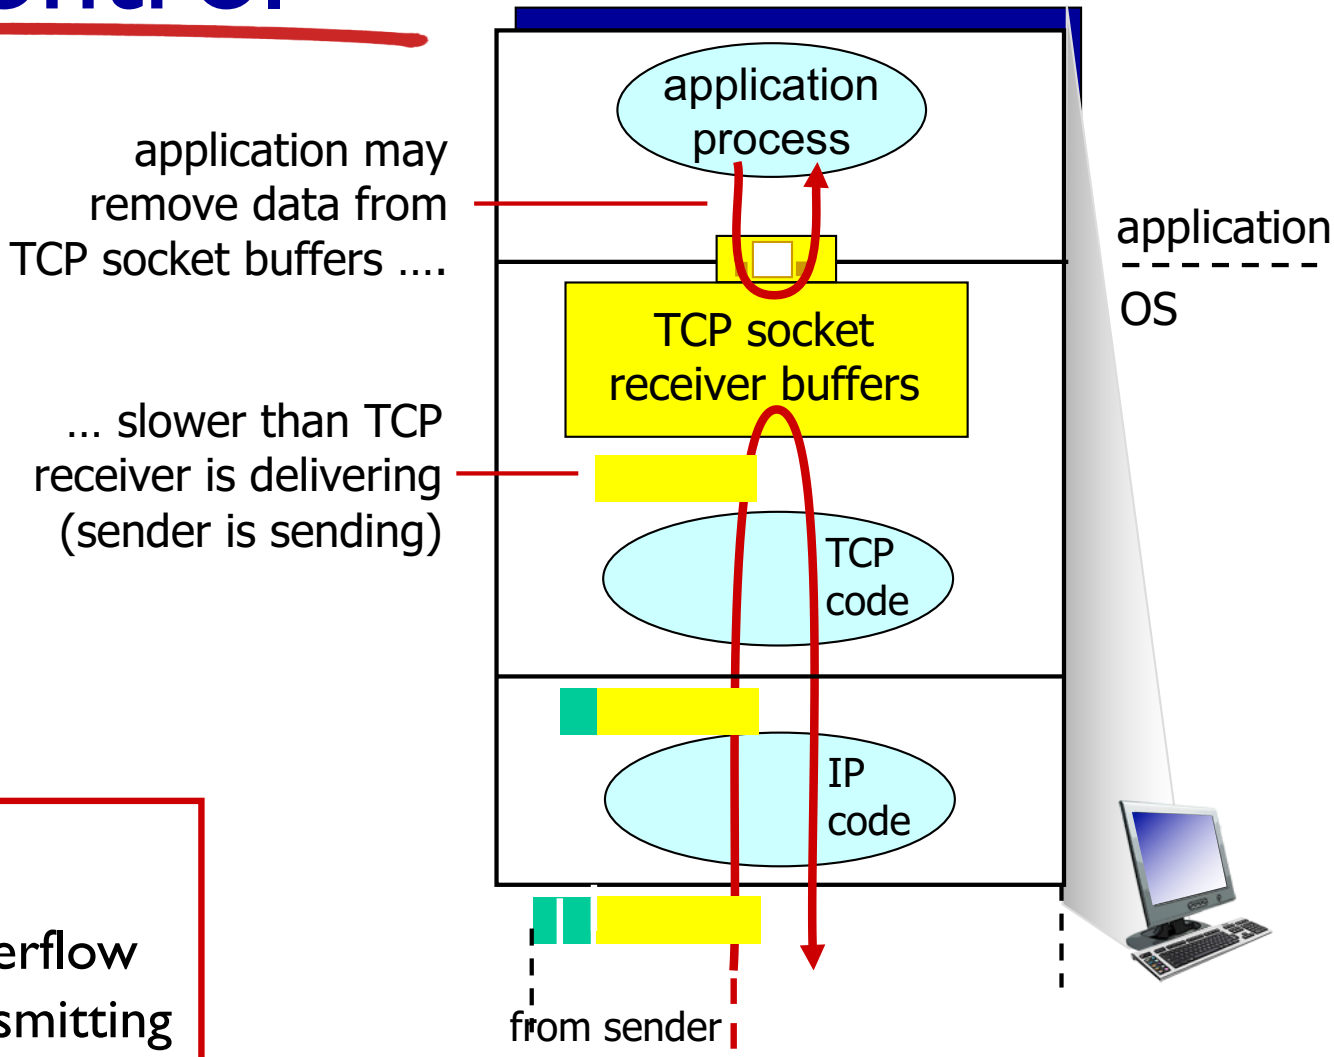

flow control

receiver controls sender,
so that sender won't overflow
receiver's buffer by transmitting
too much, too fast

TCP flow control

- receiver "advertises" free buffer space by including **`rwnd`** value in TCP header of receiver-to-sender segments
  - **`RcvBuffer`** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **`RcvBuffer`**
- sender limits amount of unacked ("in-flight") data to receiver's **`rwnd`** value
- guarantees receive buffer will not overflow

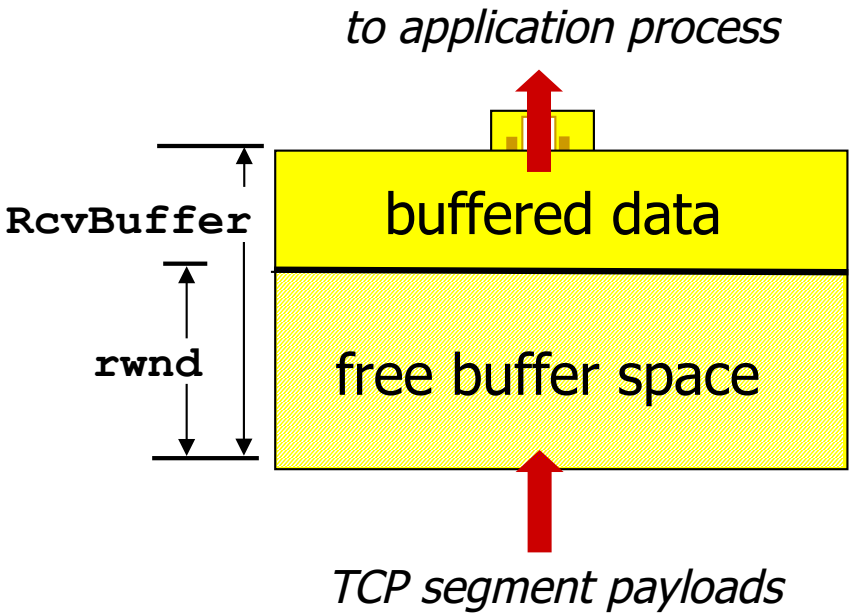

receiver-side buffering

TCP segment structure

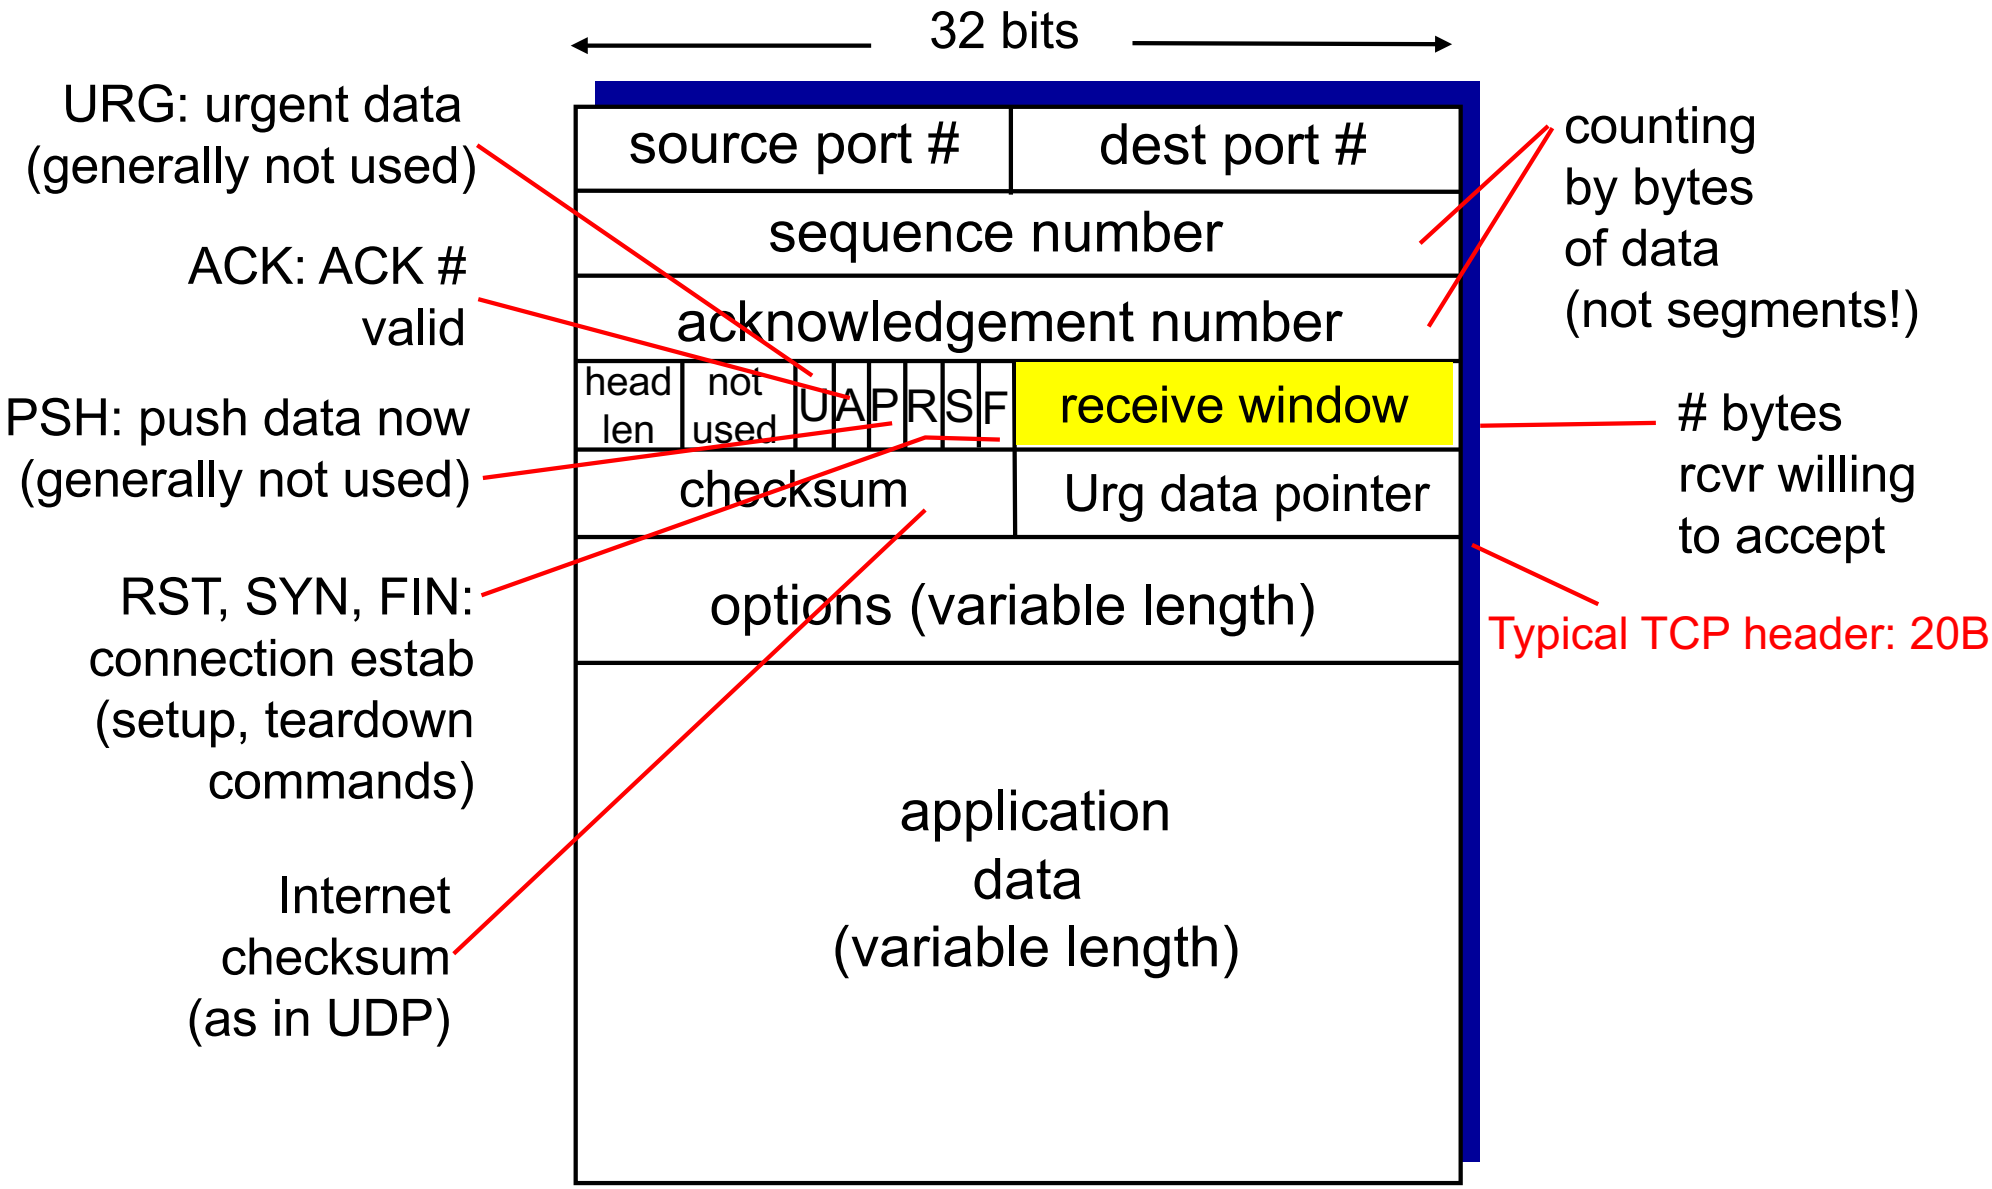

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

TCP segment structure

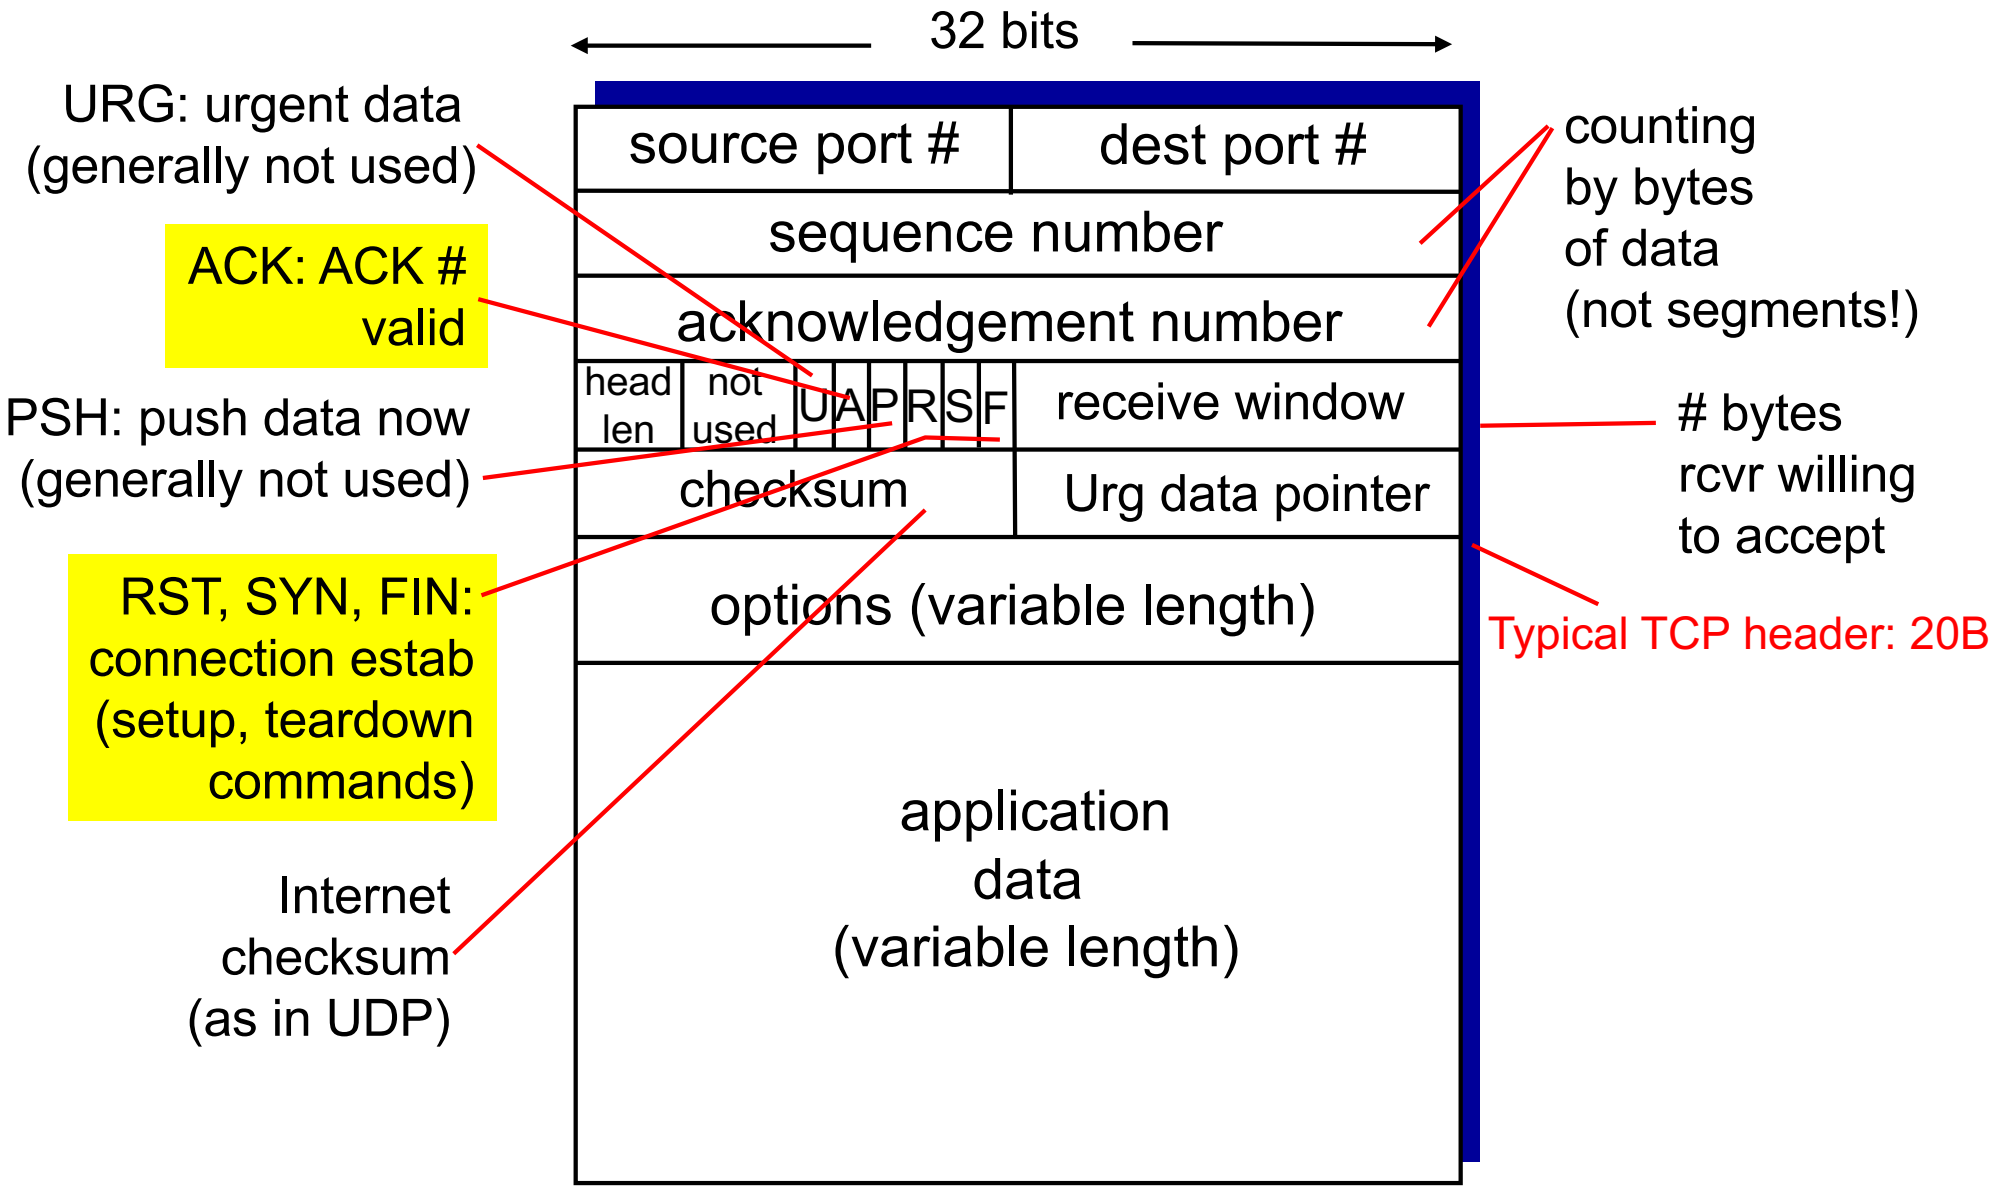

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

TCP segment structure

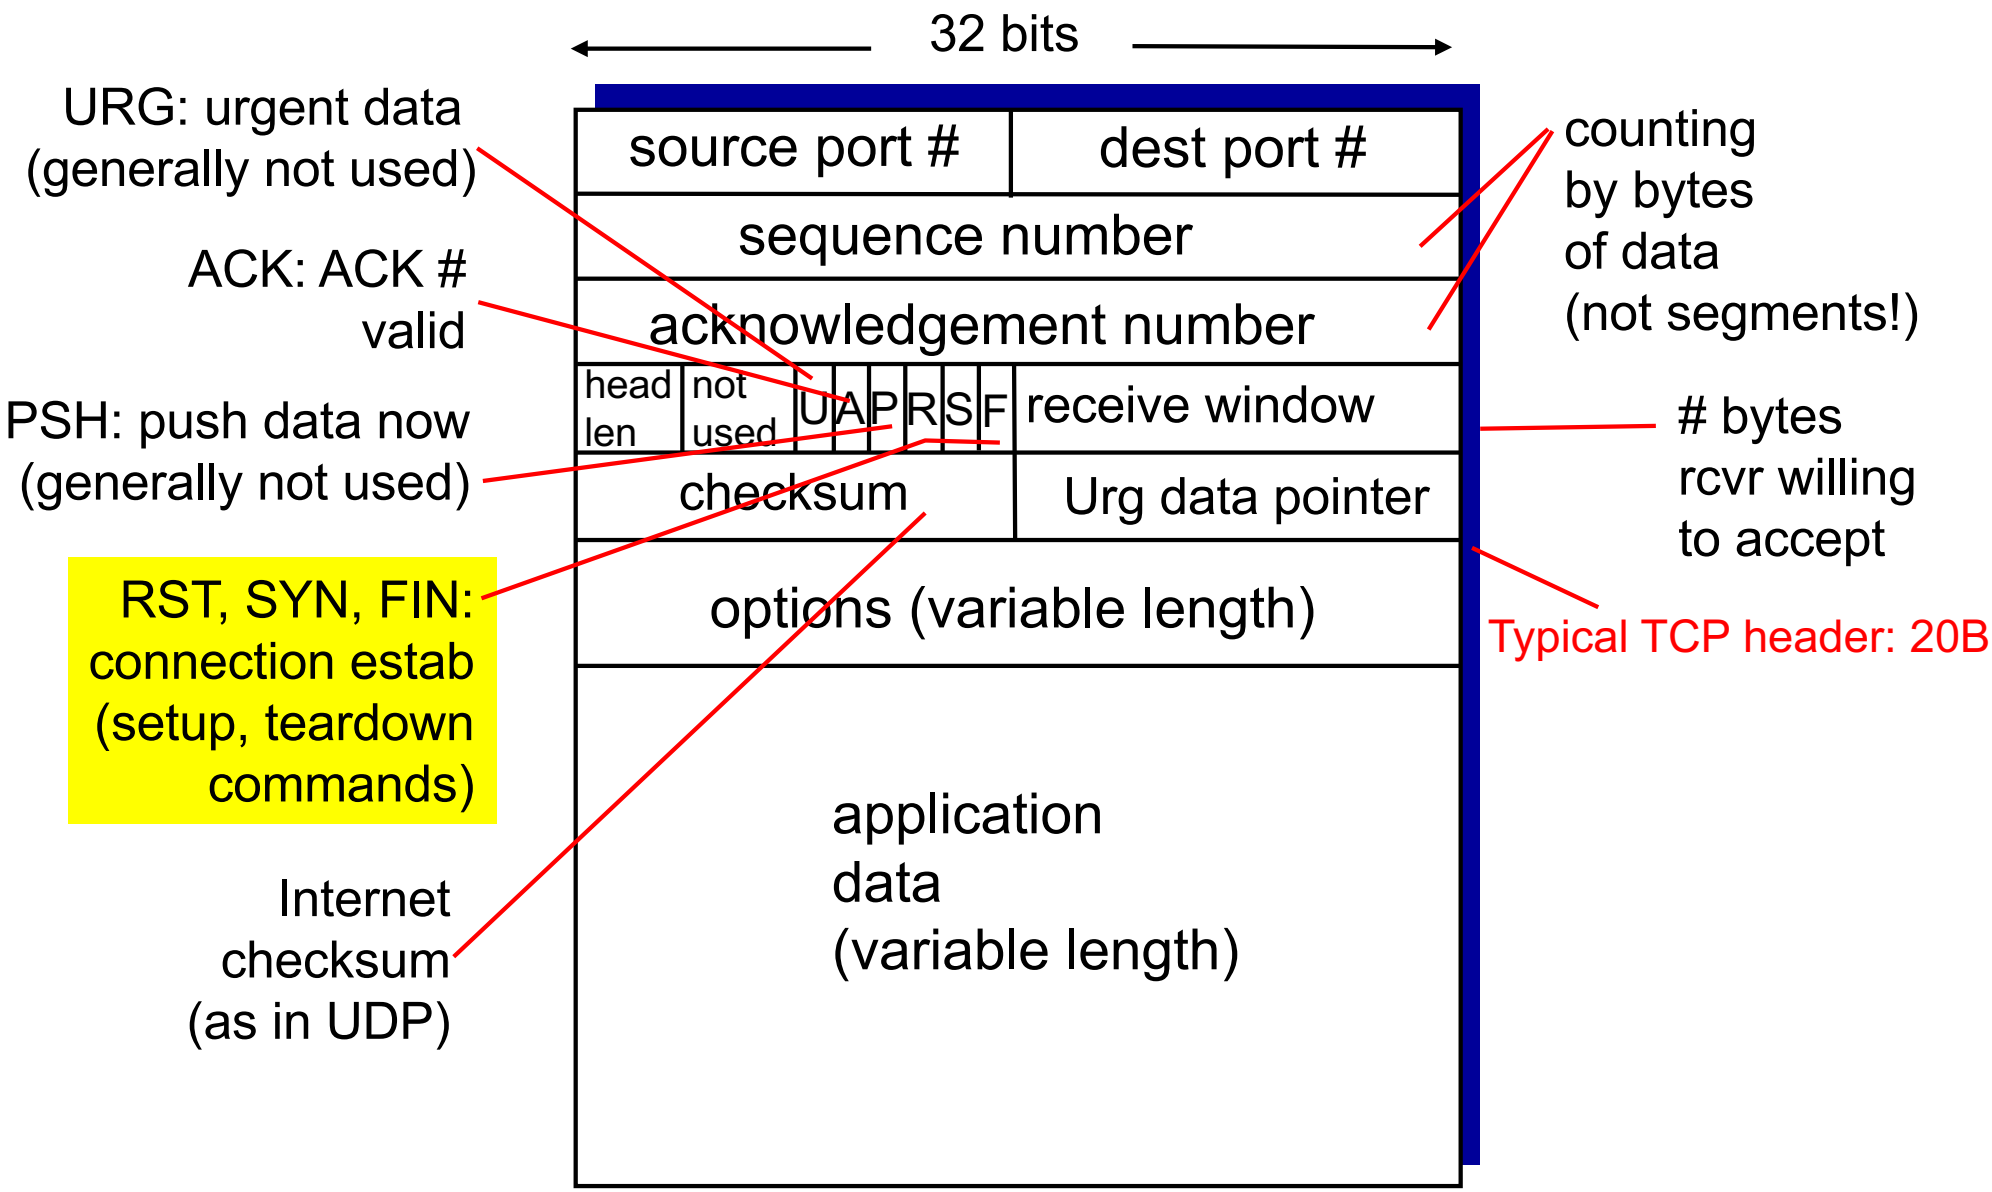

Connection Management

before exchanging data, sender/receiver "handshake":

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters

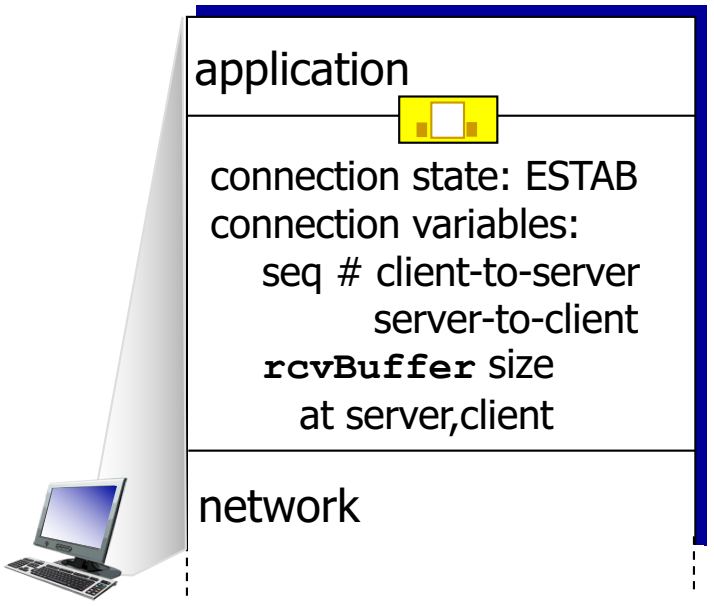

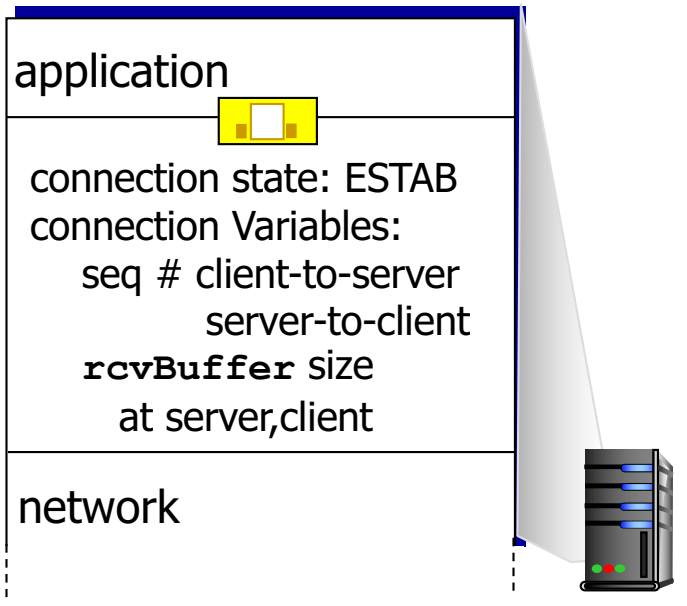


```
Socket clientSocket =
  newSocket("hostname","port
   number");
```

```
Socket connectionSocket =
   welcomeSocket.accept();
```

Agreeing to establish a connection

2-way handshake:

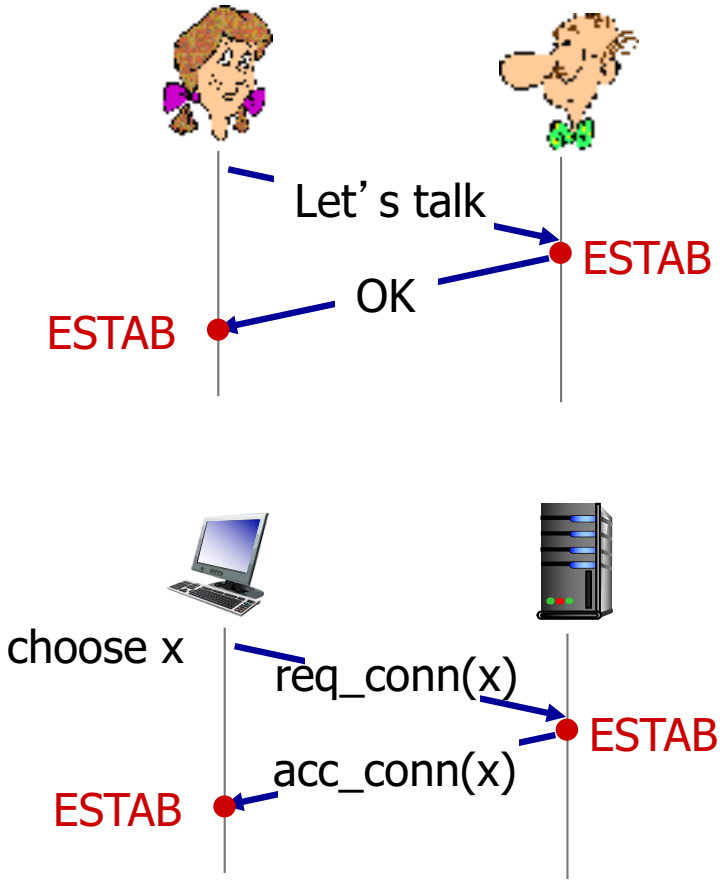

Q: will 2-way handshake always work in network?

- variable delays
- retransmitted messages (e.g. req_conn(x)) due to message loss
- message reordering
- can't "see" other side

Agreeing to establish a connection

2-way handshake failure scenarios:

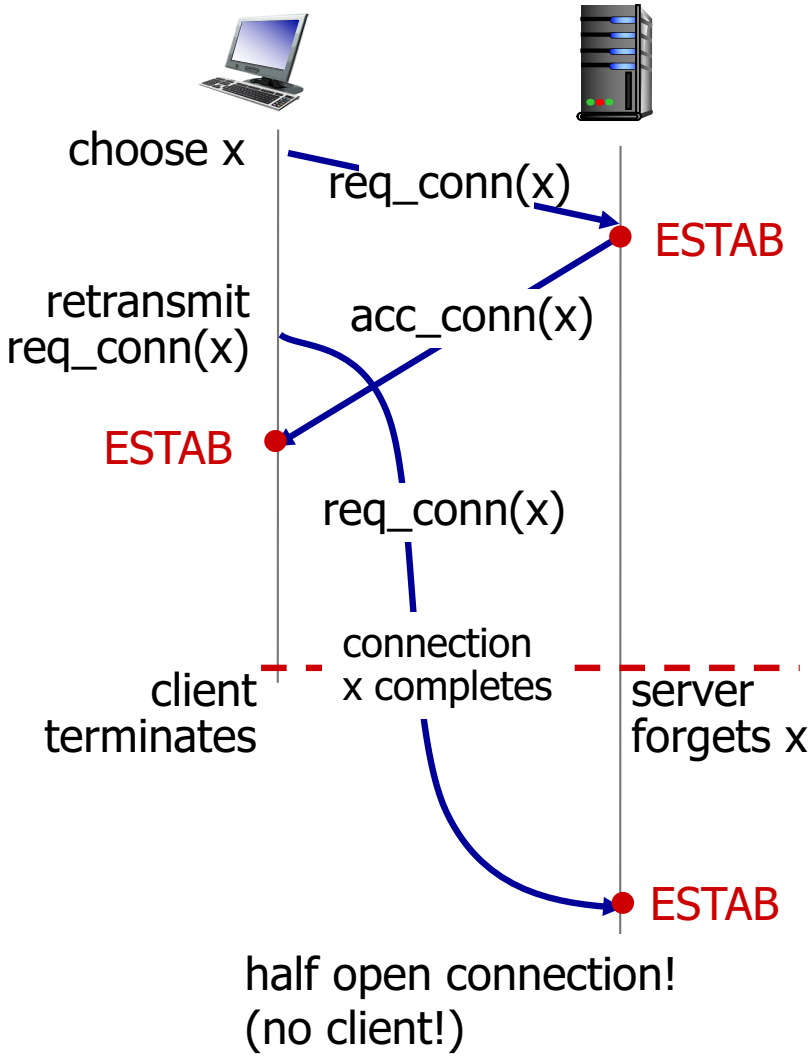

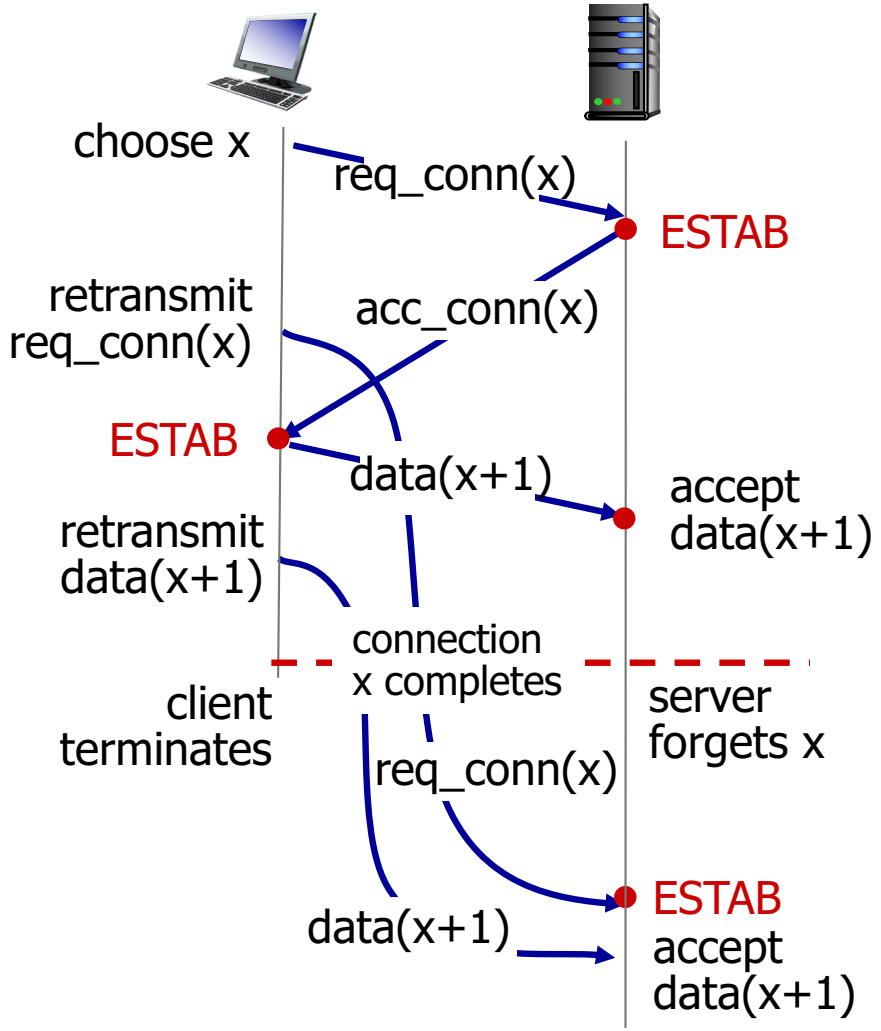

TCP 3-way handshake

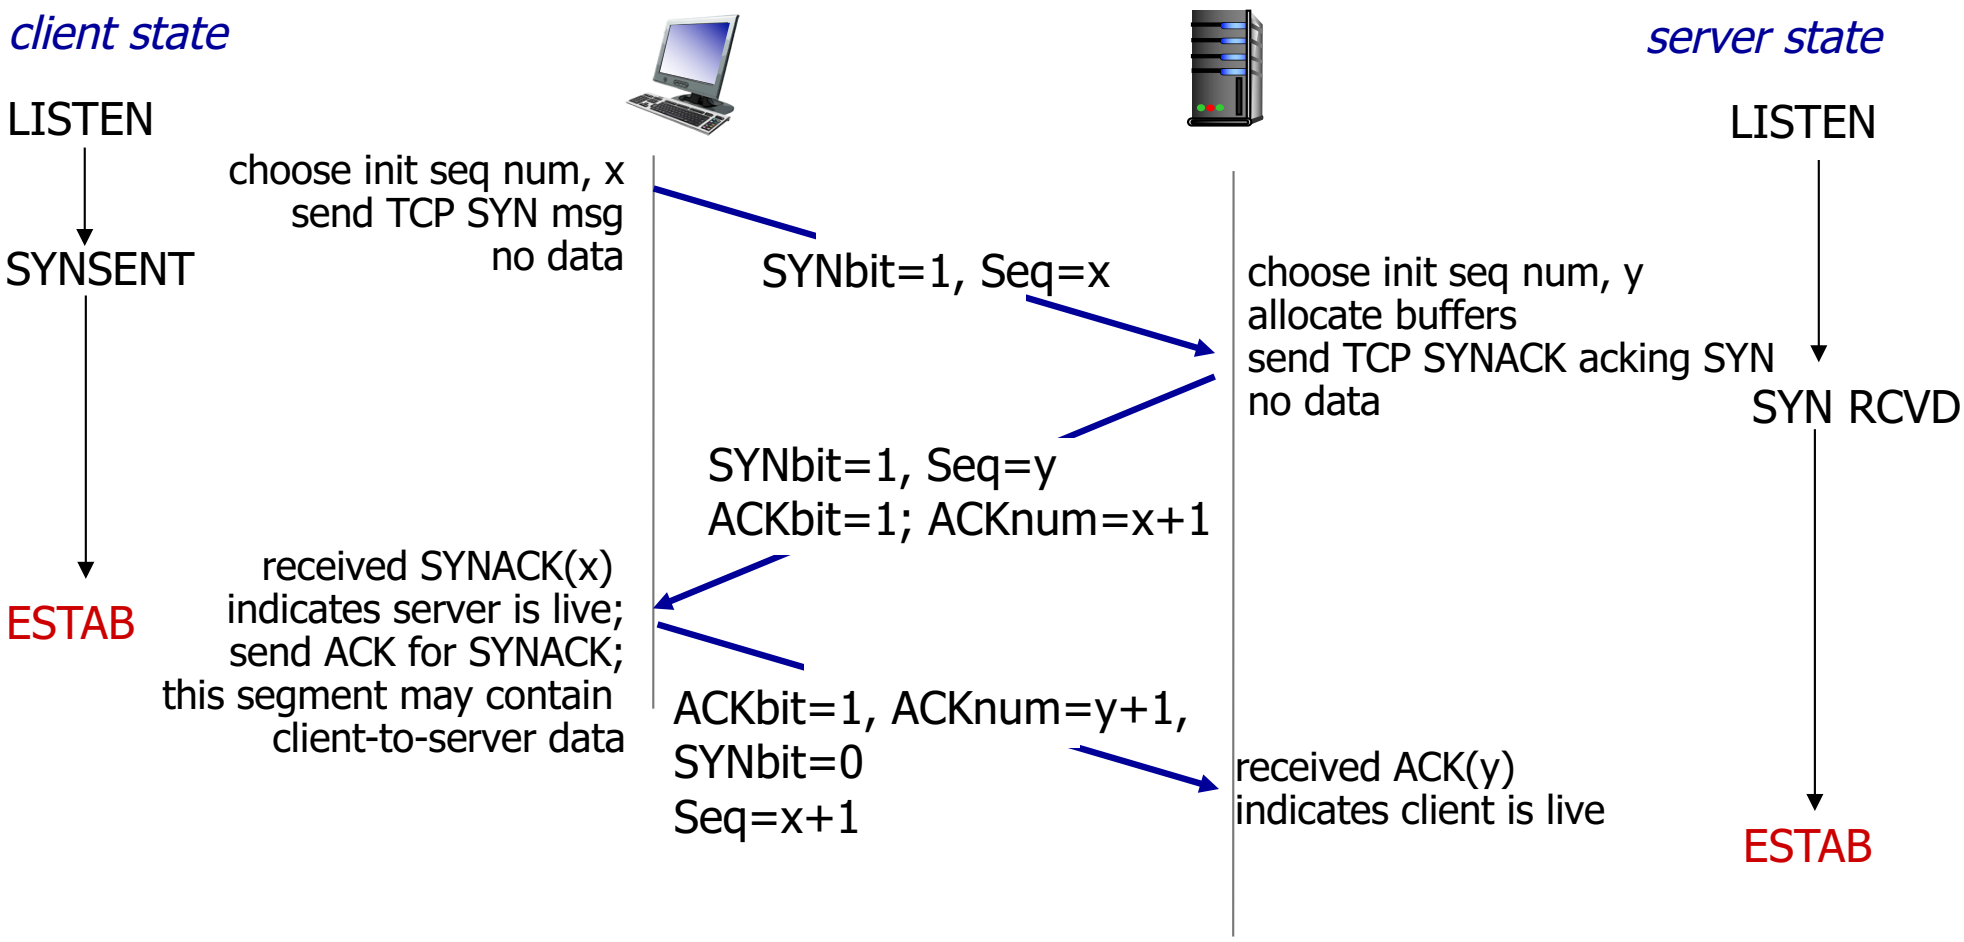

TCP 3-way handshake: FSM

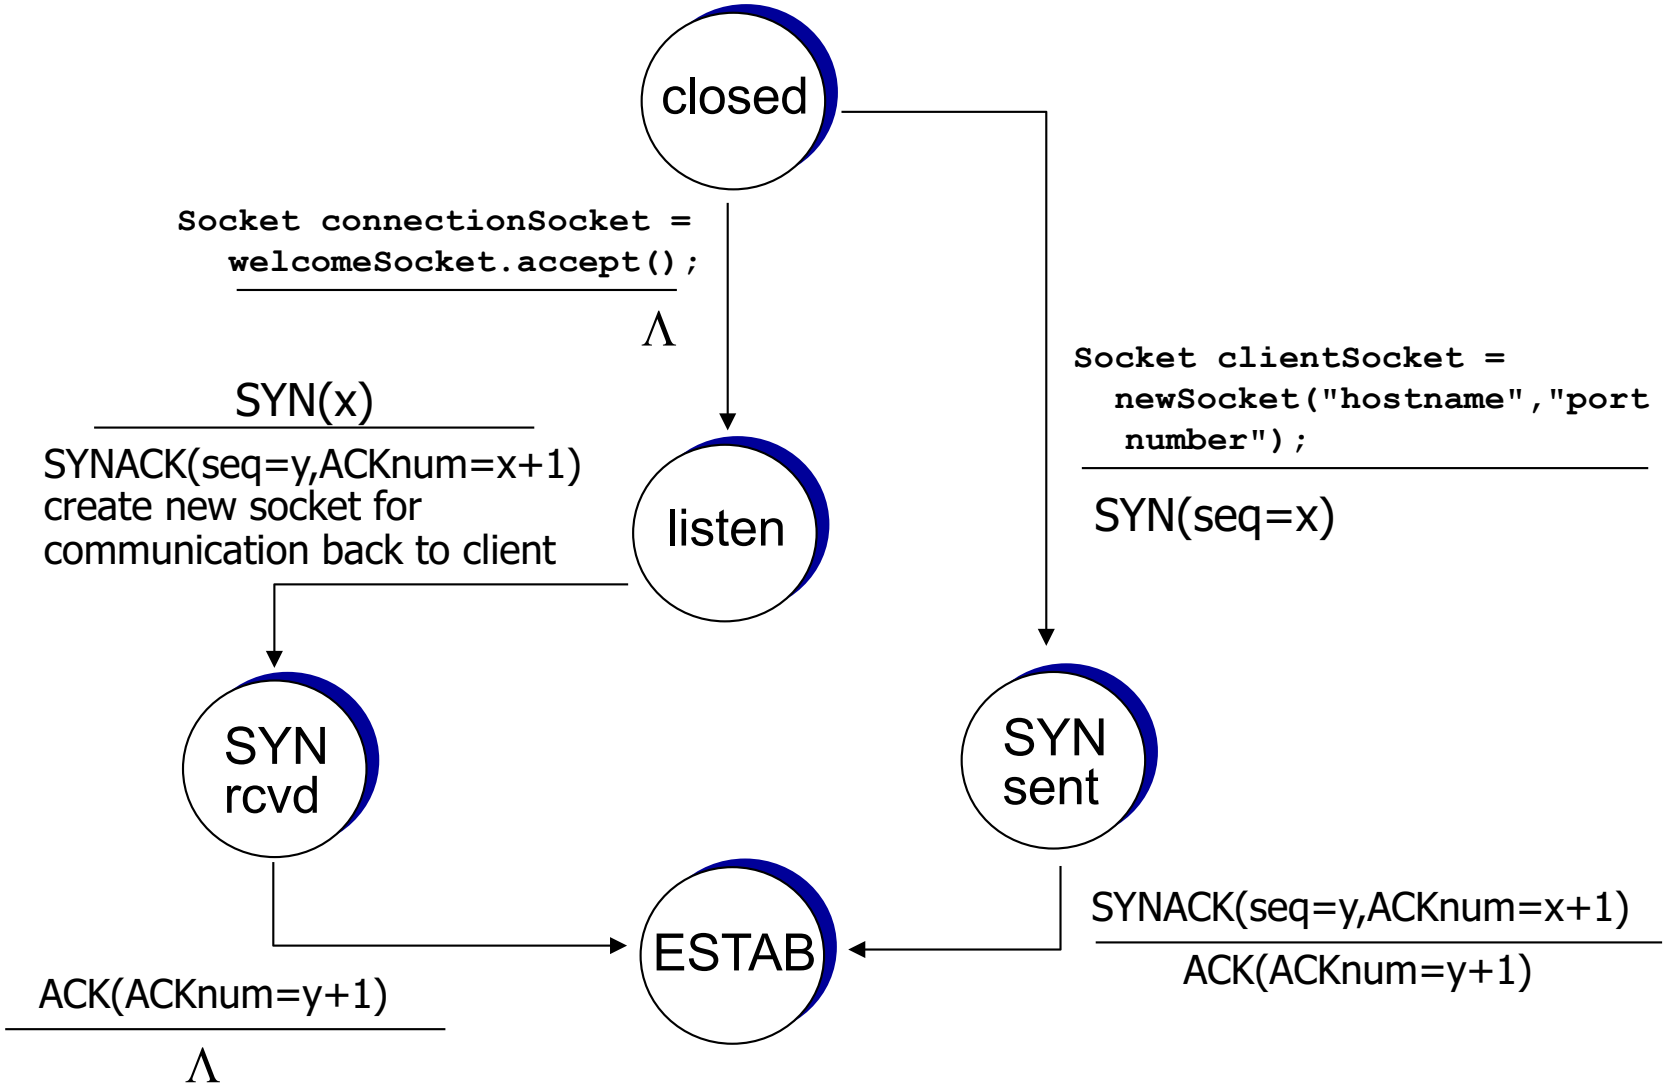

Server may not accept the connection

- Why?
  - Server may not be accepting TCP connections to that port
  - Server may be out of resources
- What happens?
  - Sends RST
  - No connection established, no resources allocated
  - But client learns that port is open
- UDP servers do not have connections
  - just listen to a socket on a dest port#
  - If server receives a UDP packet with dest port# that does not match an existing UDP socket → Sends ICMP message back

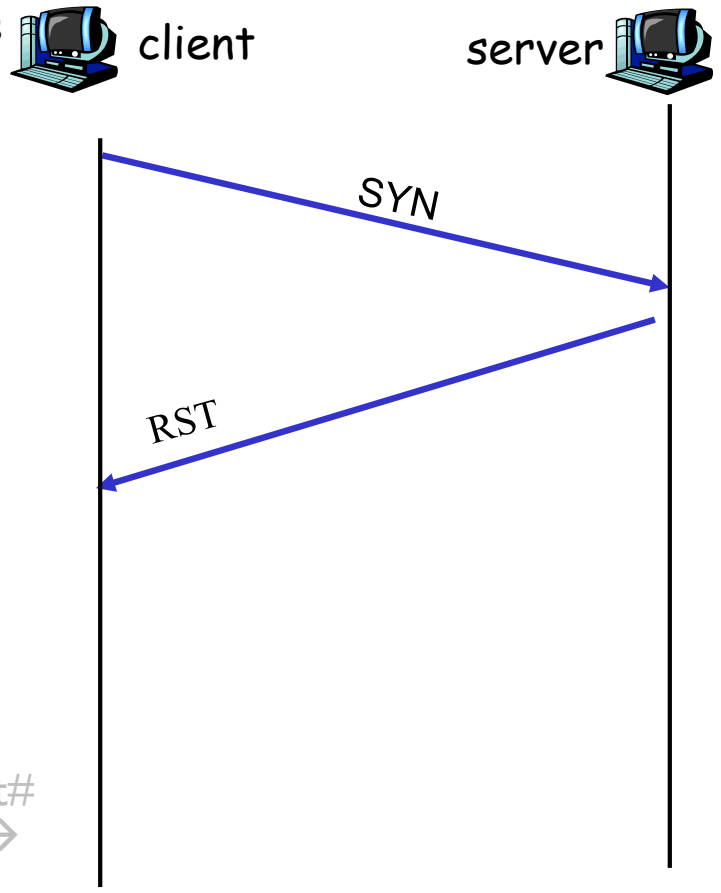

Attack 0: Scanning ports

- www.nmap.org

- Scanning TCP ports
  - Send TCP SYN
  - receive SYNACK, RST, nothing

- Scanning UDP ports
  - Receive ICNP messages

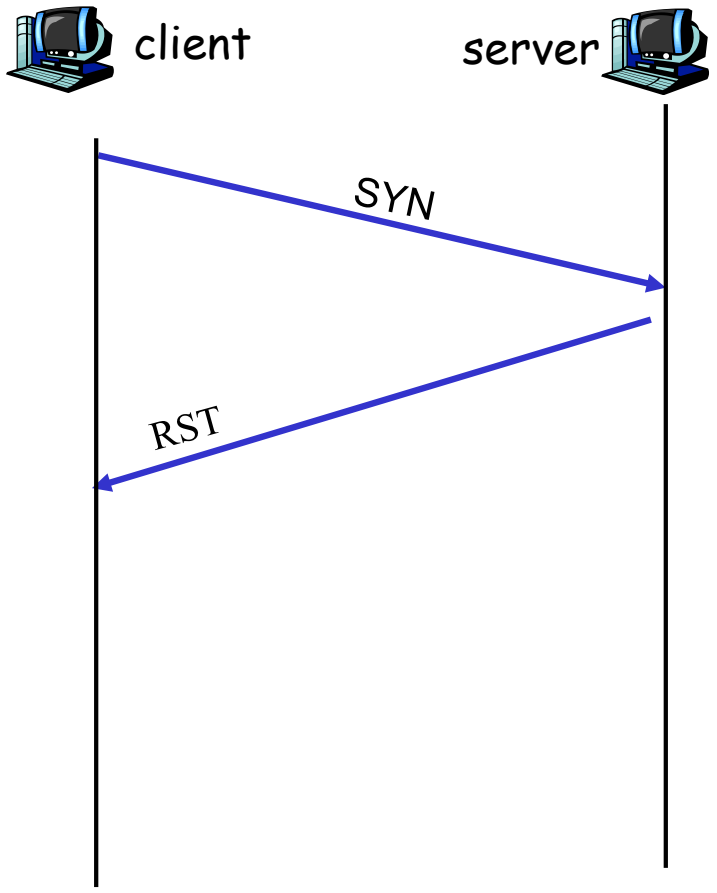

Attack I: SYN FLOOD attacks

- Client never completes the handshake
- Server allocates resources
- SYN FLOOD: the oldest (D)DoS attack
  - Server does not have resources to server legit clients
- A Solution: SYN Cookies
  - Server creates a cookie:=a hash of the client IP, port on the SYN segment (and of a secret number known to the server
  - Server sends SYNACK with that initial seqno=cookie, but does NOT allocates resources (half open connection)
  - If client is legit, it will send an ACK with initialseqno; server can verify that ACK corresponds to SYN; only then sender allocates resources (full open connection)

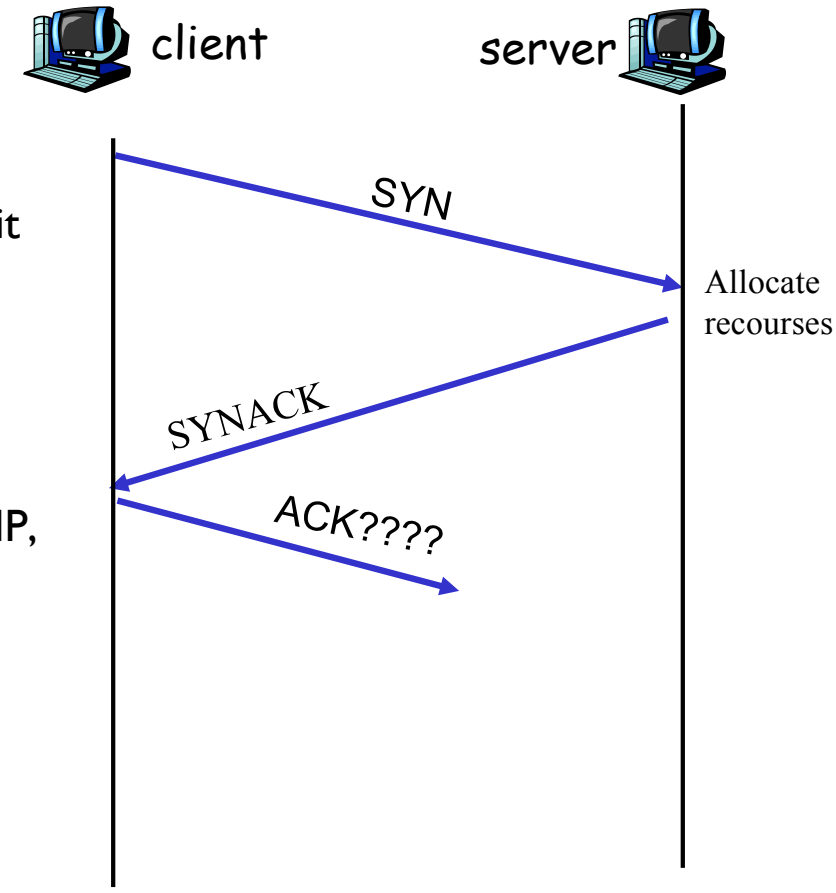

Attack II: Spoofing

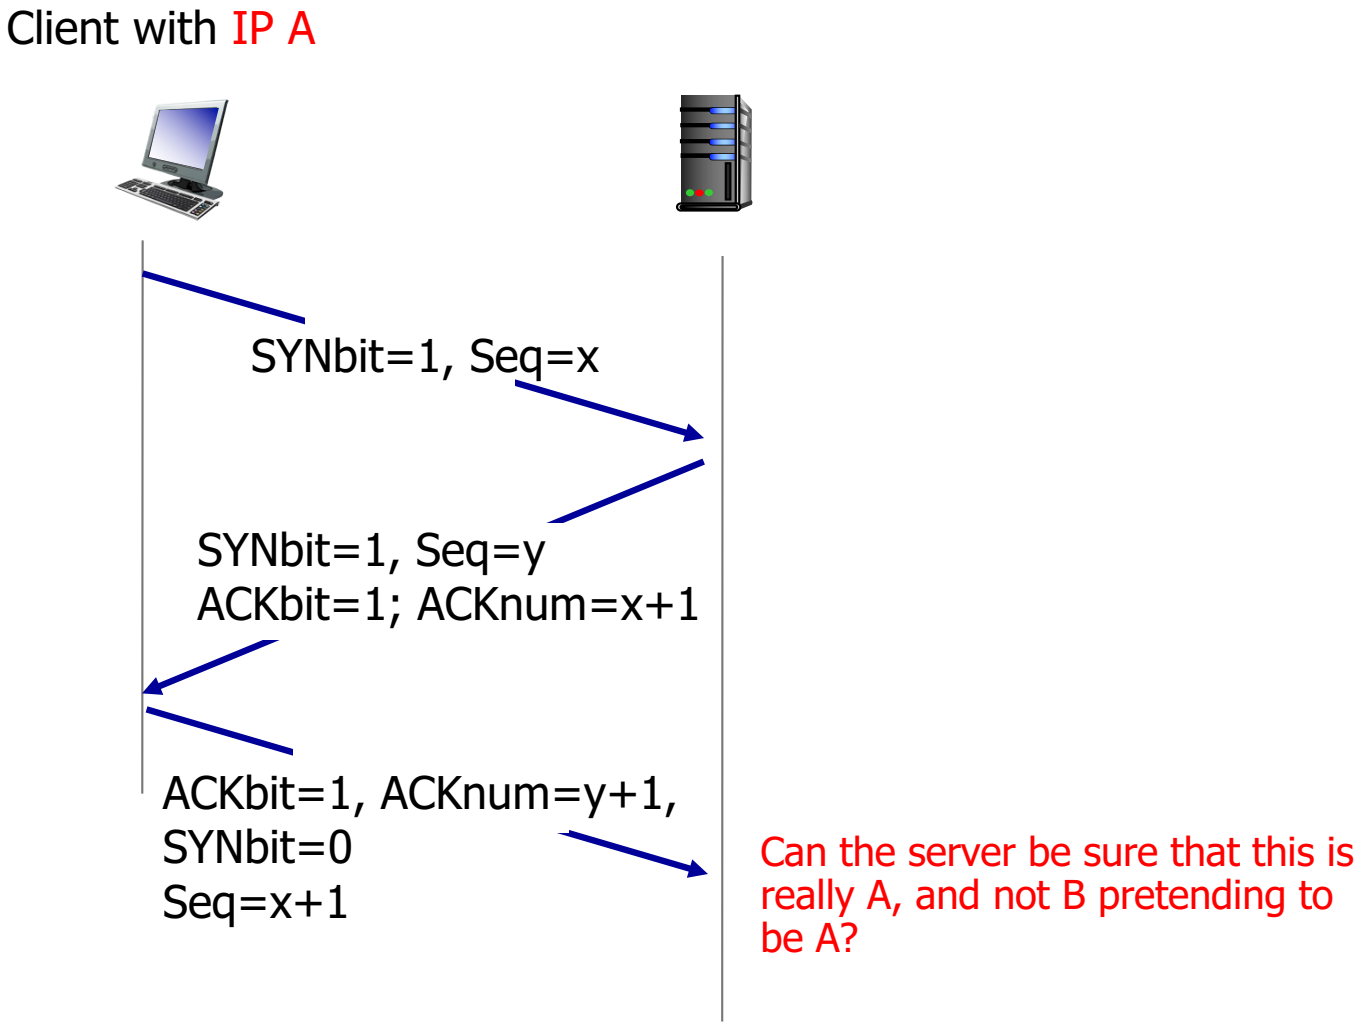

Attack II: Spoofing

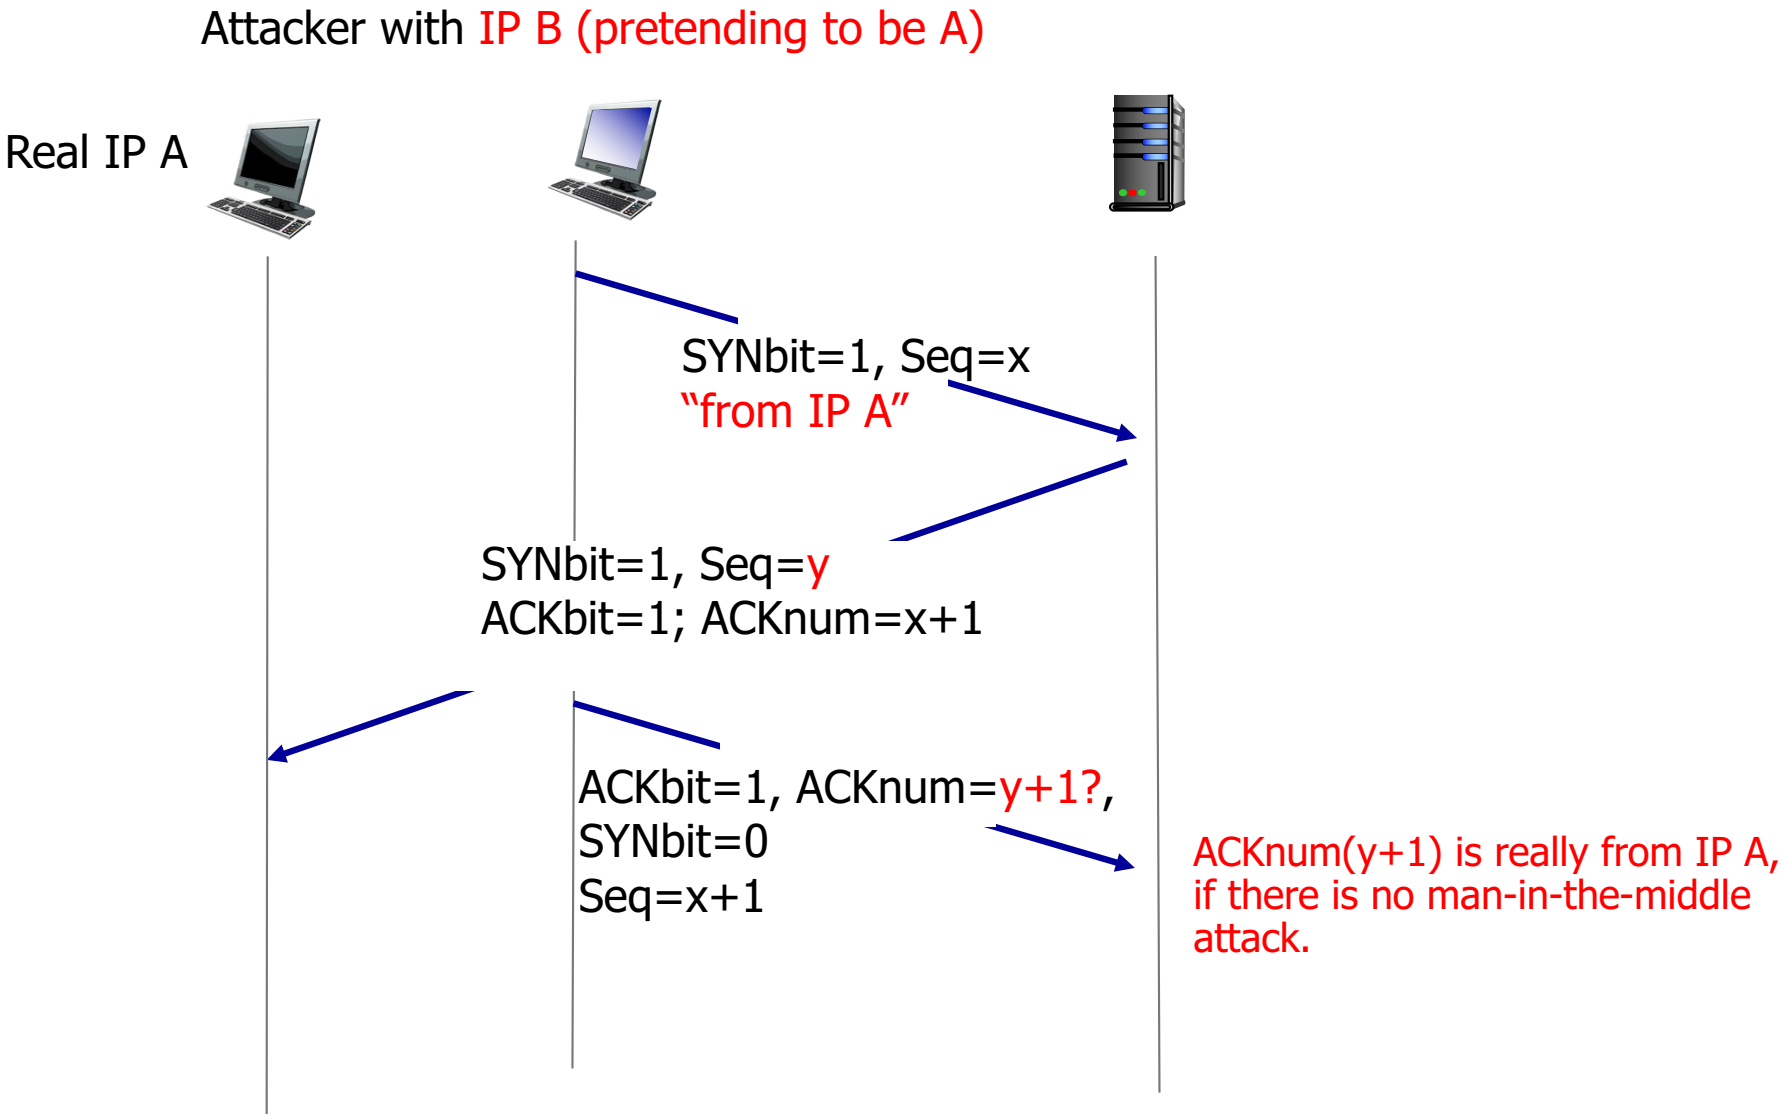

TCP: closing a connection

- remember: this is a duplex connection
- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
  - either of the two can initiate the closing
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

TCP: closing a connection

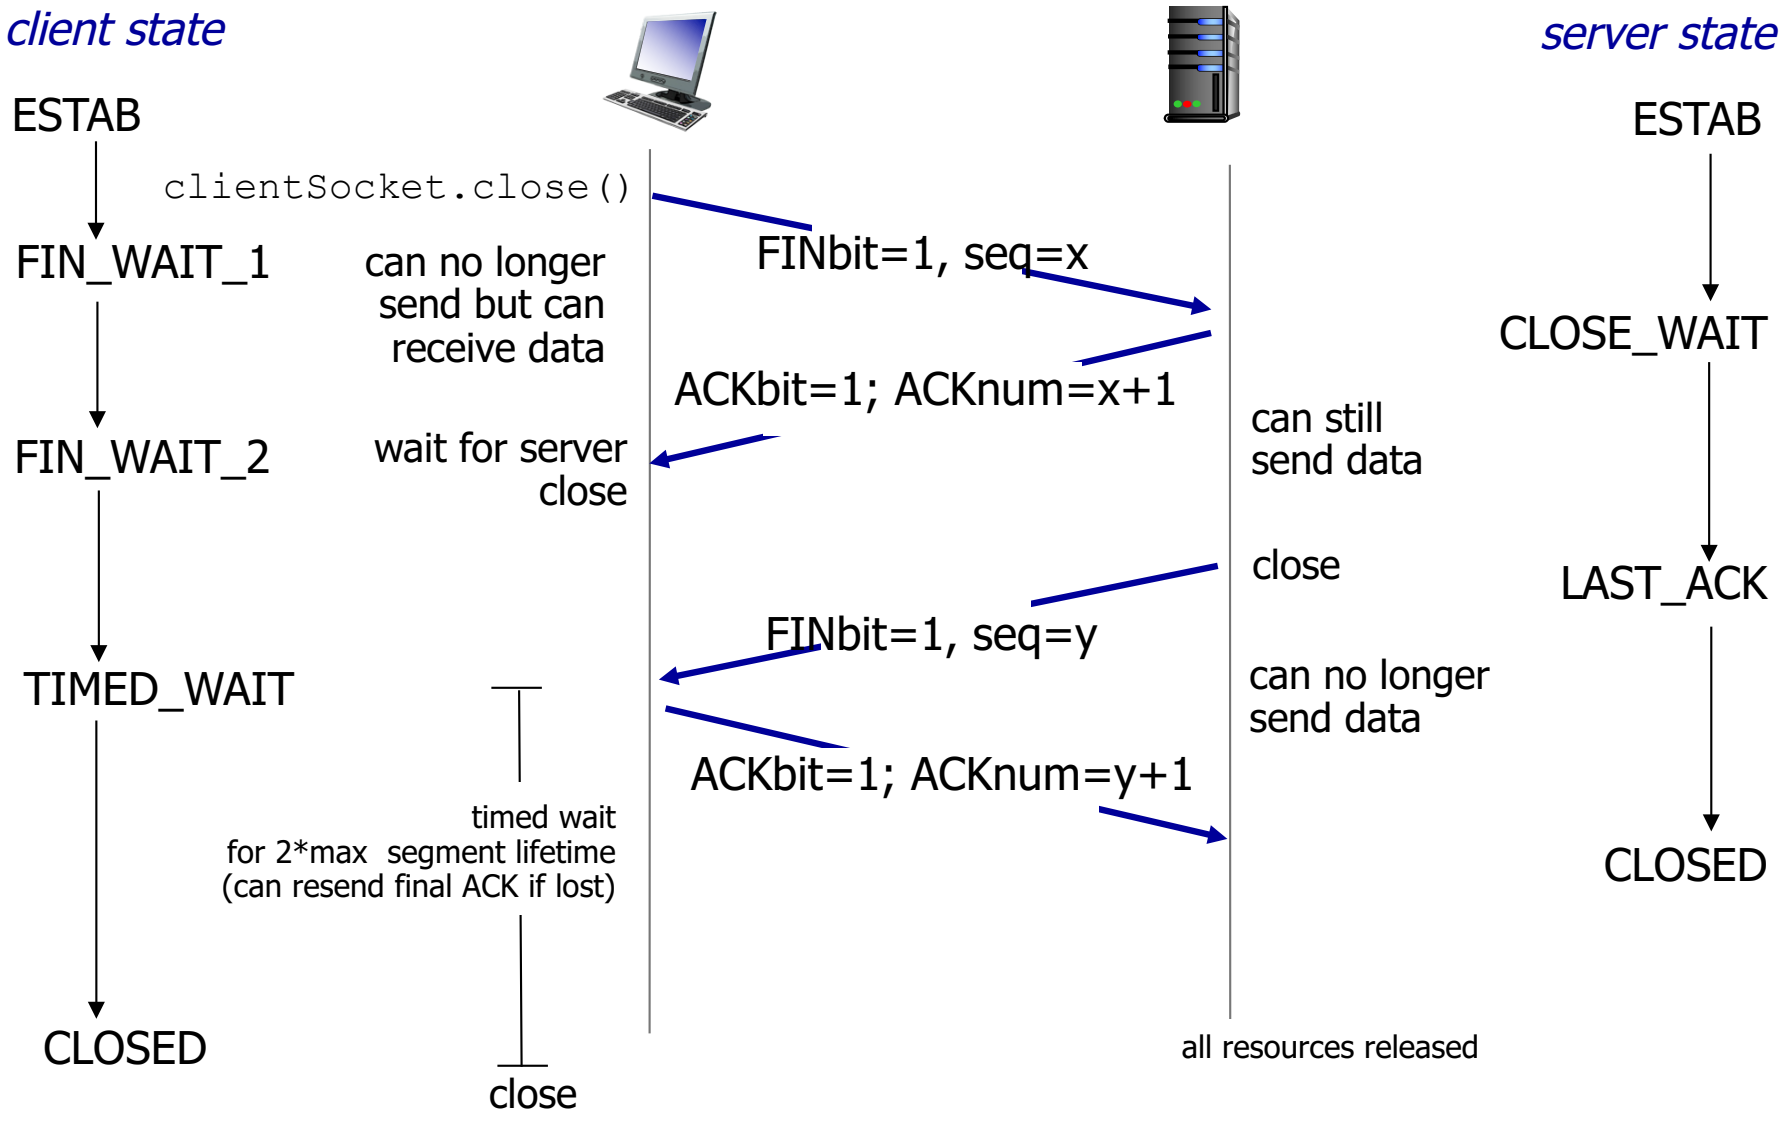

TCP Connection States

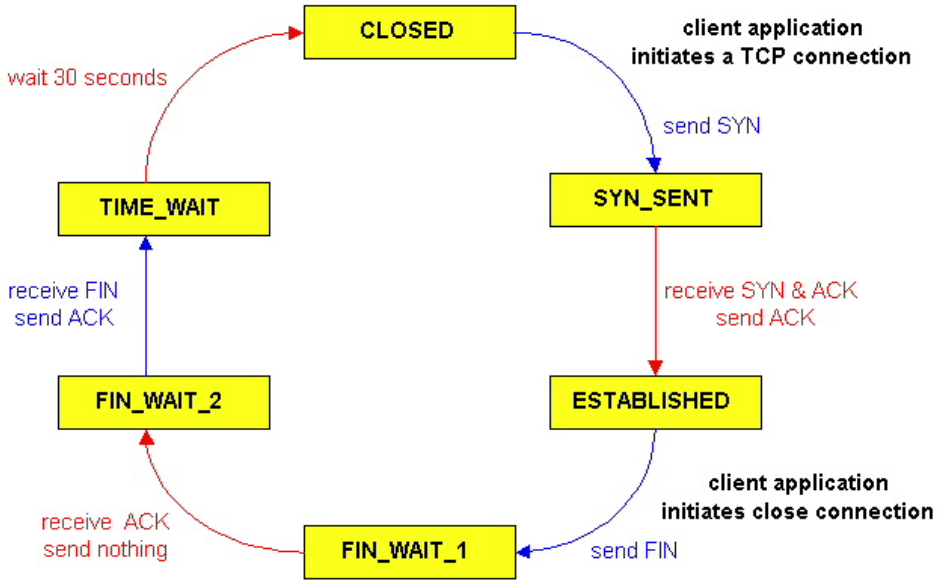

TCP client lifecycle

TCP server lifecycle