

# Chapter 4: outline

## 4.1 Overview of Network layer

- data plane
- control plane

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

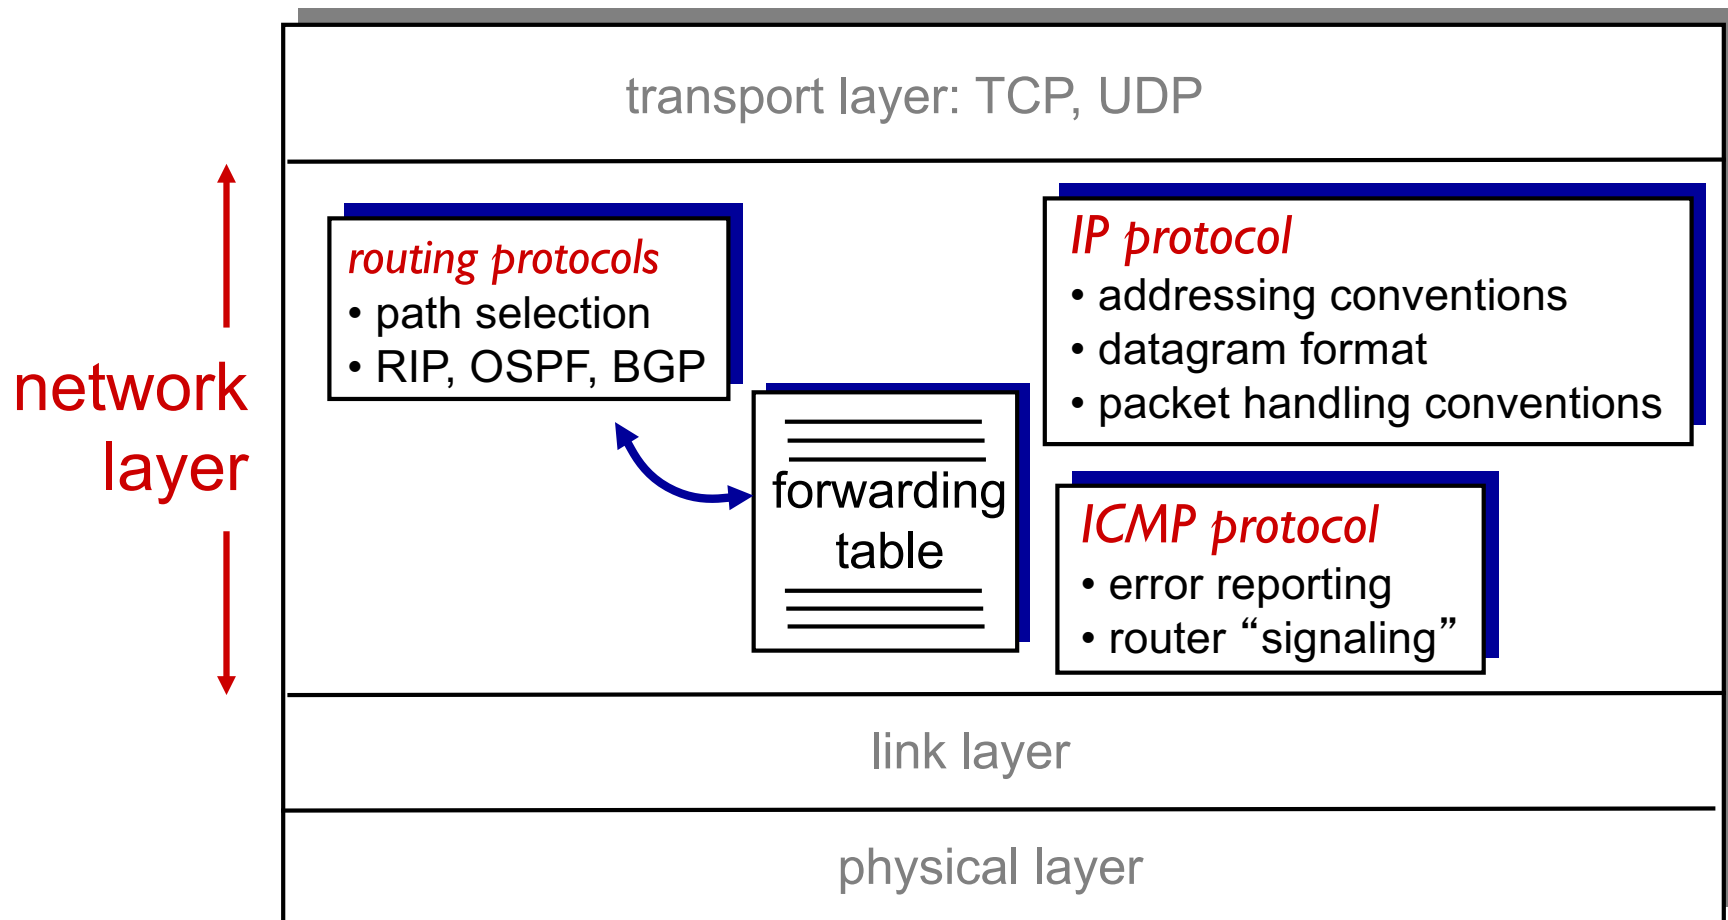
- datagram format
- fragmentation
- IPv4 addressing
- network address translation
- IPv6

## 4.4 Generalized Forward and SDN

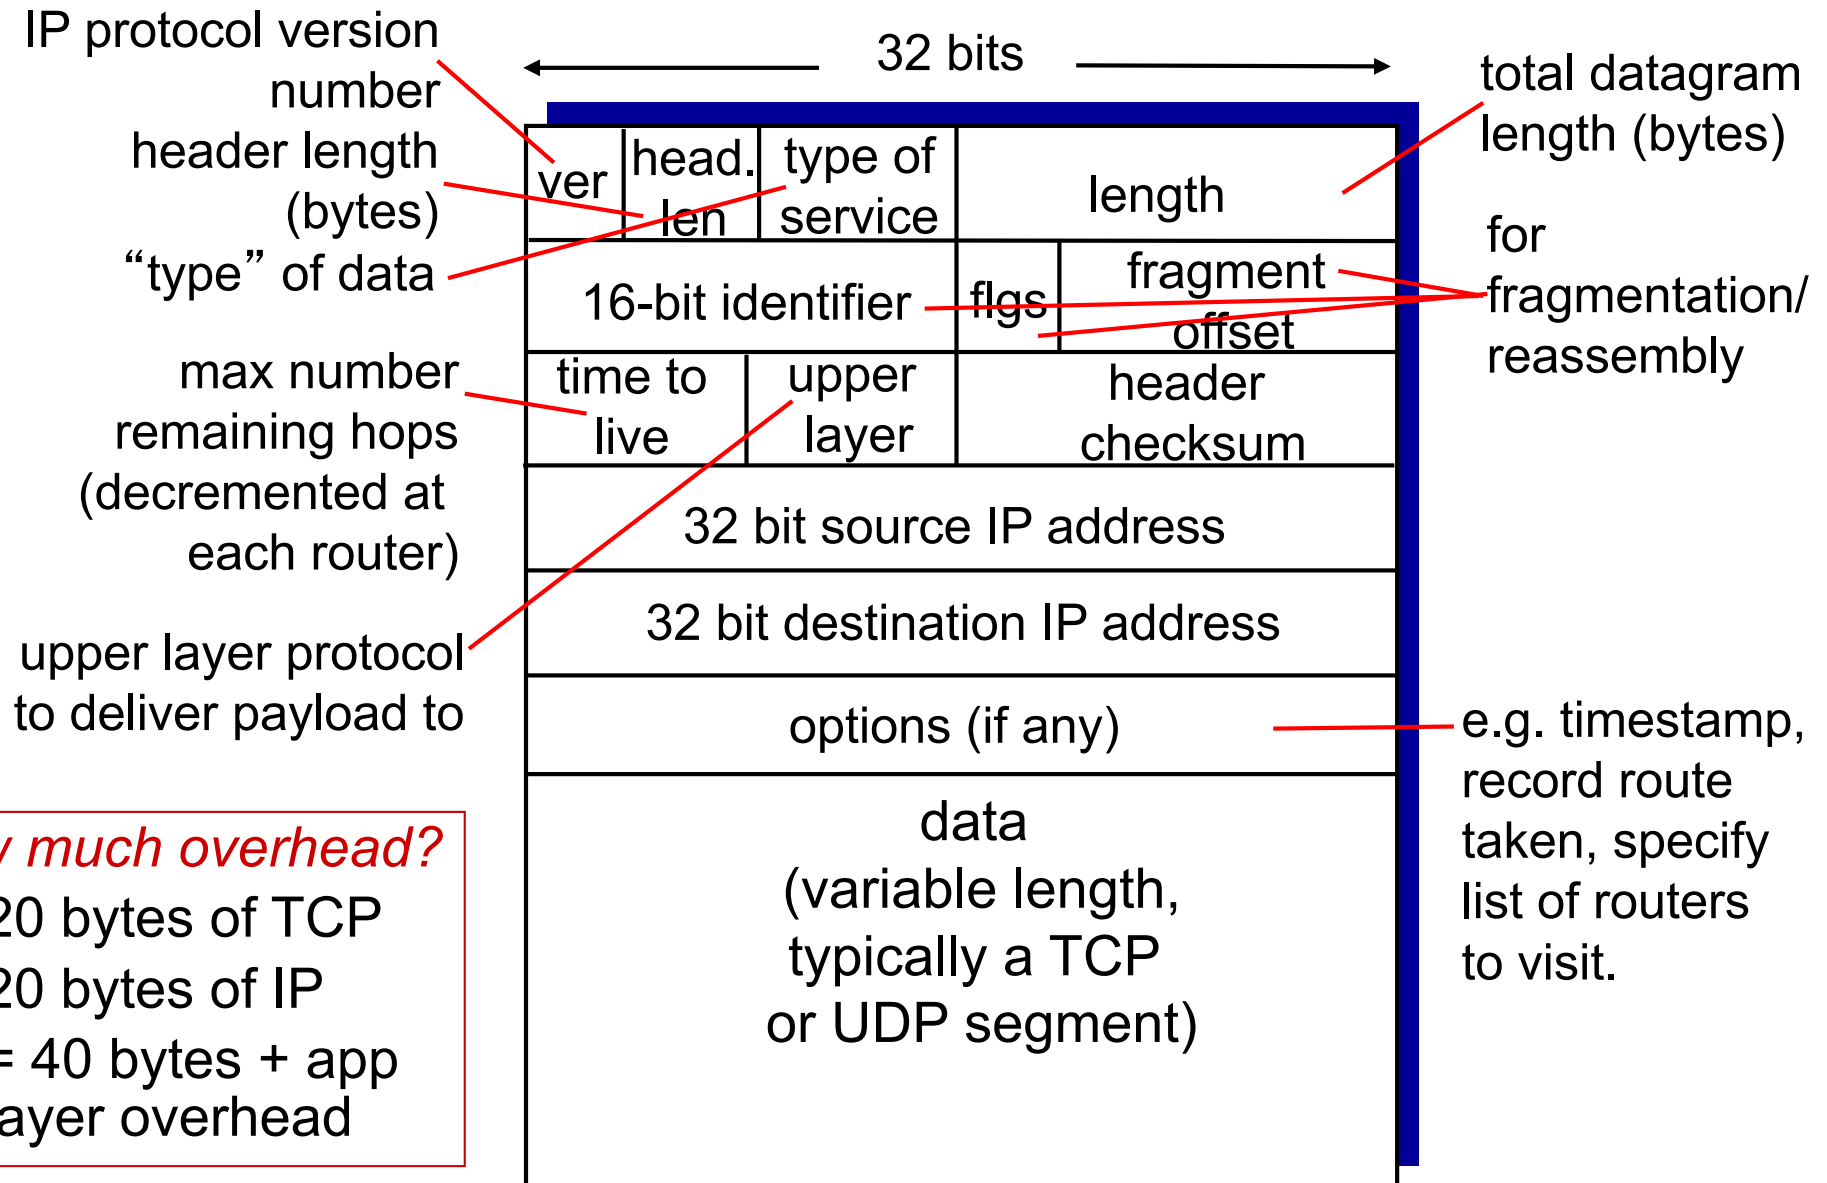
- match
- action
- OpenFlow examples of match-plus-action in action

# The Internet network layer

host, router network layer functions:

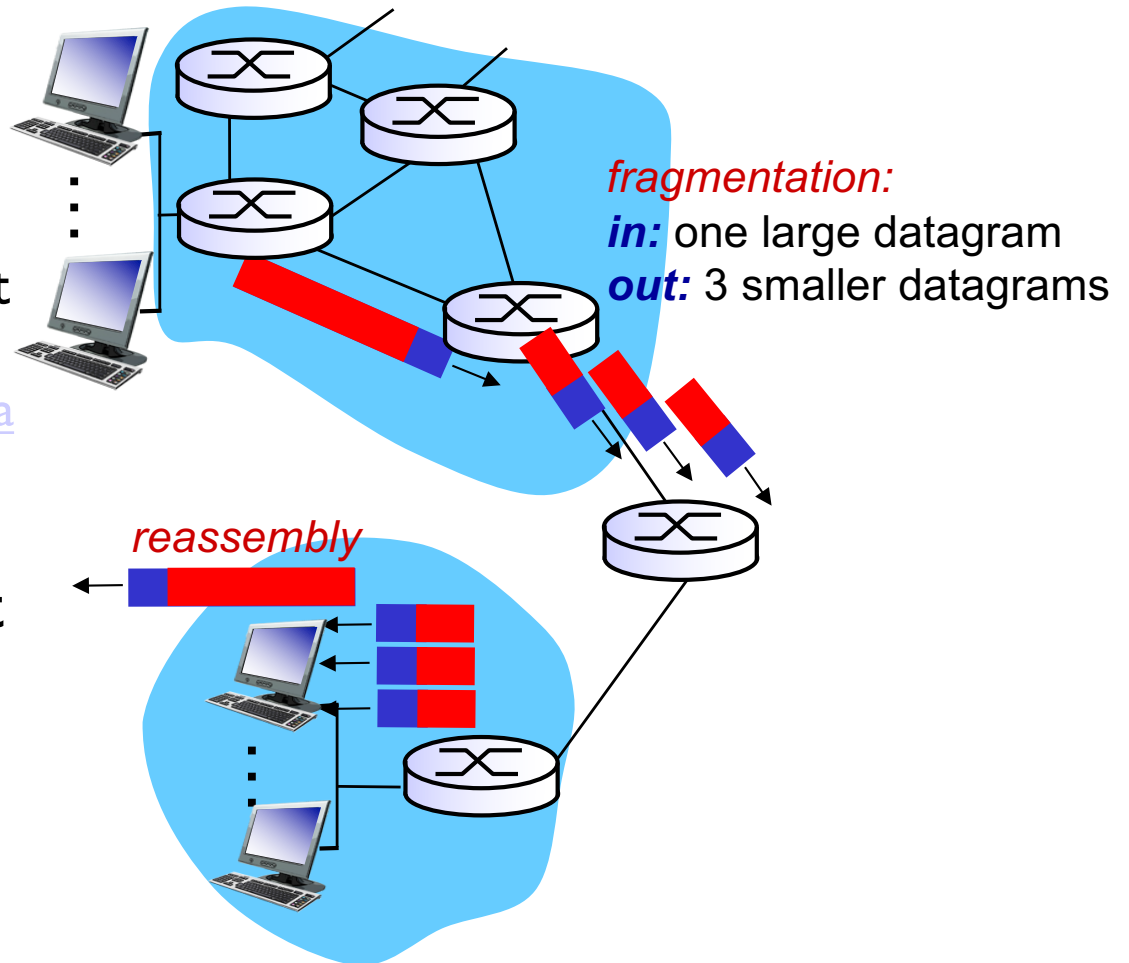


# IP datagram format



# IP fragmentation, reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame
  - different link types, different MTUs
  - [https://en.wikipedia.org/wiki/Maximum\\_transmission\\_unit](https://en.wikipedia.org/wiki/Maximum_transmission_unit)
- large IP datagram divided (“fragmented”) within net
  - one datagram becomes several datagrams
  - “reassembled” only at final destination
  - IP header bits used to identify, order related fragments



# IP fragmentation, reassembly

ID: unique per packet,  
incremented by sending  
host

offset: multiple of 8.

*example:*

- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

	length	ID	fragflag	offset	
	=4000	=x	=0	=0	

*one large datagram becomes  
several smaller datagrams*

1480 bytes in  
data field

offset =  
 $1480/8$

offset =  
 $2*1480/8$

	length	ID	fragflag	offset	
	=1500	=x	=1	=0	

	length	ID	fragflag	offset	
	=1500	=x	=1	=185	

	length	ID	fragflag	offset	
	=1040	=x	=0	=370	

Original datagram:  $20+3980=4000$   
 1<sup>st</sup> fragment:  $20+1480=1500$  starts at byte 0  
 2<sup>nd</sup> fragment:  $20+1480=1500$  starts at byte  $185*8=1480$   
 3<sup>rd</sup> fragment:  $20+1020=1040$  starts at byte  $(185+185)*8=2960$   
 Data:  $1480+1480+1020=3980$

# Chapter 4: outline

## 4.1 Overview of Network layer

- data plane
- control plane

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol


- datagram format
- fragmentation
- IPv4 addressing
- network address translation
- IPv6

## 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

# IP addressing: introduction

- **IP address:** 32-bit identifier for host, router *interface*
- **interface:** connection between host/router and physical link
  - router's typically have multiple interfaces
  - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- **IP addresses associated with each interface**



$$223.1.1.1 = \underbrace{11011111}_{223} \underbrace{00000001}_1 \underbrace{00000001}_1 \underbrace{00000001}_1$$


# IP addressing: introduction

*Q: how are interfaces actually connected?*

*A: we'll learn about that in chapter 5, 6.*

*A: wired Ethernet interfaces connected by Ethernet switches*

*For now:* don't need to worry about how one interface is connected to another (with no intervening router)




*A: wireless WiFi interfaces connected by WiFi base station*



# Subnets

- IP address:
  - subnet part - high order bits
  - host part - low order bits
- *What is a subnet ?*
  - device interfaces with same subnet part of IP address
  - can physically reach each other *without intervening router*




network consisting of 3 subnets

# Subnets

## *recipe*


- to determine the subnets, detach each interface from its host or router, creating islands of isolated networks
- each isolated network is called a *subnet*



subnet mask: /24

# Subnets

how many?



# Historic Classful Network Architecture:

Class	Starting with (bits)	Range of first byte (decimal)	Network id format	Host id format	Number of networks	Number of hosts
A	0	0-127	a	b.c.d	$2^7 = 128$	$2^{24} = 16777216$
B	10	128-191	a.b	c.d	$2^{14} = 16384$	$2^{16} = 65536$
C	110	192-223	a.b.c	d	$2^{21} = 2097152$	$2^8 = 256$

← subnet part → ← host part →  
11001000 00010111 00010000 00000000


Class C network (“/24”): 200.23.16.0

Example IP in that network: 200.23.16.1

# (Since 1993) IP addressing: CIDR

## CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



Q: what did we gain?

A: more efficient use of the IP address space


# Some special IP addresses

- All 1's: means "all hosts on this subnet"



200.23.17.255/23

- All 0's: means "this subnet"



200.23.16.0/23

# More IP addresses

- Reserved IP addresses for special purposes
  - [https://en.wikipedia.org/wiki/Reserved\\_IP\\_addresses#IPv4](https://en.wikipedia.org/wiki/Reserved_IP_addresses#IPv4)
  - 127.0.0.1: local host
  - Multicast
    - 224.0.0.0/8–239.0.0.0/8
  - Private networks:
    - not routed, typically used through NATs.
    - 24- bit block, /8 prefix, 1xA: 10.0.0.0-10.255.255.255
    - 20-bit block, /12 prefix, 16xB: 172.16.0.0- 172.31.255.255
    - 16-bit block, /16 prefix, 256xC: 192.168.0.0-192.168.255.255
  - Assigned to special institutions

# IP addresses: how to get one?

**Q:** How does *a host* get IP address?

- hard-coded by system admin in a file
  - Windows: control-panel->network->configuration->tcp/ip->properties
  - UNIX: /etc/rc.config
  - Mac: /etc/resolv.conf
- **DHCP: Dynamic Host Configuration Protocol:**  
dynamically get address from as server
  - “plug-and-play”



# DHCP: Dynamic Host Configuration Protocol


*goal:* allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network (more shortly)

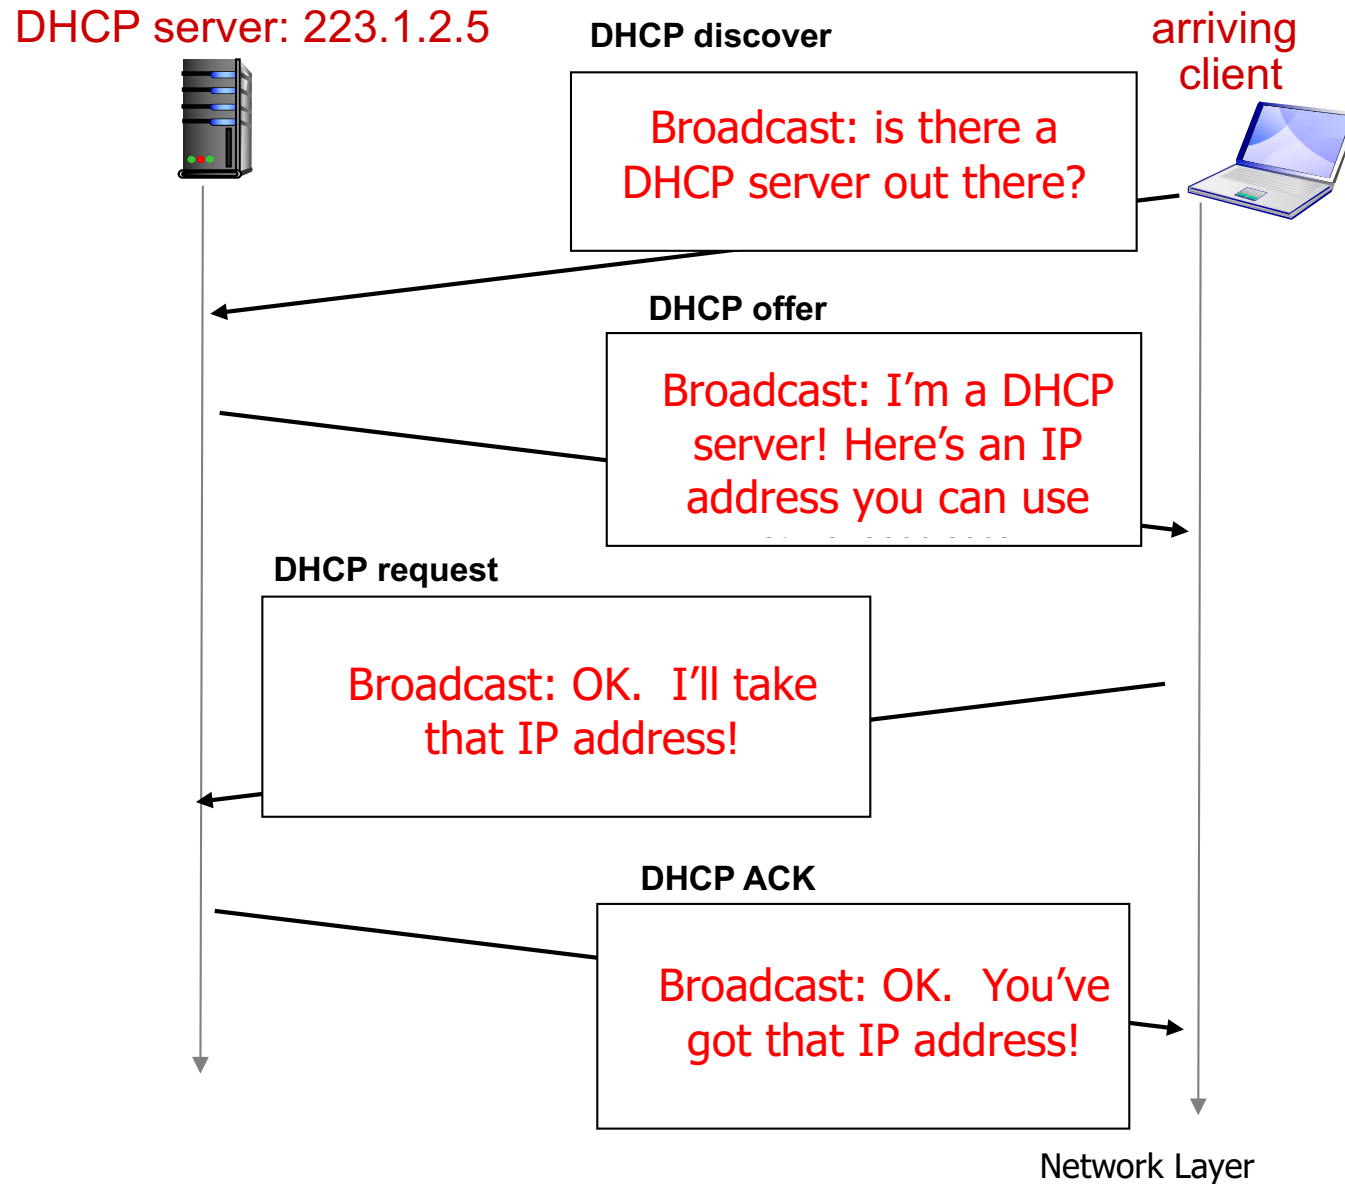
## *DHCP overview:*

- host broadcasts “DHCP discover” msg [optional]
- DHCP server responds with “DHCP offer” msg [optional]
- host requests IP address: “DHCP request” msg
- DHCP server sends address: “DHCP ack” msg

# DHCP client-server scenario



# DHCP client-server scenario



# DHCP client-server scenario

DHCP server: 223.1.2.5

**DHCP discover**

src : 0.0.0.0, 68  
dest.: 255.255.255.255,67  
yiaddr: 0.0.0.0  
transaction ID: 654

arriving  
client



**DHCP offer**

src: 223.1.2.5, 67  
dest: 255.255.255.255, 68  
yiaddr: 223.1.2.4  
transaction ID: 654  
lifetime: 3600 secs

**DHCP request**

src: 0.0.0.0, 68  
dest:: 255.255.255.255, 67  
yiaddr: 223.1.2.4  
transaction ID: 655  
lifetime: 3600 secs

**DHCP ACK**

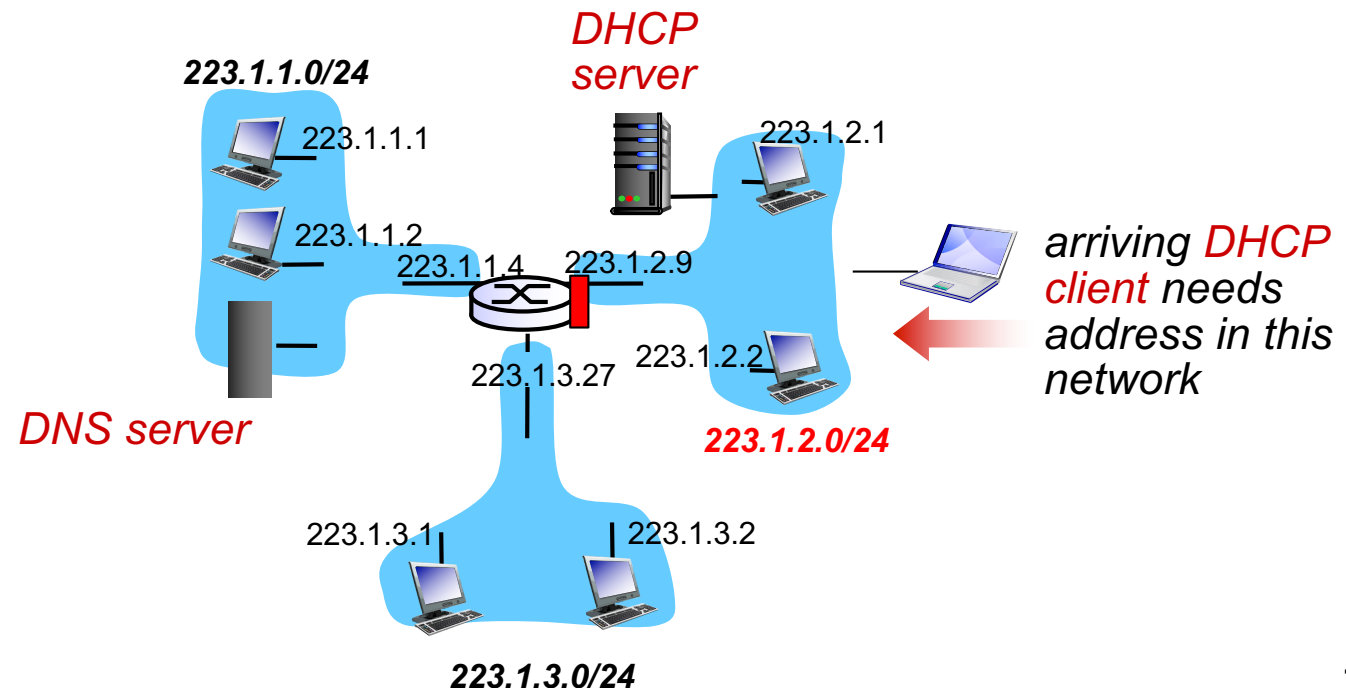
src: 223.1.2.5, 67  
dest: 255.255.255.255, 68  
yiaddr: 223.1.2.4  
transaction ID: 655  
lifetime: 3600 secs

Network Layer


# DHCP: more than IP addresses

DHCP can return more than just allocated IP address on subnet:

- address of first-hop router for client (**223.1.2.1**)
- name and IP address of DNS server
- network mask (indicating network versus host portion of address) **/24**




# DHCP: example



- connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- Ethernet frame broadcast (dest: FFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demuxed to IP demuxed, UDP demuxed to DHCP

# DHCP: example



- DCP server formulates DHCP ACK containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
- encapsulation of DHCP server, frame forwarded to client, demuxing up to DHCP at client
- client now knows its IP address, name and IP address of DSN server, IP address of its first-hop router

# DHCP: Wireshark output (home LAN)

Message type: **Boot Request (1)**

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

**Transaction ID: 0x6b3a11b7**

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

Client IP address: 0.0.0.0 (0.0.0.0)

Your (client) IP address: 0.0.0.0 (0.0.0.0)

Next server IP address: 0.0.0.0 (0.0.0.0)

Relay agent IP address: 0.0.0.0 (0.0.0.0)

**Client MAC address: Wistron\_23:68:8a (00:16:d3:23:68:8a)**

Server host name not given

Boot file name not given

Magic cookie: (OK)

Option: (t=53,l=1) **DHCP Message Type = DHCP Request**

Option: (61) Client identifier

Length: 7; Value: 010016D323688A;

Hardware type: Ethernet

Client MAC address: Wistron\_23:68:8a (00:16:d3:23:68:8a)

Option: (t=50,l=4) Requested IP Address = 192.168.1.101

Option: (t=12,l=5) Host Name = "nomad"

**Option: (55) Parameter Request List**

Length: 11; Value: 010F03062C2E2F1F21F92B

**1 = Subnet Mask; 15 = Domain Name**

**3 = Router; 6 = Domain Name Server**

**44 = NetBIOS over TCP/IP Name Server**

.....

request

Message type: **Boot Reply (2)**

Hardware type: Ethernet

Hardware address length: 6

Hops: 0

**Transaction ID: 0x6b3a11b7**

Seconds elapsed: 0

Bootp flags: 0x0000 (Unicast)

**Client IP address: 192.168.1.101 (192.168.1.101)**

Your (client) IP address: 0.0.0.0 (0.0.0.0)

**Next server IP address: 192.168.1.1 (192.168.1.1)**

Relay agent IP address: 0.0.0.0 (0.0.0.0)

Client MAC address: Wistron\_23:68:8a (00:16:d3:23:68:8a)

Server host name not given

Boot file name not given

Magic cookie: (OK)

**Option: (t=53,l=1) DHCP Message Type = DHCP ACK**

**Option: (t=54,l=4) Server Identifier = 192.168.1.1**

**Option: (t=1,l=4) Subnet Mask = 255.255.255.0**

**Option: (t=3,l=4) Router = 192.168.1.1**

**Option: (6) Domain Name Server**

Length: 12; Value: 445747E2445749F244574092;

IP Address: 68.87.71.226;

IP Address: 68.87.73.242;

IP Address: 68.87.64.146

**Option: (t=15,l=20) Domain Name = "hsd1.ma.comcast.net."**

reply



## Recap: how does an end-host gets an IP address

**Q:** How does a *host* get an IP address?

- hard-coded by system admin in a file
  - Windows: control-panel->network->configuration->tcp/ip->properties
  - UNIX: /etc/rc.config
- **DHCP: D**ynamic **H**ost **C**onfiguration **P**rotocol:  
dynamically get address from as server
  - “plug-and-play”

# IP addresses: how to get one?

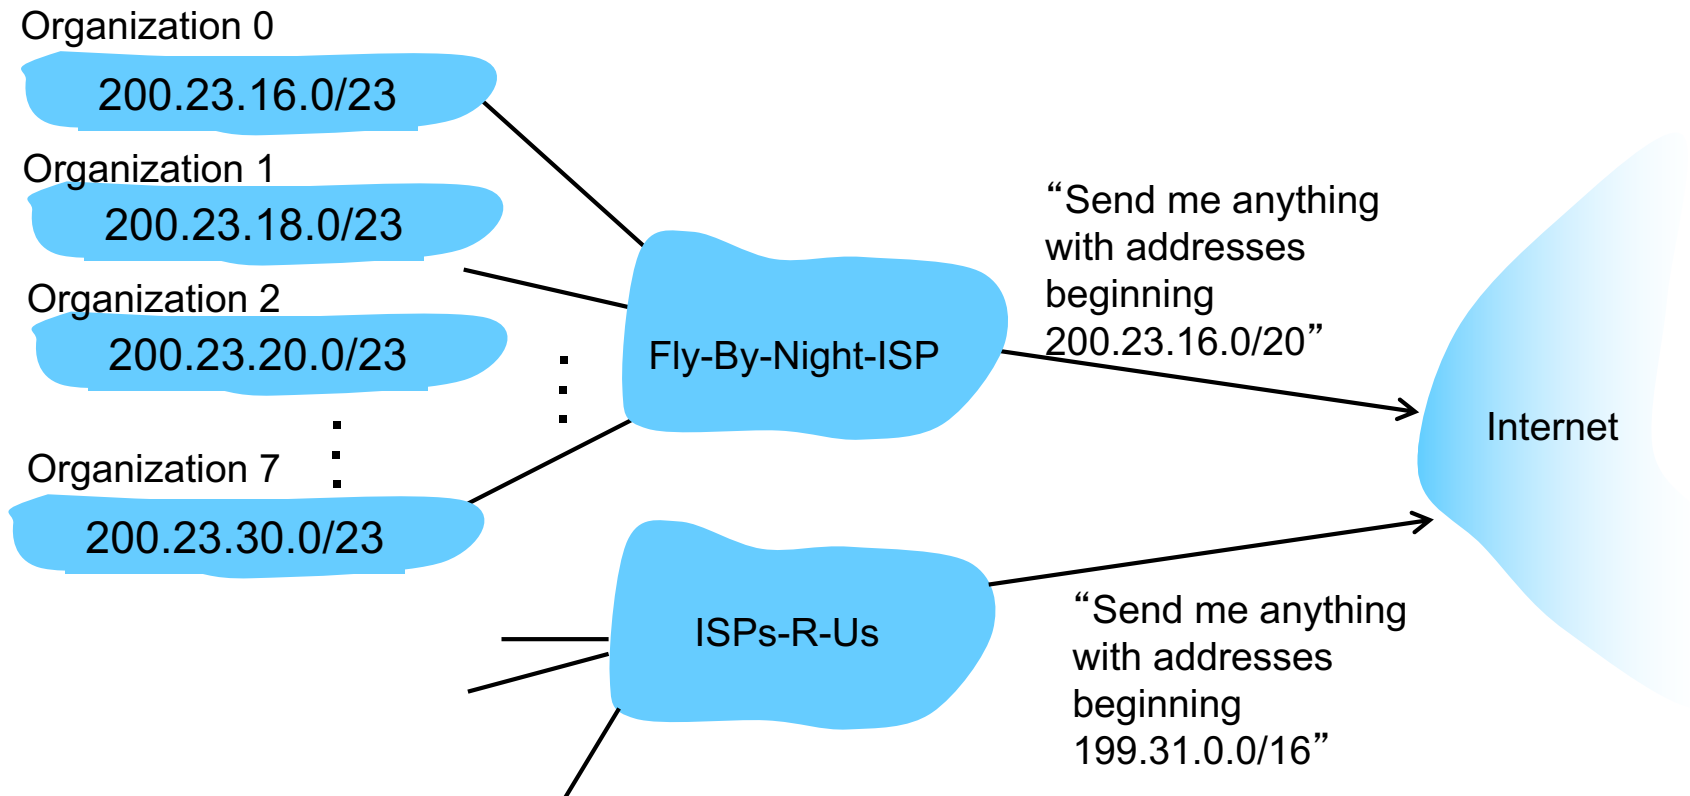
**Q:** how does a *network* get subnet part of IP addr?

**A:** gets allocated portion of its provider ISP' s address space

ISP's block	<u>11001000 00010111 00010000</u> 00000000	200.23.16.0/20
Organization 0	<u>11001000 00010111 00010000</u> 00000000	200.23.16.0/23
Organization 1	<u>11001000 00010111 00010010</u> 00000000	200.23.18.0/23
Organization 2	<u>11001000 00010111 00010100</u> 00000000	200.23.20.0/23
...	.....	....
Organization 7	<u>11001000 00010111 00011110</u> 00000000	200.23.30.0/23

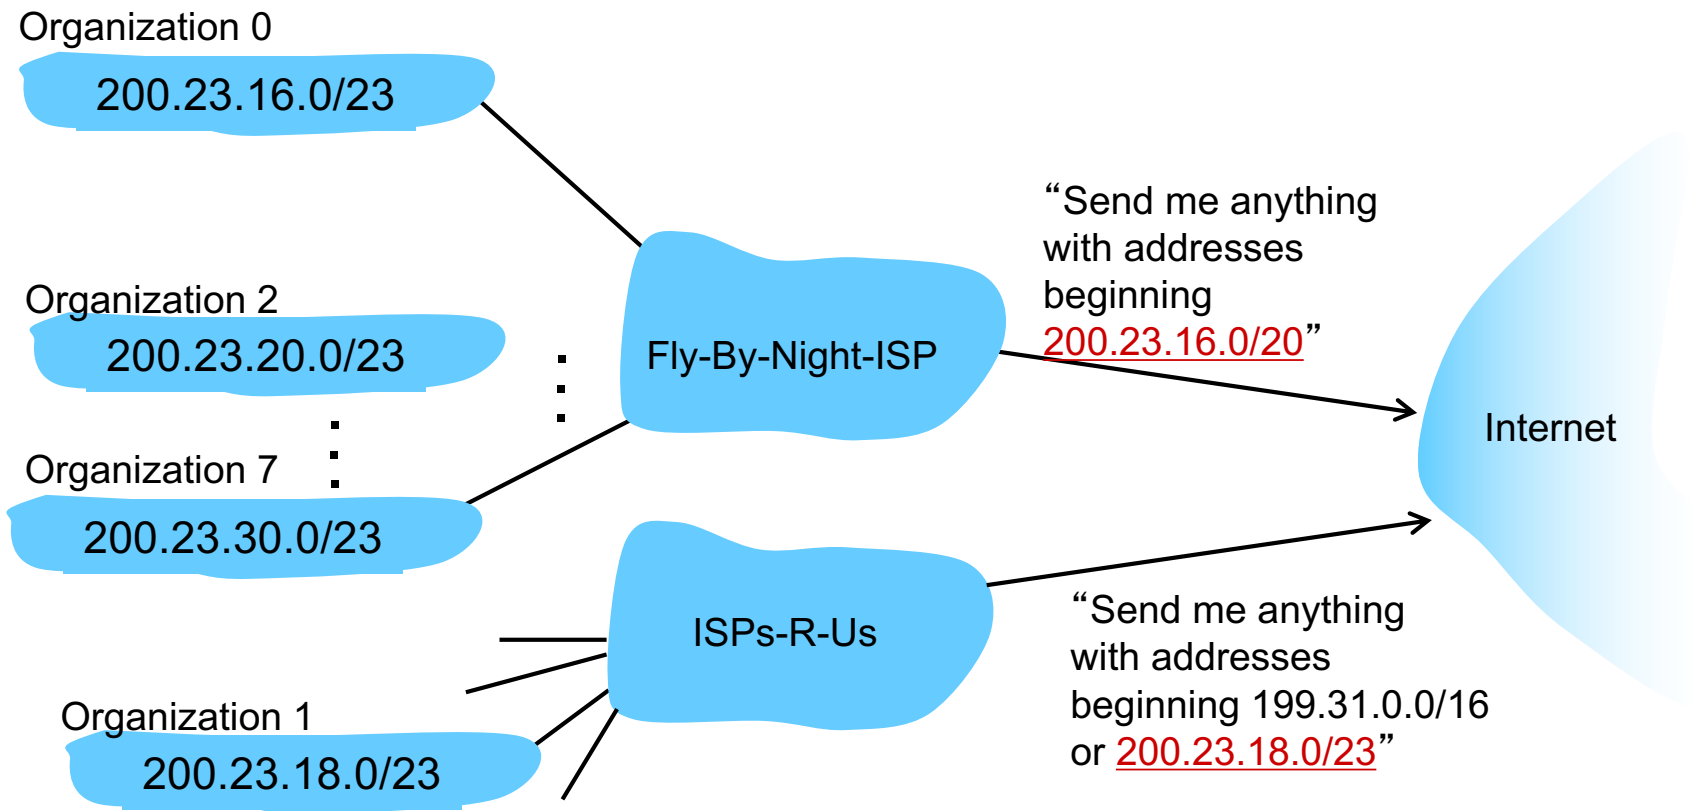
# Hierarchical addressing: route aggregation

hierarchical addressing allows efficient advertisement of routing information:



# Hierarchical addressing: more specific routes

- ISPs-R-U's has a more specific route to Organization 1
- Longest Prefix Match preferred



# IP addressing: the last word...

**Q:** how does *an ISP* get block of addresses?

**A:** ICANN: Internet Corporation for Assigned Names and Numbers <http://www.icann.org/>

- allocates addresses
  - <http://www.iana.org/assignments/ipv4-address-space/ipv4-address-space.xhtml>
- manages DNS
- assigns domain names, resolves disputes
- delegates to regional registries (RIRs)


# IP addresses: a scarce resource!

- More users than IP addresses
  - In the 1970's:  $2^{32}$  was plenty of addresses
  - ~2010: 1B computers, 5B phones
- Bad utilization
  - Some prefixes are full
  - Some are reserved and unused!
- How are they allocated today?
  - [http://en.wikipedia.org/wiki/List\\_of\\_assigned\\_/8\\_IPv4\\_address\\_blocks](http://en.wikipedia.org/wiki/List_of_assigned_/8_IPv4_address_blocks)
  - ARIN: [https://en.wikipedia.org/wiki/American\\_Registry\\_for\\_Internet\\_Numbers](https://en.wikipedia.org/wiki/American_Registry_for_Internet_Numbers)
  - <http://www.caida.org/research/id-consumption/whois-map/>

# IP addresses: a scarce resource!

- More users than IP addresses:
  - In the 1970's:  $2^{32}$  was plenty of addresses
  - Today: >1B computers, >7B phones
- Bad utilization of existing IP addresses
  - Some prefixes are full
  - Some are reserved and unused!
  - CAIDA's map
- **Solutions:**
  - **DHCP:** get an IP dynamically, when you use it
  - **NAT:** entire subnet uses one public IP externally
  - **IPv6:** increase the IP address from 32 to 64 bits.

# NAT: network address translation



*all* datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)



# NAT: network address translation

*motivation:* local network uses just one IP address as far as outside world is concerned:

- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

# NAT: network address translation

*implementation:* NAT router must:

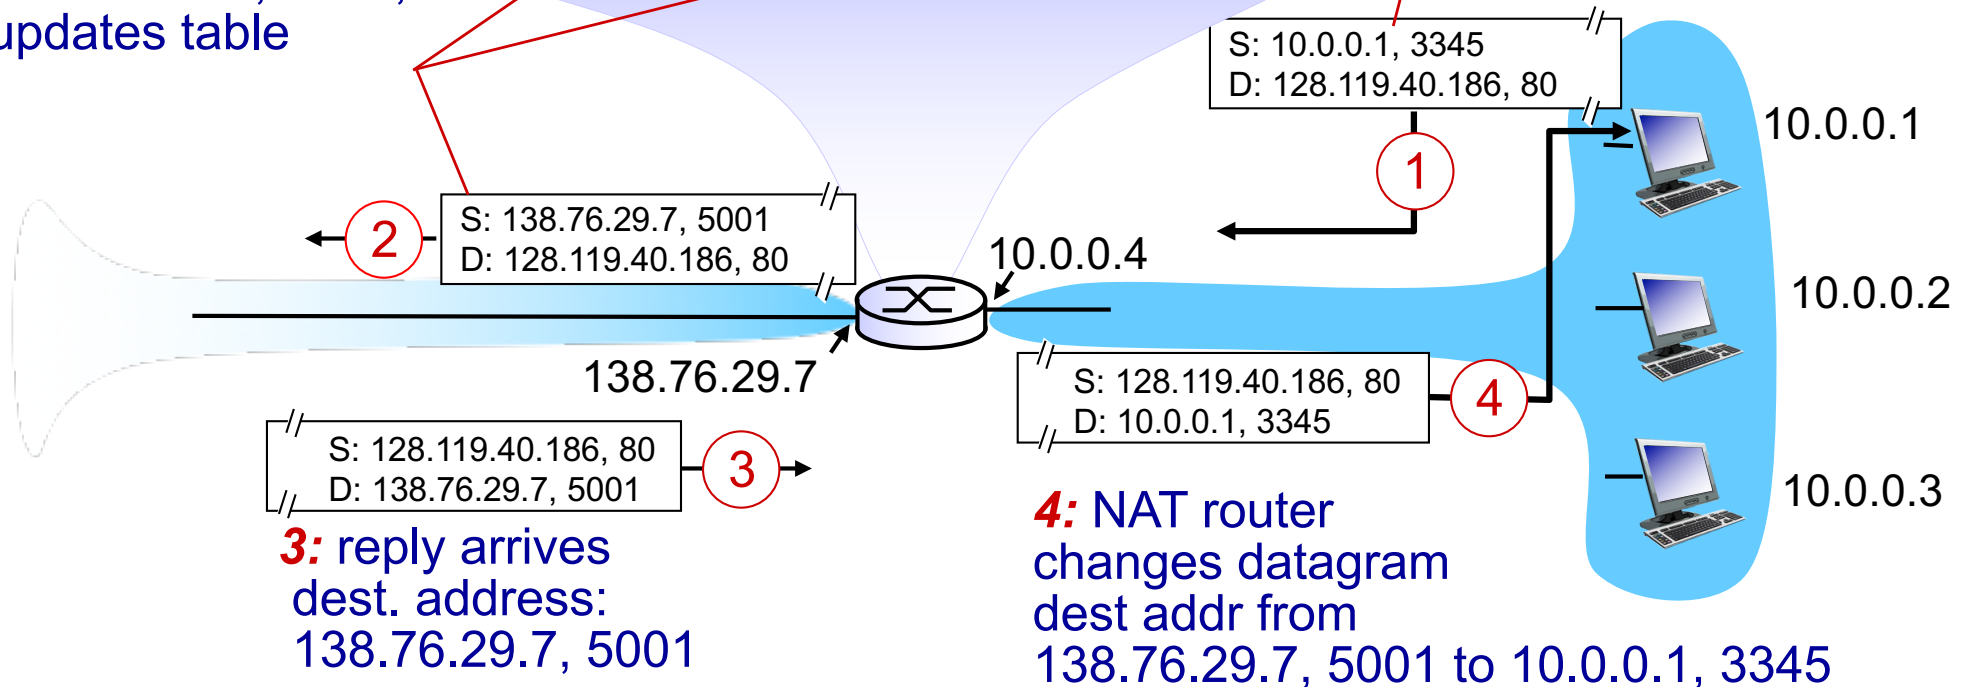
- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #) . . . remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

# NAT: network address translation

**2:** NAT router changes datagram source addr from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table

NAT translation table	
WAN side addr	LAN side addr
138.76.29.7, 5001	10.0.0.1, 3345
.....	.....

**1:** host 10.0.0.1 sends datagram to 128.119.40.186, 80




\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

# NAT: network address translation

- 16-bit port-number field:
  - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - routers should only process up to layer 3
  - address shortage should be solved by IPv6
  - violates end-to-end argument
    - NAT possibility must be taken into account by app designers, e.g., P2P applications
  - NAT traversal: what if client wants to connect to server behind NAT?

# NAT traversal problem

- client wants to connect to server with address 10.0.0.1
  - server address 10.0.0.1 local to LAN (client can't use it as destination addr)
  - only one externally visible NATed address: 138.76.29.7
- **solution 1:** statically configure NAT to forward incoming connection requests at given port to server
  - e.g., (138.76.29.7, port 2500) always forwarded to 10.0.0.1 port 25000




# NAT traversal problem

- *solution 2*: Universal Plug and Play (UPnP) Internet Gateway Device (IGD) Protocol. Allows NATed host to:

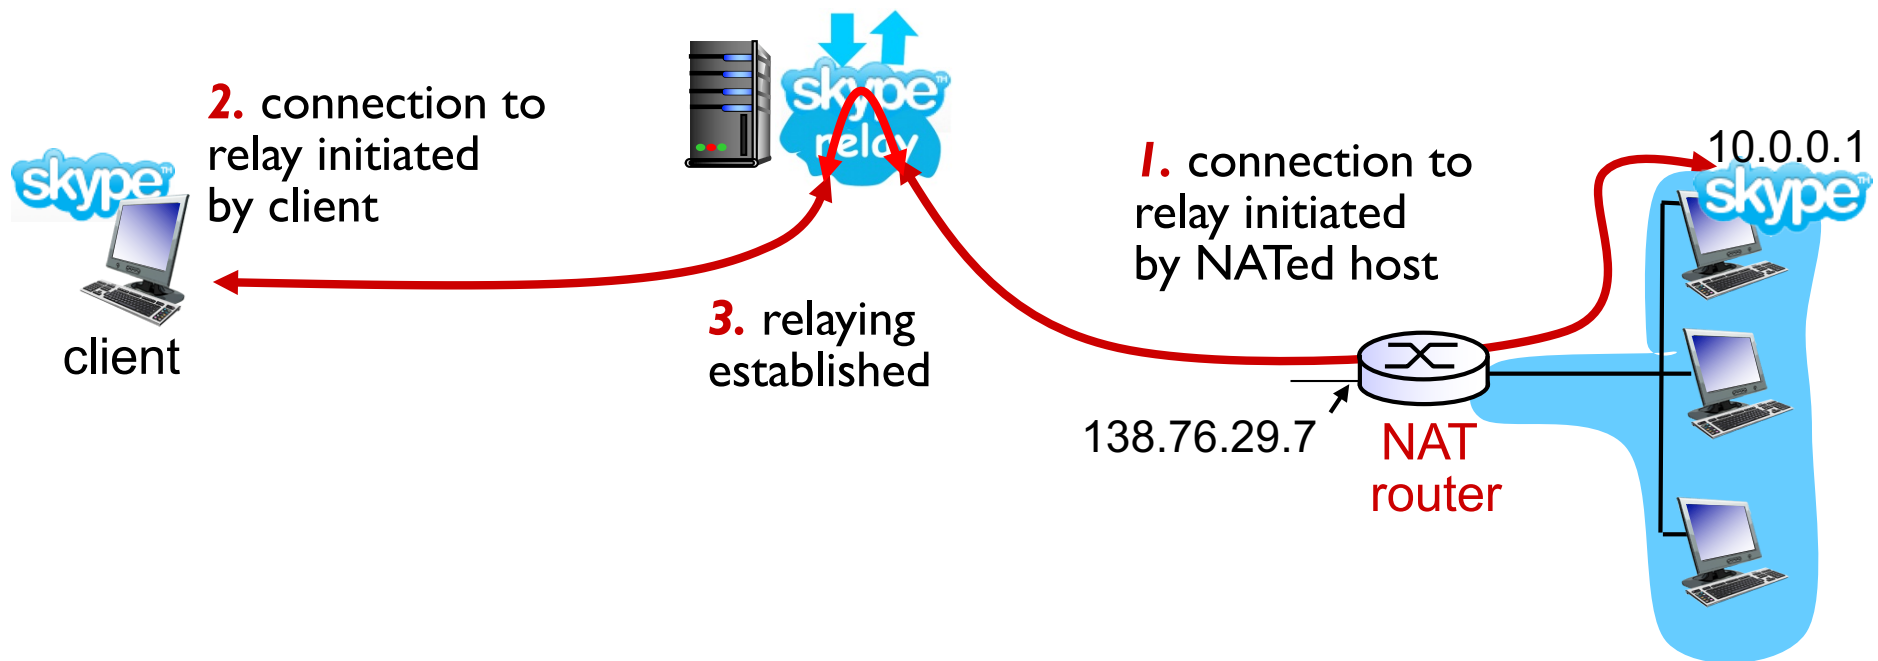
- ❖ learn public IP address (138.76.29.7)
- ❖ add/remove port mappings (with lease times)

i.e., automate static NAT port map configuration



# NAT traversal problem

- **solution 3:** relaying (used in Skype)
  - NATed client establishes connection to relay
  - external client connects to relay
  - relay bridges packets between to connections



# Chapter 4: outline

## 4.1 Overview of Network layer

- data plane
- control plane

## 4.2 What's inside a router

## 4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing
- network address translation
- IPv6

## 4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action



# IPv6: motivation

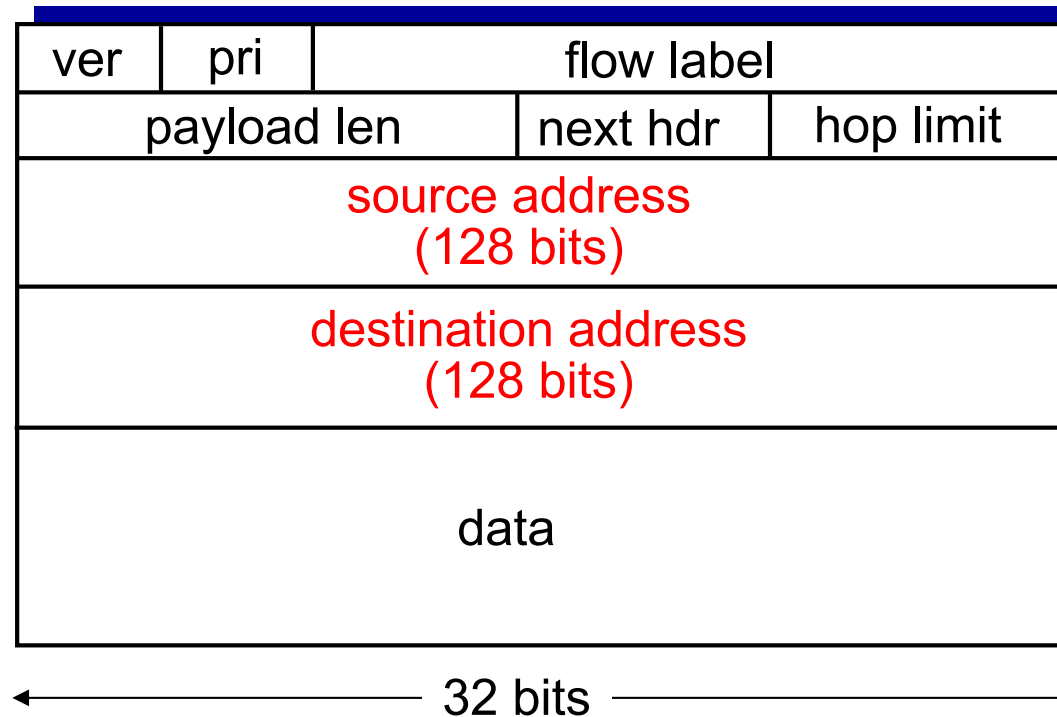
- *initial motivation*: 32-bit address space soon to be completely allocated.
- additional motivation:
  - header format helps speed processing/forwarding
  - header changes to facilitate QoS

## *IPv6 datagram format:*

- fixed-length 40 byte header
- no fragmentation allowed

# IPv6 datagram format

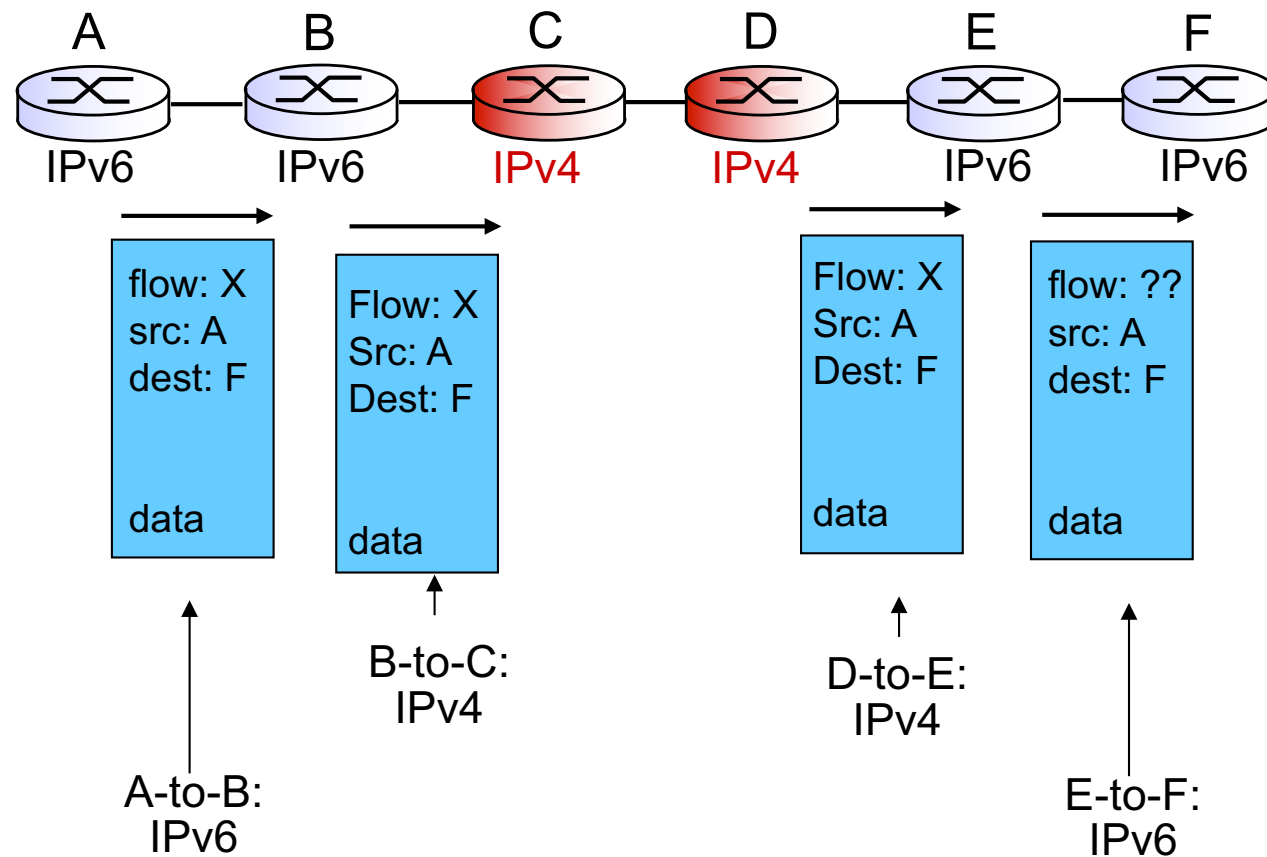
- *priority*: identify priority among datagrams in flow
- *flow Label*: identify datagrams in same “flow.” (concept of “flow” not well defined).
- *next header*: identify upper layer protocol for data
- *checksum*: removed entirely to reduce processing time at each hop
- *options*: allowed, but outside of header, indicated by “Next Header” field
- *ICMPv6*: new ICMP (additional message types, e.g. “Packet Too Big”, multicast group management)



# Transition from IPv4 to IPv6

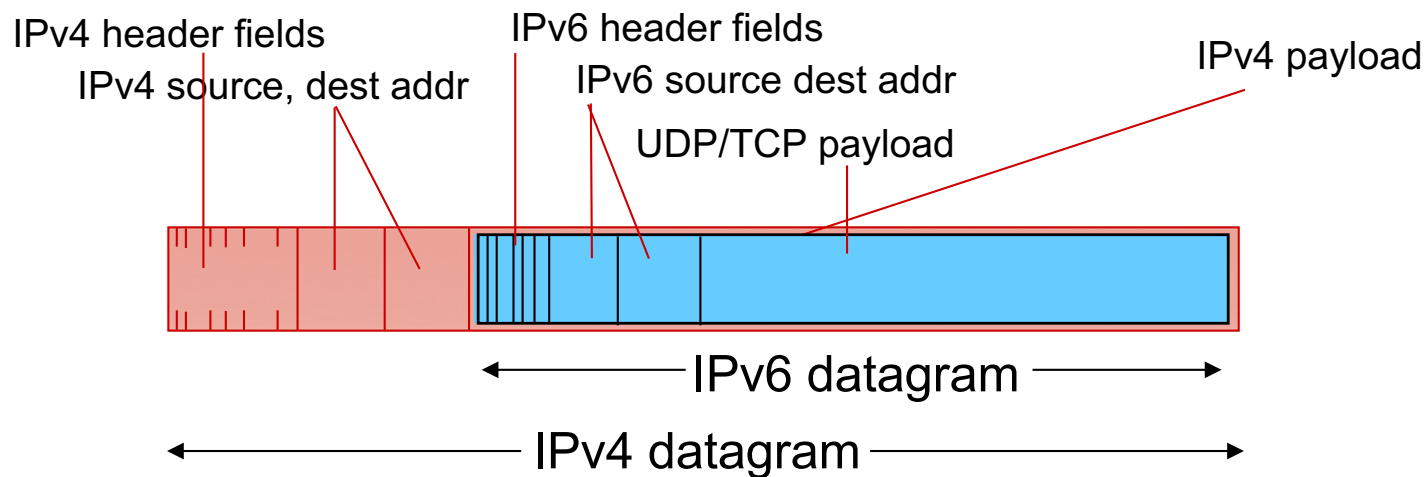
- Not all routers can be upgraded simultaneously
  - no “flag days”: Could we shut down the Internet on Friday on IPv4, upgrade all routers to IPv6, and restart on Monday?
- how will network operate with mixed IPv4 and IPv6 routers?
  - *Dual stack: IPv4/IPv6 [RFC4213]*
    - Have both *ipv4* (which defeats the purpose) and *ipv6* addresses
    - Talk *ipv4* to *ipv4* routers and *ipv6* to *ipv6* routers
    - Determine if other routers are *ipv4* or *ipv6* – capable through DNS: DNS can return different addresses based on the capabilities of requesting node
    - If *ipv4*-capable talks to *ipv6*-capable: convert *ipv6* → *ipv4*,
      - *ipv4* → *ipv6* will not recover some of the IPv6 fields

# Dual Stack




# Transition from IPv4 to IPv6

- Not all routers can be upgraded simultaneously
  - no “flag days”: Could we shut down the Internet on Friday on IPv4, upgrade all routers to IPv6, and restart on Monday?
- how will network operate with mixed IPv4 and IPv6 routers?
  - *Dual stack*: IPv4/IPv6 [RFC4213]
  - *Tunneling*: IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers




# Tunneling

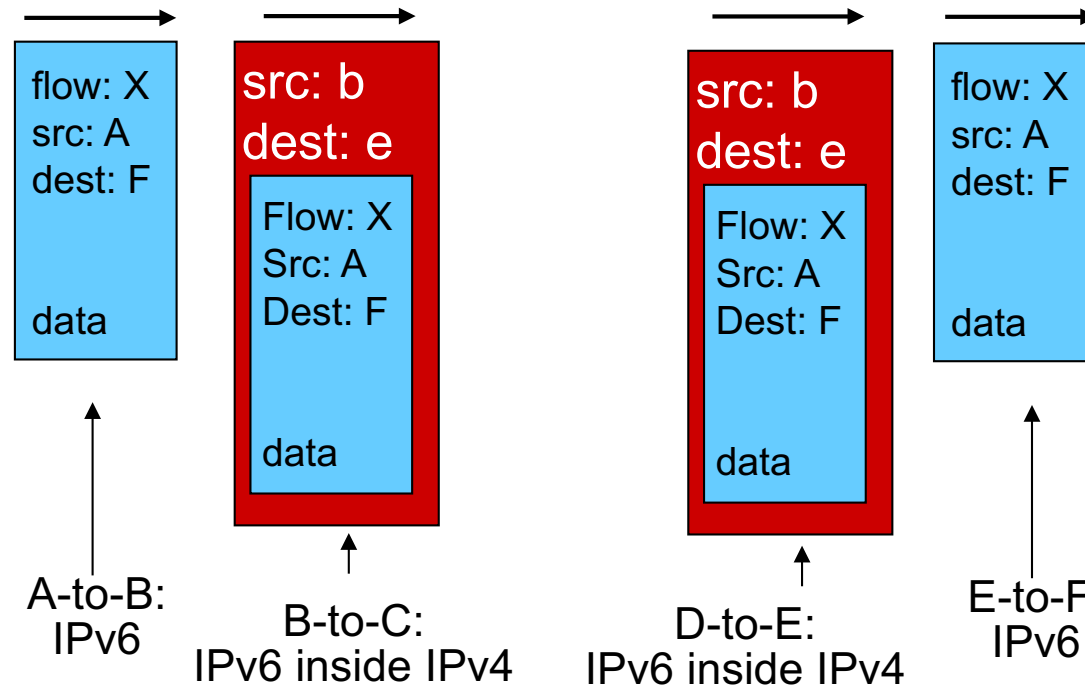



# Tunneling

logical view:



physical view:



# IPv6: adoption

- Google: 8% of clients access services via IPv6
- NIST: 1/3 of all US government domains are IPv6 capable
- Mobile, IoT:
- *Long (long!) time for deployment, use*
  - 20 years already and counting!
  - think of application-level changes in last 20 years: WWW, Facebook, streaming media, Skype, ...
  - *Why?*