

Chapter 5: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector

5.3 intra-AS routing in the Internet: OSPF


5.4 routing among the ISPs: BGP

5.5 The SDN control plane

5.6 ICMP: The Internet Control Message Protocol

5.7 Network management and SNMP

Bellman-Ford Algorithm (Dynamic Programming)



Let $d_v(t) :=$ cost of shortest path from v to t


Then $d_v(t) = \min_w \{ c(v,w) + d_w(t) \}$

\min taken over all neighbors w of v

$c(v,w)$ cost from v to neighbor w

$d_w(t)$ cost from neighbor w to destination t

Single destination “t” can be implicit



Let $d_v(t) :=$ cost of shortest path from v to t

Then $d_v(t) = \min_w \{ c(v,w) + d_w(t) \}$

cost from neighbor w to destination t
cost from v to neighbor w

\min taken over all neighbors w of v

Shortest path from v to t

- Each node v maintains local information
 - 1- $d(v)$: length of **shortest** path from v to t .
 - 2- w : the next hop after v , on the shortest path to t .
-
- Each node makes local decisions

$$d(v) = \min_{w:(v,w) \in E} \{ d(w) + c_{vw} \}$$





Figure 6.22 The minimum-cost path P from v to t using at most i edges.

Bellman-Ford equation example: $d_u(z)$




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

Compute $d_u(z)$ according to B-F:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

- $d_u(z)=4$ via next hop x
- Entry at Forwarding Table at u: next hop to **z** is **x**

Single node's operation



- v 's local view
 - v knows the costs $c(v,w)$ to all his neighbors w
 - v **exchanges updates** [when? how?] with this neighbors about their $d_w(t)$ (and his $d_v(t)$)
- v locally computes: $d_v(t) = \min_w \{c(v,w) + d_w(t)\}$
- v 's forwarding table entry
 - **$(t, w^*, d_v(t))$**
 - w^* : the next hop after v , on the shortest path to t
 - $d_v(t)$: length of shortest path from v to t
- Does the same for all possible destinations $t \in N$
 - v maintains distance **vector** $\mathbf{d}_v = [d_v(t): t \in N]$ to all $t \in N$.

Distance vector algorithm

key idea:

- from time-to-time [when?], each node v sends its own distance vector estimate to neighbors
- When v receives new DV estimate from neighbor, it updates its own DV using Bellman-Ford equation:

$$d_v(t) \leftarrow \min_w \{c(v,w) + d_w(t)\} \text{ for each node } t \in N$$

- ❖ under minor, natural conditions, the estimate $d_v(t)$ converge to the actual least cost $d_v(t)$

Distance vector algorithm

iterative, synchronous:

each node:

In each iteration



Exchange DV with all neighbors



Recompute estimates to
all destinations



Until the estimate does not
change

Example of Synchronous Execution: **distance to x**

node x table


	cost to
from	x
x	0

node y table

	cost to
from	x
y	2

node z table

	cost to
from	x
z	7



time

Network Layer

Example of Synchronous Execution: **distance to x**

node x table


	cost to
from	x
x	0

node y table

	cost to
from	x
y	2

node z table

	cost to
from	x
z	7



time

Network Layer

Example of Synchronous Execution: **distance to x**

$$D_y(x) = \min\{c(y,x) + D_x(x), c(y,z) + D_z(x)\}$$

$$= \min\{2+0, 1+7\} = 2$$

$$D_z(x) = \min\{c(z,x) + D_x(x), c(z,y) + D_y(x)\}$$

$$= \min\{7+0, 1+2\} = 3$$


node x table		cost to	
	X		
from x	0		

node y table		cost to	
	X		
from y	2		

node z table		cost to	
	X		
from z	7		

		cost to	
	X		
from x	0		
from y	2		
from z	7		

		cost to	
	X		
from x	0		
from y	2		
from z	3		



time

Network Layer

Example of Synchronous Execution: **distance to x**

node x table

	cost to
from	x
x	0

node y table

	cost to
from	x
y	2

node z table

	cost to
from	x
z	7

	cost to
from	x
x	0
y	2
z	7

	cost to
from	x
x	0
y	2
z	7

	cost to
from	x
x	0
y	2
z	3


	cost to
from	x
x	0
y	2
z	3

	cost to
from	x
x	0
y	2
z	3

	cost to	cost to	cost to
from	x	y	z
x	0		
y	2		
z	3		

$$D_z(x) = \min\{c(z,x) + D_x(x), c(z,y) + D_y(x)\}$$

$$= \min\{7+0, 1+2\} = 3$$



time

Network Layer

Example of Synchronous Execution: **all dest**

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

**node x
table**

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞


**node y
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

**node z
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0



		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0



time

Network Layer

Example of Synchronous Execution: **all dest**



Distance vector algorithm

iterative, asynchronous:

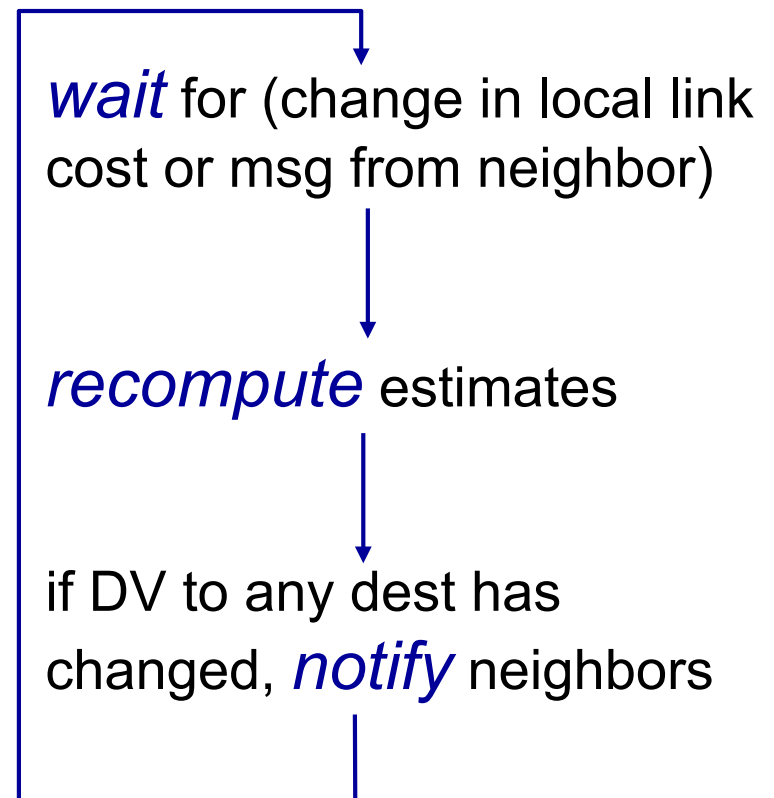
each local iteration
caused by:

- local link cost change
- DV update message from neighbor


Distributed, push-based:

- each node notifies neighbors *only* when its DV changes
 - neighbors then notify their neighbors if necessary

each node:



"Counting to Infinity" Example




- Originally:
 - B to C: via C, cost 1,
 - A to C: via B, cost 2
- After link failure:
 - B to C: via A, cost 3
 - A to C: via B, cost 4
 - B to C: via A, cost 5
 - A to C: via B, cost 6,
 - ++
- Fixes: (1) max hop count (2) poisonous reverse (3) keep path state.

Distance vector: link cost changes

link cost changes:

- ❖ node detects local link cost change
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors



“good
news
travels
fast”

t_0 : y detects link-cost change, updates its DV, informs its neighbors.

t_1 : z receives update from y , updates its table, computes new least cost to x , sends its neighbors its DV.


t_2 : y receives z 's update, updates its distance table. y 's least costs do *not* change, so y does *not* send a message to z .

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

Distance vector: link cost changes

link cost changes:

- ❖ node detects local link cost change
- ❖ *bad news travels slow* - “count to infinity” problem!
- ❖ 44 iterations before convergence
 - ❖ y to x: 4 via x
 - ❖ z to x: 5 via y
 - ❖ change of $c(x,y)$: $4 \rightarrow 60$
 - ❖ y to x: 6 via z
 - ❖ z to x: 7 via y
 - ❖ ...
 - ❖ Counting to 50



Solutions:

- ❖ *Poisoned Reverse*: if Z routes through Y to get to X, then Z advertises to Y that its (Z's) distance to X is infinite (so Y won't route to X via Z)
- ❖ Will this completely solve count to infinity problem?
 - ❖ No, if loops involve more than 2 nodes.
- ❖ Other fixes: (1) max hop count (2) keep path state.

Comparison of Link State and Distance Vector

- **In LS:** each node broadcasts to all other nodes information about its neighborhood (neighbors+link costs)
- **In DV:** each node talks only to its neighbors and tells them its estimates of path costs to all other destinations

message complexity

- ❖ with n nodes, E links
- ❖ **LS:** $O(nE)$ msgs sent
- ❖ **DV:** exchange between neighbors
 - convergence time varies

speed of convergence

- ❖ **LS:**
 - Naïve: $O(n^2)$
 - Efficient implementation: $O(m \log E)$
 - oscillations for load-dependent weights
- ❖ **DV:** convergence time varies
 - Efficient implementation: $O(mn)$ time, $O(m+n)$ memory
 - changes may cause routing loops
 - count-to-infinity problem

robustness: what if router malfunctions?



LS:

- node can advertise incorrect *link* cost
- each node computes only its own table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate thru network

Practice examples



More practice on routing algorithms

- Find shortest paths from s to all other nodes, using Dijkstra
- Find shortest paths from all nodes to T, using Bellman-Ford
 - Synchronous execution, asynchronous execution
 - **Consider that a link fails or changes cost. What happens?**
- Discussion session. Interactive exercises. HW4

